

X10: a High-Productivity Approach to High Performance Programming

Rajkishore Barik
Christopher Donawa
Matteo Frigo
Allan Kielstra
Vivek Sarkar


HPC Challenge Class 2 Award Submission

This work has been supported in part by the Defense Advanced Research Projects Agency (DARPA) under contract No. NBCH30390004.




Motivation: Productivity Challenges caused by Future Hardware Trends

Clusters → Global Address Space

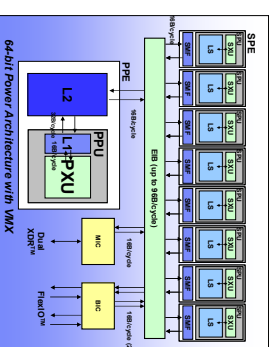


Challenge: Develop new language, compiler and tools technologies to support productive portable parallel abstractions for future hardware


Homogeneous Multi-core



Heterogeneous Accelerators



X10 Programming Model



Storage classes:

- **Activity-local**
- **Place-local**
- **Partitioned global**
- **Immutable**

- Dynamic parallelism with a *Partitioned Global Address Space*
- *Places* encapsulate binding of activities and globally addressable data
- All concurrency is expressed as *asynchronous activities* – subsumes threads, structured parallelism, messaging, DMA transfers (beyond SPMD)
- *Atomic sections* enforce mutual exclusion of co-located data
 - No place-remote accesses permitted in atomic section
- *Immutable* data offers opportunity for single-assignment parallelism




3

Deadlock safety: any X10 program written with *async*, *atomic*, *finish*, *foreach*, *attach*, and *clocks* can never deadlock


High-Productivity, High-Performance Programming with X10




X10 Deployment




Homogeneous Multi-core



Heterogeneous Accelerators



Clusters

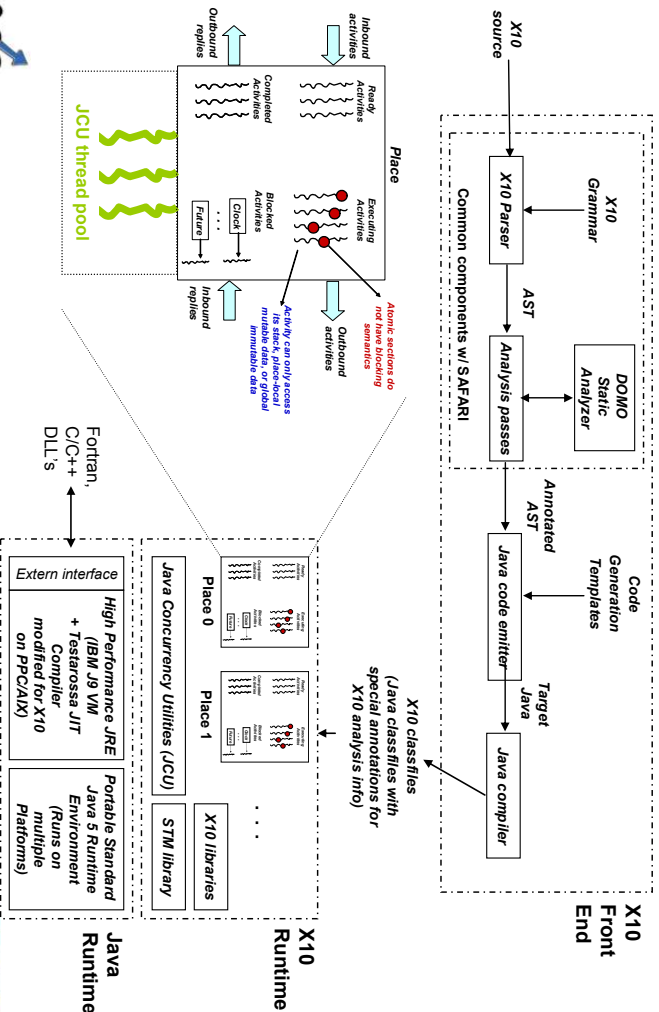


4

High-Productivity, High-Performance Programming with X10



Current Status: Multi-core SMP Implementation for X10



5

High-Productivity, High-Performance Programming with X10



System Configuration used for Performance Results

- Hardware
 - STREAM (C/OpenMP & X10), RandomAccess (C/OpenMP & X10), FFT (X10)
 - 64-core POWER5+, p595+, 2.3 GHz, 512 GB (r28n01.pbn.ilst.com)
 - FFT (Cilk version)
 - 16-core POWER5+, p570, 1.9 GHz
 - All runs performed with page size = 4KB and SMT turned off
- Operating System
 - AIX v5.3
- Compiler
 - xlc v7.0.0.5 w/ -O3 option (also qsmc=omp for OpenMP compilation)
- X10
 - Dynamic compilation options: -J-Xjit:count=0,opt:level=veryHot
 - X10 activities use serial libraries written in C and linked with X10 runtime
 - Data size limitation: current X10 runtime is limited to a max heap size of 2GB
 - All results reported are for runs that passed validation
 - Caveat: these results should *not* be treated as official benchmark measurements of the above systems



6

High-Productivity, High-Performance Programming with X10



STREAM

OpenMP / C version

```
#pragma omp parallel for
for (j=0; j<N; j++) {
    b[j] = scalar*c[j];
}
```

Hybrid X10 + Serial C version

```
finish ateach(point p : dist.factory.unique()) {
    final region myR = (D | here).region;
    scale(b, scalar,c,myR.rank(0).low(),myR.rank(0).high()+1);
}
```



7

High-Productivity, High-Performance Programming with X10



STREAM

OpenMP / C version

```
#pragma omp parallel for
for (j=0; j<N; j++) {
    b[j] = scalar*c[j];
}
```

Traversing array region
can be error-prone

Implicitly assumes Uniform
Memory Access model
(no distributed arrays)

SL OC counts are comparable

Hybrid X10 + Serial C version

```
finish ateach(point p : dist.factory.unique()) {
    final region myR = (D | here).region;
    scale(b, scalar,c,myR.rank(0).low(),myR.rank(0).high()+1);
}
```

scale() is a sequential C function

Restrict operator simplifies
computation of local region




8

High-Productivity, High-Performance Programming with X10



Performance Results for STREAM

Array size = 2²⁶ elements
Combined memory for 3 arrays = 1.5GB



High-Productivity, High-Performance Programming with X10



RandomAccess

OpenMP / C version

```
#define NUPDATE (4 * TableSize)
for (i=0; i<NUPDATE/128; i++) {
#pragma omp parallel for
    for (j=0; j<128; j++) {
        ran[j] = (ran[j] << 1) ^ ((s64int) ran[j] < 0 ? POLY : 0);
        Table[ran[j]] & (TableSize-1)] ^= ran[j];
    }
}
```

Hybrid X10 + Serial C version

```
finish ateach(point p : dist.factory.unique()) {
    final region myR = (D | here).region;
    for (int i=0; i<(4 * TableSize)/W; i++) {
        innerLoop(Table, TableSize, ran, myR.ran(0).low(), myR.ran(0).high()+1);
    }
}
```



High-Productivity, High-Performance Programming with X10



RandomAccess

OpenMP / C version

```
#define NUPDATE (4 * TableSize)
for (i=0; i<NUPDATE/128; i++) {
#pragma omp parallel for
for (j=0; j<128; j++) {
    ran[j] = (ran[j] << 1) ^ ((s64Int) ran[j] < 0 ? POLY : 0);
    Table[ran[j] & (TableSize-1)] ^= ran[j];
}
}
```

Inner parallel loop is a source of inefficiency in OpenMP version

SLOC counts are comparable

Multi-place version designed to run unchanged on an SMP or a cluster

Hybrid X10 + Serial C version

```
finish ateach(point p : dist.factory.unique()) {
final region myR = (D | here).region;
for (int i=0; i<(4 * TableSize)/W; i++) {
    innerLoop(Table, TableSize, ran, myR.rank(0).low(), myR.rank(0).high()+1);
}
}
```


innerLoop() is a sequential C function

Restrict operator simplifies computation of local region

PERCS 11 High-Productivity, High-Performance Programming with X10

Performance Results for RandomAccess

Array size = 1.8GB



FFT: Transpose example

Cilk / C version (Recursive version)

```
#define SUB(A, i, j) [(i)*SQRN+(j)]
cilk void transpose(fftw_complex *A, int n)
{
    if (n > 1) {
        int n2 = n/2;
        spawn transpose(A, n2);
        spawn transpose(&SUB(A, n2, n2), n-n2);
        spawn transpose_and_swap(A, 0, n2, n2, n);
    } else {
        /* X1 transpose is a NOP */
    }
}
```

Implicit sync at function boundary

Hybrid X10 + Serial C version (Non-recursive version)


```
int nBlocks = SQRN / bSize;
int p = 0;
finish for (int x = 0; x < nBlocks; ++x) {
    for (int c = r; c < nBlocks; ++c) { // Triangular loop
        final int topLeftta_r = (bSize * x);
        final int topLeftta_c = (bSize * c);
        final int topLefttb_r = (bSize * c);
        final int topLefttb_c = (bSize * x);
        async (place.factory.place(p++))
            transpose_and_swap(A, topLeftta_r, topLeftta_c, topLefttb_r, topLefttb_c, bSize);
    }
}
```

"finish" operator is used to wait for termination of all subactivities (async's)

transpose_and_swap() is a sequential C function

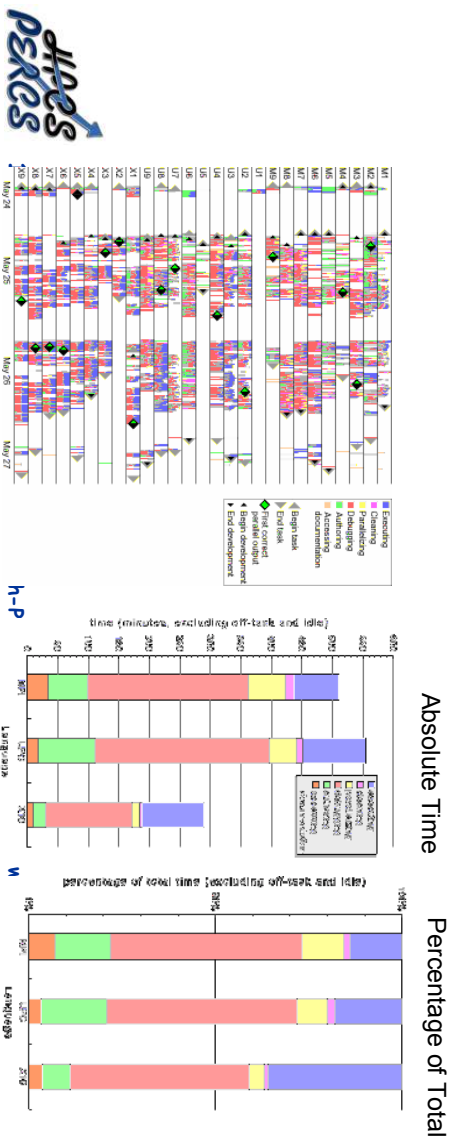
Performance Results for FFT (w/ memoized sine/cosine twiddle factors)

$$N = 2^{24} \text{ (SQRN} = 2^{12}\text{)}$$



Summary

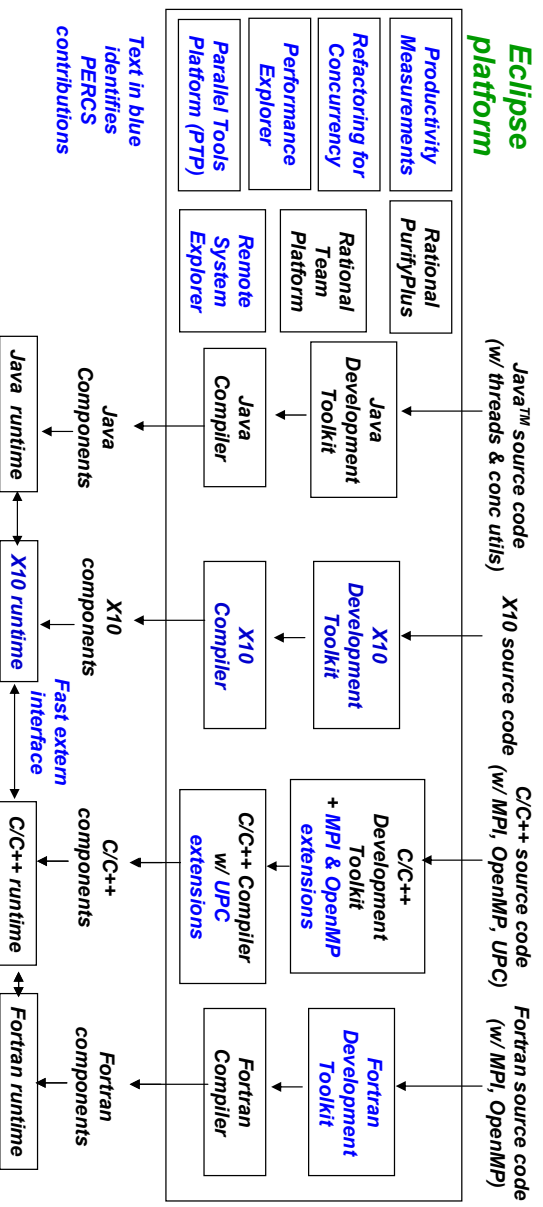
- X10 programming model provides core concurrency and distribution constructs for new era of parallel processing
- Results show competitive performance for Hybrid X10+C relative to OpenMP/C and Cilk
- Past studies have shown other productivity benefits of X10
- To find out more, come to the X10 exhibit in the Exotic Technologies area!



BACKUP SLIDES START HERE

X10 context: PERCS Programming Model, Tools and Compilers

(PERCS = Productive Easy-to-use Reliable Computer System)

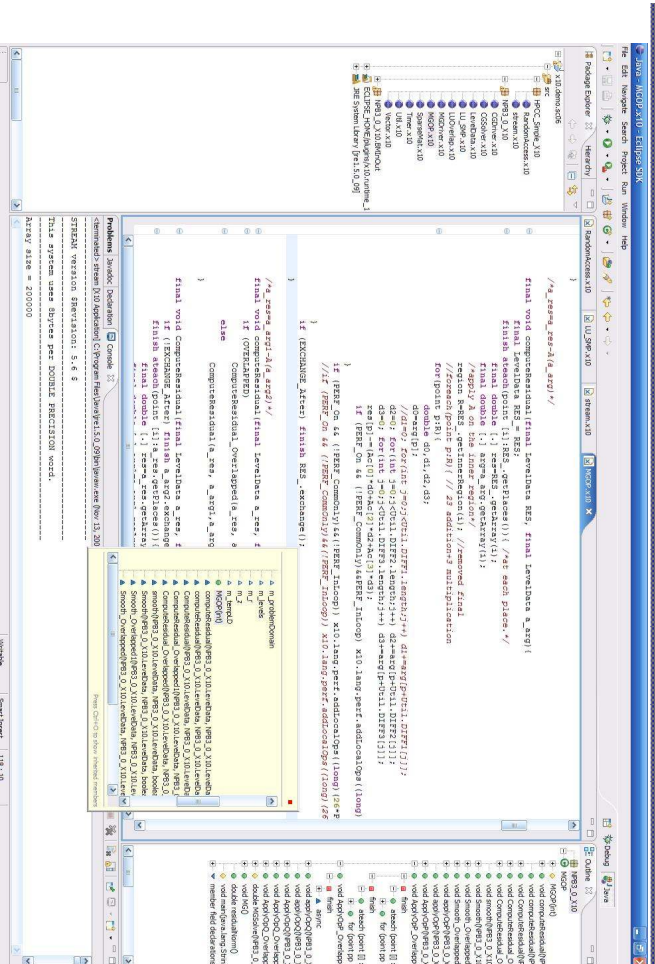


17

High-Productivity, High-Performance Programming with X10



X10 Eclipse Development Toolkit

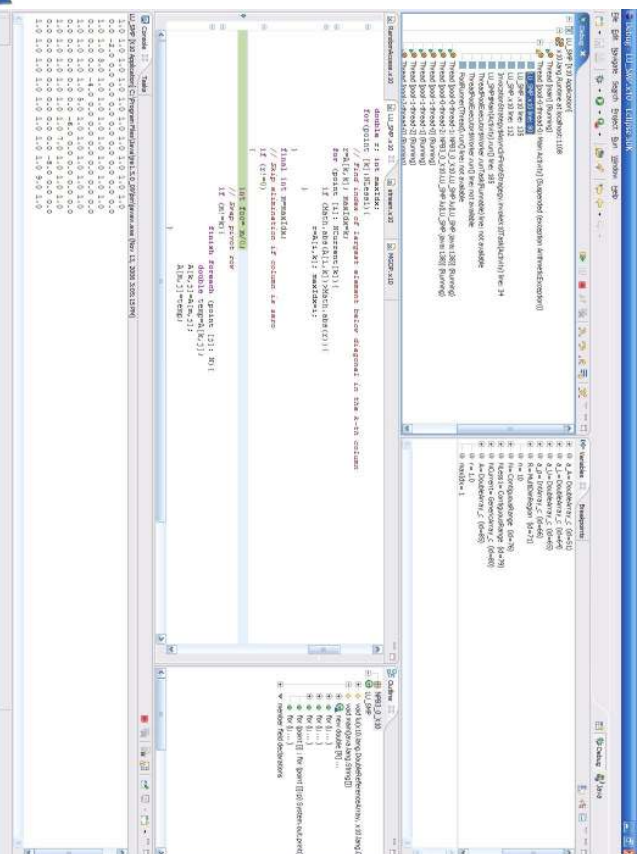


18

High-Productivity, High-Performance Programming with X10



X10 Eclipse Debugging Toolkit



19

High-Productivity, High-Performance Programming with X10



X10 Language

- **async** [(Place)] [clocked(c...)] *Stm*
 - Run *Stm* asynchronously at Place
- **finish** *Stm*
 - Execute *s*, wait for all asyncs to terminate (generalizes join)
- **foreach** (point *P* : *Reg*) *Stm*
 - Run *Stm* asynchronously for each point in *region*
- **ateach** (point *P* : *Dist*) *Stm*
 - Run *Stm* asynchronously for each point in *dist*, in its place.
- **atomic** *Stm*
 - Execute *Stm* atomically
- new *T*
 - Allocate object at this place (here)
- new *T*[*d*] / new *T* value [*d*]
 - Array of base type *T* and distribution *d*
- **future** [(Place)] [clocked(c...)] *Expr*
 - Compute *Expr* asynchronously at Place
- **F. force**()
 - Block until future *F* has been computed
- **extern**
 - Lightweight interface to native code



20

High-Productivity, High-Performance Programming with X10



Deadlock safety: any X10 program written with above constructs (excluding future) can never deadlock

- Can be extended to restricted cases of using future

X10 Arrays, Regions, Distributions

ArrayExpr:

```
new ArrayType ( Formal ) { Stm }
Distribution Expr
ArrayExpr [ Region ]
ArrayExpr | Distribution
ArrayExpr || ArrayExpr
ArrayExpr.overlay(ArrayExpr)
ArrayExpr.scan( fun l, ArgList )
ArrayExpr.reduce( fun l, ArgList )
ArrayExpr.lift( fun l, ArgList )
```

Region:

```
Expr : Expr
[ Range, ..., Range ]
Region && Region
Region || Region
Region - Region
BuiltinRegion
-- 1-D region
-- Multidimensional Region
-- Intersection
-- Union
-- Set difference
```

Dist:

```
Region -> Place
Distribution | Place
Distribution | Region
Distribution || Distribution
Distribution - Distribution
Distribution.overlay ( Distribution )
BuiltinDistribution
-- Constant distribution
-- Restriction
-- Restriction
-- Union
-- Set difference
```

ArrayType:

```
Type [Kind] [ ]
Type [Kind] [ region(N) ]
Type [Kind] [ Region ]
Type [Kind] [ Distribution ]
```

Language supports type safety, memory safety, place safety, clock safety.