

Report on the Programming Language X10

Version 2.0.2

DRAFT — March 11, 2010

Vijay Saraswat

Please send comments to vsaraswa@us.ibm.com

March 11, 2010

This report provides a description of the programming language X10. X10 is a class-based object-oriented programming language designed for high-performance, high-productivity computing on high-end computers supporting $\approx 10^5$ hardware threads and $\approx 10^{15}$ operations per second.

X10 is based on state-of-the-art object-oriented programming languages and deviates from them only as necessary to support its design goals. The language is intended to have a simple and clear semantics and be readily accessible to mainstream OO programmers. It is intended to support a wide variety of concurrent programming idioms.

The X10 design team consists of Bard Bloom, Ganesh Bikshandi, David Cunningham, Robert Fuhrer, David Grove, Sreedhar Kodali, Nathaniel Nystrom, Igor Peshansky, Vijay Saraswat, Olivier Tardieu.

Past members include Shivali Agarwal, David Bacon, Raj Barik, Bob Blainey, Philippe Charles, Perry Cheng, Christopher Donawa, Julian Dolby, Kemal Ebcioglu, Patrick Gallop, Christian Grothoff, Allan Kielstra, Sriram Krishnamoorthy, Bruce Lucas, Vivek Sarkar, Armando Solar-Lezama, S. Alexander Spoon, Sayantan Sur, Christoph von Praun, Pradeep Varma, Krishna Nandivada Venkata, Jan Vitek, and Tong Wen.

For extended discussions and support we would like to thank: Gheorghe Almasi, Robert Blackmore, Robert Callahan, Calin Cascaval, Norman Cohen, Elmootaz Elnozahy, John Field, Kevin Gildea, Chulho Kim, Orren Krieger, Doug Lea, John McCalpin, Paul McKenney, Andrew Myers, Filip Pizlo, Ram Rajamony, R. K. Shyamasundar, V. T. Rajan, Frank Tip, Mandana Vaziri, and Hanhong Xue.

Special thanks is due to Bard Bloom who meticulously read every section and provided detailed feedback. This document is much better because of his diligence and hard work.

We thank Jonathan Rhees and William Clinger with help in obtaining the \LaTeX style file and macros used in producing the Scheme report, on which this document is based. We acknowledge the influence of the JavaTM Language Specification [5].

This document revises Version 1.7 of the Report, released in September 2008. It documents the language corresponding to Version 2.0 of the implementation. Version 1.7 of the report was co-authored by Nathaniel Nystrom. The design of structs in X10 was led by Olivier Tardieu and Nathaniel Nystrom.

Earlier implementations benefited from significant contributions by Raj Barik, Philippe Charles, David Cunningham, Christopher Donawa, Robert Fuhrer, Christian Grothoff, Nathaniel Nystrom, Igor Peshansky, Vijay Saraswat, Vivek Sarkar, Olivier Tardieu, Pradeep Varma, Krishna Nandivada Venkata, and Christoph von Praun. Tong Wen has written many application programs in X10. Guojing Cong has helped in the development of many applications. The implementation of generics in X10 was influenced by the implementation of PolyJ [2] by Andrew Myers and Michael Clarkson.

Contents

1	Introduction	9
2	Overview of X10	12
2.1	Object-oriented features	12
2.2	The sequential core	14
2.3	Places and activities	15
2.4	Clocks	17
2.5	Arrays, regions and distributions	18
2.6	Annotations	18
2.7	Translating MPI programs to X10	19
2.8	Summary and future work	19
2.8.1	Design for scalability	19
2.8.2	Design for productivity	19
2.8.3	Conclusion	20
3	Lexical structure	22
4	Types	24
4.1	Classes and interfaces	26
4.1.1	Class types	26
4.1.2	Interface types	27
4.1.3	Properties	27
4.2	Type parameters	28
4.2.1	Generic types	28
4.3	Type definitions	30
4.4	Constrained types	31
4.4.1	Constraints	31
4.4.2	Place constraints	32
4.4.3	Constraint semantics	33

4.4.4	Consistency of dependent types	34
4.5	Function types	36
4.6	Annotated types	37
4.7	Subtyping and type equivalence	37
4.8	Least common ancestor of types	38
4.9	Coercions and conversions	39
4.9.1	Coercions	39
4.9.2	Conversions	40
4.10	Built-in types	41
4.10.1	The class <code>Object</code>	41
4.10.2	The class <code>String</code>	42
4.10.3	Array types	42
4.10.4	Rails	43
4.10.5	Future types	43
4.11	Type inference	43
4.11.1	Variable declarations	44
4.11.2	Return types	44
4.11.3	Type arguments	44
5	Variables	46
5.1	Final variables	47
5.2	Initial values of variables	47
5.3	Destructuring syntax	48
5.4	Formal parameters	48
5.5	Local variables	48
5.6	Fields	49
5.7	Properties	50
6	Objects	51
6.1	Basic Design	51
6.2	Examples	54
6.2.1	Programming Methodology	55
7	Names and packages	56
8	Interfaces	57
9	Classes	58

9.1	Type invariants	60
9.2	<code>implements</code> and <code>extends</code> clauses	62
9.3	Constructor definitions	62
9.4	<code>proto</code> qualifier on types	64
9.5	Field definitions	68
9.5.1	Field hiding	68
9.5.2	Field qualifiers	69
9.6	Method definitions	69
9.6.1	Property methods	70
9.6.2	Method overloading, overriding, hiding, shadowing and obscuring	71
9.6.3	Method qualifiers	72
9.7	Static initialization	74
10	Structs	76
10.1	Struct declaration	77
10.2	Boxing of structs	78
10.3	Implementation of <code>Any</code> methods	78
10.4	“Primitives”	79
10.5	Generic programming with structs	79
10.6	Programming Methodology	79
10.6.1	Compatibility Note	80
10.6.2	Examples	80
11	Functions	82
11.1	Overview	82
11.2	Function Literals	83
11.2.1	Outer variable access	84
11.3	Methods selectors	85
11.4	Operator functions	86
11.5	Functions as objects of type <code>Any</code>	87
12	Expressions	88
12.1	Literals	88
12.2	<code>this</code>	89
12.3	Local variables	89
12.4	Field access	89
12.5	Function Literals	90

12.6	Calls	90
12.7	Assignment	91
12.8	Increment and decrement	92
12.9	Numeric promotion	92
12.10	Unary plus and unary minus	92
12.11	Bitwise complement	93
12.12	Binary arithmetic operations	93
12.13	Binary shift operations	93
12.14	Binary bitwise operations	94
12.15	String concatenation	94
12.16	Logical negation	94
12.17	Boolean logical operations	94
12.18	Boolean conditional operations	95
12.19	Relational operations	95
12.20	Conditional expressions	95
12.21	Stable equality	96
12.22	Allocation	96
12.23	Casts	97
12.24	instanceof	97
12.25	Subtyping expressions	98
12.26	Contains expressions	98
12.27	Rail constructors	98
13	Statements	100
13.1	Empty statement	100
13.2	Local variable declaration	100
13.3	Block statement	101
13.4	Expression statement	101
13.5	Labeled statement	101
13.6	Break statement	101
13.7	Continue statement	102
13.8	If statement	103
13.9	Switch statement	103
13.10	While statement	104
13.11	Do-while statement	104
13.12	For statement	104
13.13	Throw statement	105
13.14	Try-catch statement	106

13.15	Return statement	106
14	Places	107
14.1	Place expressions	108
14.2	here	108
15	Activities	110
15.1	The X10 rooted exception model	111
15.2	Spawning an activity	111
15.3	Place changes	113
15.4	Finish	113
15.5	Initial activity	114
15.6	Foreach statements	115
15.7	Ateach statements	115
15.8	Futures	115
15.9	At expressions	117
15.10	Shared variables	117
15.11	Atomic blocks	117
15.11.1	Unconditional atomic blocks	117
15.11.2	Conditional atomic blocks	119
16	Clocks	122
16.1	Clock operations	123
16.1.1	Creating new clocks	123
16.1.2	Registering new activities on clocks	124
16.1.3	Resuming clocks	124
16.1.4	Advancing clocks	125
16.1.5	Dropping clocks	125
16.2	Program equivalences	125
17	Arrays	126
17.1	Points	126
17.2	Regions	127
17.2.1	Operations on regions	128
17.3	Distributions	129
17.3.1	Operations returning distributions	129
17.3.2	User-defined distributions	131
17.3.3	Operations on distributions	131

17.3.4	Example	131
17.4	Array initializer	132
17.5	Operations on arrays	133
17.5.1	Element operations	133
17.5.2	Constant promotion	133
17.5.3	Restriction of an array	133
17.5.4	Assembling an array	134
17.5.5	Global operations	134
18	Annotations and compiler plugins	136
18.1	Annotation syntax	136
18.2	Annotation declarations	138
18.3	Compiler plugins	138
19	Linking with native code	141
	Alphabetic index of definitions of concepts, keywords, and procedures	146
A	Change Log	153
A.1	Changes from X10 v2.0	153
A.2	Changes from X10 v1.7	153

1 Introduction

Background

Larger computational problems require more powerful computers capable of performing a larger number of operations per second. The era of increasing performance by simply increasing clocking frequency now seems to be behind us. It is becoming increasingly difficult to manage chip power and heat. Instead, computer designers are starting to look at *scale out* systems in which the system's computational capacity is increased by adding additional nodes of comparable power to existing nodes, and connecting nodes with a high-speed communication network.

A central problem with scale out systems is a definition of the *memory model*, that is, a model of the interaction between shared memory and simultaneous (read, write) operations on that memory by multiple processors. The traditional “one operation at a time, to completion” model that underlies Lamport's notion of *sequential consistency* (SC) proves too expensive to implement in hardware, at scale. Various models of *relaxed consistency* have proven too difficult for programmers to work with.

One response to this problem has been to move to a *fragmented memory model*. Multiple processors are made to interact via a relatively language-neutral message-passing format such as MPI [10]. This model has enjoyed some success: several high-performance applications have been written in this style. Unfortunately, this model leads to a *loss of programmer productivity*: the message-passing format is integrated into the host language by means of an application-programming interface (API), the programmer must explicitly represent and manage the interaction between multiple processes and choreograph their data exchange; large data-structures (such as distributed arrays, graphs, hash-tables) that are conceptually unitary must be thought of as fragmented across different nodes; all processors must generally execute the same code (in an SPMD fashion) etc.

One response to this problem has been the advent of the *partitioned global address space* (PGAS) model underlying languages such as UPC, Titanium and Co-Array Fortran [3, 11]. These languages permit the programmer to think of a single computation running across multiple processors, sharing a common address space. All data resides at some processors, which is said to have *affinity* to the data. Each processor may operate directly on the data it contains but must use some indirect mechanism to access or update data at other processors. Some kind of global *barriers* are used to ensure that processors remain roughly in lock-step.

X10 is a modern object-oriented programming language in the PGAS family. The fundamental goal of X10 is to enable scalable, high-performance, high-productivity transformational programming for high-end computers—for traditional numerical computation workloads (such as weather simulation, molecular dynamics, particle transport problems etc) as well as commercial server workloads.

X10 is based on state-of-the-art object-oriented programming ideas primarily to take advantage of their proven flexibility and ease-of-use for a wide spectrum of programming problems. X10 takes advantage of several years of research (e.g., in the context of the Java Grande forum, [7, 1]) on how to adapt such languages to the context of high-performance numerical computing. Thus X10 provides support for user-defined *struct types* (such as `Int`, `Float`, `Complex` etc), supports a very flexible form of multi-dimensional arrays (based on ideas in ZPL [4]) and supports IEEE-standard floating point arithmetic. Some capabilities for supporting operator overloading are also provided.

X10 introduces a flexible treatment of concurrency, distribution and locality, within an integrated type system. X10 extends the PGAS model with *asynchrony* (yielding the *APGAS* programming model). X10 introduces *places* as an abstraction for a computational context with a locally synchronous view of shared memory. An X10 computation runs over a large collection of places. Each place hosts some data and runs one or more *activities*. Activities are extremely lightweight threads of execution. An activity may synchronously (and *atomically*) use one or more memory locations in the place in which it resides, leveraging current symmetric multiprocessor (SMP) technology. To access or update memory at other places, it must spawn activities asynchronously (either explicitly or implicitly). X10 provides weaker ordering guarantees for inter-place data access, enabling applications to scale. *Immutable* data needs no consistency management and may be freely copied by the implementation between places. One or more *clocks* may be used to order activities running in multiple places. Arrays may be distributed across multiple places. Arrays support parallel collective operations. A novel

exception flow model ensures that exceptions thrown by asynchronous activities can be caught at a suitable parent activity. The type system tracks which memory accesses are local. The programmer may introduce place casts which verify the access is local at run time. Linking with native code is supported.

X10 v2.0 builds on v1.7 to support the following features: *structs* (i.e., “headerless”, inlinable objects), type rules for preventing escape of `this` from a constructor, the introduction of a global object model, permitting user-specified (immutable) fields to be replicated with the object reference. `value` classes are no longer supported; their functionality is accomplished by using structs or global fields and methods.

Several representative idioms for concurrency and communication have already found pleasant expression in X10. We intend to develop several full-scale applications to get better experience with the language, and revisit the design in the light of this experience.

2 Overview of X10

X10 is a statically typed object-oriented language, extending a sequential core language with *places*, *activities*, *clocks*, (distributed, multi-dimensional) *arrays* and *struct* types. All these changes are motivated by the desire to use the new language for high-end, high-performance, high-productivity computing.

2.1 Object-oriented features

The sequential core of X10 is a *container-based* object-oriented language similar to Java or Scala. Programmers write X10 code by defining containers for behavior called *interfaces* (§8), *classes* (§9) and *structs* (§10).

Interfaces An X10 interface specifies a collection of abstract methods (a *signature*). Interfaces may extend one or more interfaces, can be extended by other interfaces and implemented by classes and structs. To implement an interface a class or struct must declare that it implements that interface, and provide definitions for all the methods defined in that interface.

Classes and Structs There are two kinds of concrete containers: *classes* (§9) and *structs* (§10). Containers may have type parameters, have fields, methods and inner containing types and can implement one or more interfaces. A method or an immutable field of a container may be marked `global`. For a given entity (instance of a container), global fields and methods may be accessed from any place. Otherwise they may be accessed only in the place in which they were created.

Classes are organized in a single-inheritance tree. Classes may have mutable fields. Instances of classes with mutable fields may not be freely copied from

place to place. Non-global methods may be invoked on such an object only by an activity in the same place as the object. The `null` reference is a value of any reference type.

In contrast, structs are “headerless” objects that can be onlined in the containing object. Structs cannot inherit code – hence a variable of a struct type contains an instance of precisely that struct. Structs may not be recursive, i.e. there can be no cycles in the graph whose nodes are structs and edges $s \rightarrow t$ reflect that struct s has a field of (struct) type t . Thus a struct instance may be represented with exactly as much space as necessary to store the fields of that struct (modulo alignment considerations). Structs can be copied freely.

X10 has no primitive classes. However, the standard library `x10.lang` supplies structs `Boolean`, `Byte`, `Short`, `Char`, `Int`, `Long`, `Float`, `Double`, `Complex` and `String`. The user may defined additional arithmetic structs using the facilities of the language.

Functions. X10 provides functions (§11) to allow code to be used as values. The body of a function may capture variables in the function’s environment (technically, then, functions are *closures*). Functions are used implicitly in *asyns*, *futures*, and *array initializers*. For example, the following method uses a function to increment elements of an array.

```
def incr(A: Array[Int]): Array[Int] {  
    val f = (x: Int) => x+1; // e.g., f(1) == 2  
    return A.lift(f);  
}
```

Dependent types Classes and interfaces may declare *properties*: immutable object members bound at object construction. Types may be defined by constraining a class or interface’s properties. Properties enable the definition of *dependent types*.

For example, the following code declares a class for a two-dimensional `Point` class with an `add` method.

```
struct Point(x: Int, y: Int) {  
    def add(p: Point) = Point(x+p.x,y+p.y);  
}
```

The class has integer value properties `x` and `y`. The `add` method creates and returns a new point by element-wise addition. The dependent type `Point{x==0}` is the type of all points with `x` set to 0; that is, all points along the *y*-axis.

Generic types Classes and interfaces may have type parameters, permitting the definition of *generic types*. Type parameters may be instantiated by any type – a container type or a function type.

For example, the following code declares a simple `List` class with a type parameter `T`.

```
class List[T] {
  var head: T;
  var tail: List[T];
  def this(h: T, t: List[T]) { head = h; tail = t; }
  def append(x: T) {
    if (this.tail == null)
      this.tail = new List(x, null);
    else
      this.tail.add(x);
  }
}
```

The constructor (`def this`) initializes the fields of the new object. The `add` method appends an element to the list. `List` is a generic type. When instances of `List` are allocated, the type parameter `T` must be bound to a concrete type. `List[Int]` is the type of lists of element type `Int`, `List[String]` is the type of lists of element type `String`.

2.2 The sequential core

Control flow. X10 supports standard sequential control flow constructs: `if` statements, `while` loops, `for` loops, `switch` statements, etc. X10 also supports exceptions: exceptions are raised by `throw` statements and are handled by `try-catch` statements.

Primitive operations. The language provides syntax for performing binary and unary operations on values. The programmer may specify code by using operator definitions.

Allocation. Objects are allocated with the `new` operator (§12.22), which takes a class name and type and value arguments to pass to the class’s constructor. The constructor must ensure that all properties of the class and its superclasses are bound. Structs are created through the invocation of a constructor without using the `new` operator and live on the heap only in the fields of objects.

Coercions and conversions X10 supports implicit and explicit coercions and conversions (§4.9).

Values of one type can be converted to another type using the `as` operation:

```
val x: Int = 65535;
val y: Byte = x as Byte; // convert to Byte,
                        // retaining the lower 8 bits
```

The `as` operation does not necessarily preserve equality and for numeric values may result in a loss of precision.

References may be coerced to another type, preserving object identity. A run-time check is performed to ensure the reference is to an object of the target type. If not, a `ClassCastException` is thrown.

```
// C and D are immediate subclasses of B.
val x: B = new C();
val y: C = x as C; // run-time check succeeds
val z: D = x as D; // run-time check fails
```

2.3 Places and activities

An X10 program is intended to run on a wide range of computers, from uniprocessors to large clusters of parallel processors supporting millions of concurrent operations. To support this scale, X10 introduces the central concept of *place* (§14). Conceptually, a place is a “virtual shared-memory multi-processor”: a computational unit with a finite (though perhaps changing) number of hardware threads and a bounded amount of shared memory, uniformly accessible by all threads.

An X10 computation acts on *data objects*(§6) through the execution of lightweight threads called *activities*(§15). Objects are of two kinds. A *scalar* object has a small, statically fixed set of fields, each of which has a distinct name. A scalar object is located at a single place and stays at that place throughout its lifetime. An *aggregate* object has many fields (the number may be known only when the object is created), uniformly accessed through an index (e.g., an integer) and may be distributed across many places. The distribution of an aggregate object remains unchanged throughout the computation. X10 assumes an underlying garbage collector will dispose of (scalar and aggregate) objects and reclaim the memory associated with them once it can be determined that these objects are no longer accessible from the current state of the computation. (There are no operations in the language to allow a programmer to explicitly release memory.)

X10 has a *unified* or *global address space*. This means that an activity can reference objects at other places. However, an activity may synchronously access data items only in the current place (the place in which the activity is running). It may atomically update one or more data items, but only in the current place. To read a remote location, an activity must spawn another activity *asynchronously* (§15.2). This operation returns immediately, leaving the spawning activity with a *future* (§15.8) for the result. Similarly, remote location can be written into only by asynchronously spawning an activity to run at that location.

Throughout its lifetime an activity executes at the same place. An activity may dynamically spawn activities in the current or remote places.

Place casts. The programmer may use the standard type cast mechanism (§12.23) to cast a value to a located type. A `BadPlaceException` is thrown if the value is not of the given type. This is the only language construct that throws a `BadPlaceException`.

Atomic blocks. X10 introduces statements of the form `atomic S` where `S` is a statement. The type system ensures that such a statement will dynamically access only local data. (The statement may throw a `BadPlaceException`—but only because of a failed place cast.) Such a statement is executed by the activity as if in a single step during which all other activities are frozen.

Asynchronous activities. An asynchronous activity is created by a statement `async (P) S` where `P` is a place expression and `S` is a statement. Such a state-

ment is executed by spawning an activity at the place designated by *P* to execute statement *S*.

An asynchronous expression of type `Future[T]` has the form `future (p) e` where *e* is an expression of type *T*. The expression *e* may reference final and shared variables declared in the lexically enclosing environment. It executes the expression *e* at the place *p* as an asynchronous activity, immediately returning with a future. The future may later be forced causing the activity to be blocked until the return value has been computed by the asynchronous activity.

2.4 Clocks

The MPI style of coordinating the activity of multiple processes with a single barrier is not suitable for the dynamic network of (possibly diverse) activities in an X10 computation. Instead, it becomes necessary to allow a computation to use multiple barriers. X10 *clocks* (§16) are designed to offer the functionality of multiple barriers in a dynamic context while still supporting determinate, deadlock-free parallel computation.

Activities may use clocks to repeatedly detect quiescence of arbitrary programmer-specified, data-dependent set of activities. Each activity is spawned with a known set of clocks and may dynamically create new clocks. At any given time an activity is *registered* with zero or more clocks. It may register newly created activities with a clock, un-register itself with a clock, suspend on a clock or require that a statement (possibly involving execution of new async activities) be executed to completion before the clock can advance. At any given step of the execution a clock is in a given phase. It advances to the next phase only when all its registered activities have *quiesced* (by executing a `next` operation on the clock). When a clock advances, all its activities may now resume execution.

Thus clocks act as *barriers* for a dynamically varying collection of activities. They generalize the barriers found in MPI style program in that an activity may use multiple clocks simultaneously. Yet programs using clocks are guaranteed not to suffer from deadlock.

Future Extensions. *In future versions of the language, clocks will be integrated into the X10 type system, permitting variables to be declared so that they are final in each phase of a clock.*

2.5 Arrays, regions and distributions

An X10 array type is a map from a *distribution* (§17.3) to a type, which may itself be an array type.

A distribution is a map from a *region* (§17.2) to places. A region is a collection of *points* or *indices*. For instance, the region `[0..200, 1..100]` specifies a collection of two-dimensional points (i, j) with i ranging from 0 to 200 and j ranging from 1 to 100. Points are used in array index expressions to pick out a particular array element.

Operations are provided to construct regions from other regions, and to iterate over regions. Standard set operations, such as union, disjunction and set difference are available for regions.

A primitive set of distributions is provided, together with operations on distributions. A *sub-distribution* of a distribution is one defined on a smaller region and agrees with the distribution at all points. The standard operations on regions are extended to distributions.

A new array can be created by restricting an existing array to a sub-distribution, by combining multiple arrays, and by performing pointwise operations on arrays with the same distribution.

X10 allows array constructors to iterate over the underlying distribution and specify a value at each item in the underlying region. Such a constructor may spawn activities at multiple places.

In future versions of the language, a programmer may specify new distributions, and new operations on distributions.

2.6 Annotations

X10 supports annotations on classes and interfaces, methods and constructors, variables, types, expressions and statements. These annotations may be processed by compiler plugins.

2.7 Translating MPI programs to X10

While X10 permits considerably greater flexibility in writing distributed programs and data structures than MPI, it is instructive to examine how to translate MPI programs to X10.

Each separate MPI process can be translated into an X10 place. Async activities may be used to read and write variables located at different processes. A single clock may be used for barrier synchronization between multiple MPI processes. X10 collective operations may be used to implement MPI collective operations. X10 is more general than MPI in (a) not requiring synchronization between two processes in order to enable one to read and write the other's values, (b) permitting the use of high-level atomic blocks within a process to obtain mutual exclusion between multiple activities running in the same node (c) permitting the use of multiple clocks to combine the expression of different physics (e.g., computations modeling blood coagulation together with computations involving the flow of blood), (d) not requiring an SPMD style of computation.

2.8 Summary and future work

2.8.1 Design for scalability

X10 is designed for scalability. An activity may atomically access only multiple locations in the current place. Unconditional atomic blocks are statically guaranteed to be non-blocking, and may be implemented using non-blocking techniques that avoid mutual exclusion bottlenecks. Data-flow synchronization permits point-to-point coordination between reader/writer activities, obviating the need for barrier-based or lock-based synchronization in many cases.

2.8.2 Design for productivity

X10 is designed for productivity.

Safety and correctness. Programs written in X10 are guaranteed to be statically *type safe*, *memory safe* and *pointer safe*. Static type safety guarantees that at run time a location contains only those values whose

dynamic type satisfies the constraints imposed by the location's static type and every run-time operation performed on the value in a location is permitted by the static type of the location.

Memory safety guarantees that an object may only access memory within its representation, and other objects it has a reference to. X10 supports no pointer arithmetic, and bound-checks array accesses dynamically if necessary. X10 uses dynamic garbage collection to collect objects no longer referenced by the computation. X10 guarantees that no object can retain a reference to an object whose memory has been reclaimed. Further, X10 guarantees that every location is initialized at run time before it is read, and every value read from a location has previously been written into that location.

Because places are reflected in the type system, static type safety also implies *place safety*: a location may contain references to only those objects whose location satisfies the restrictions of the static place type of the location.

X10 programs that use only clocks and unconditional atomic blocks are guaranteed not to deadlock. Unconditional atomic blocks are non-blocking, hence cannot introduce deadlocks (assuming the implementation is correct).

Many concurrent programs can be shown to be determinate (hence race-free) statically.

Integration. A key issue for any new programming language is how well it can be integrated with existing (external) languages, system environments, libraries and tools.

We believe that X10, like Java, will be able to support a large number of libraries and tools. An area where we expect future versions of X10 to improve on Java like languages is *native integration* (§19). Specifically, X10 will permit multi-dimensional local arrays to be operated on natively by native code.

2.8.3 Conclusion

X10 is considerably higher-level than thread-based languages in that it supports dynamically spawning very lightweight activities, the use of atomic operations for mutual exclusion, and the use of clocks for repeated quiescence detection.

Yet it is much more concrete than languages like HPF in that it forces the programmer to explicitly deal with distribution of data objects. In this the language

reflects the designers' belief that issues of locality and distribution cannot be hidden from the programmer of high-performance code in high-end computing. A performance model that distinguishes between computation and communication must be made explicit and transparent.¹ At the same time we believe that the place-based type system and support for generic programming will allow the X10 programmer to be highly productive; many of the tedious details of distribution-specific code can be handled in a generic fashion.

We expect the next version of the language to be significantly informed by experience in implementing and using the language. We expect it to have constructs to support continuous program optimization, and allow the programmer to provide guidance on clustering places to (hardware) nodes. For instance, we may introduce a notion of hierarchical clustering of places.

¹In this X10 is similar to more modern languages such as ZPL [4].

3 Lexical structure

In general, X10 follows Java rules [5, Chapter 3] for lexical structure.

Lexically a program consists of a stream of white space, comments, identifiers, keywords, literals, separators and operators.

Whitespace ASCII space, horizontal tab (HT), form feed (FF) and line terminators constitute white space.

Comments All text included within the ASCII characters “/*” and “*/” is considered a comment and ignored; nested comments are not allowed. All text from the ASCII characters “//” to the end of line is considered a comment and is ignored.

Identifiers Identifiers are defined as in Java. Identifiers consist of a single letter followed by zero or more letters or digits. Letters are defined as the characters for which the Java method `Character.isJavaIdentifierStart` returns true. Digits are defined as the ASCII characters 0 through 9.

Keywords X10 reserves the following keywords:

abstract	any	as	async
at	ateach	atomic	await
break	case	catch	class
clocked	const	continue	current
def	default	do	else
extends	extern	final	finally
finish	for	foreach	future

global			
goto	has	here	if
implements	import	instanceof	interface
native	new	next	nonblocking
or	package	pinned	private
protected	property	public	return
safe	self	sequential	shared
static			
super	switch	this	throw
throws	to	try	type
val	value	var	when
while			

Note that the primitive types are not considered keywords. The keyword `goto` is reserved, but not used.

Literals Literals are either integers, unsigned integers, floating point numbers, booleans, characters, strings, and `null`. X10 v2.0 defines literal syntax in the same way as Java does except that the suffix `'u'` or `'U'` is used for unsigned integers.

Separators X10 has the following separators and delimiters:

() { } [] ; , .

Operators X10 has the following operators:

```

==  !=  <  >  <=  >=
&&  ||  &  |  ^
<<  >>  >>>
+   -   *   /   %
++  --  !   ~
&=  |=  ^=
<<= >>= >>>=
+=  -=  *=  /=  %=
=   ?   :   =>  ->
<:  >:  @   ..

```

4 Types

X10 is a *strongly typed* object-oriented language: every variable and expression has a type that is known at compile-time. Types limit the values that variables can hold and specify the places at which these values lie.

X10 supports three kinds of runtime entities, *objects*, *structs*, and *functions*. Objects are instances of *classes* (§9). They may contain mutable fields and must stay resident in the place in which they were created. Objects are said to be *boxed* in that variables of a class type are implemented through a single memory location that contains a reference to the memory containing the declared state of the object (and other meta-information such as the list of methods of the object). Thus objects are represented through an extra level of indirection. A consequence of this flexibility is that every class type contains the value `null` corresponding to the invalid reference. `null` is often useful as a default value. Further, two objects may be compared for equality (`==`) in constant time by simply containing references to the memory used to represent the objects.

Structs are instances of *struct classes* (§10). They are immutable and may be freely copied from place to place. Further, they may be allocated inline, i.e. using only as much memory as necessary to hold the fields of the struct (and any additional memory necessary to satisfy alignment constraints for data layout).

Functions are instances of *function types*— §11) and are created using function literals $(x_1:T_1, \dots, x_n:T_n)\{c\}:T \Rightarrow e$. Functions contain no user-visible mutable or immutable state; their representation, however, contains enough memory to hold the values of the variables in the environment that are referenced in the body of the function (e). Functions may be freely copied from place to place and may be repeatedly applied to a set of arguments, provided that the precondition c is satisfied.

These runtime entities are classified into different groups using *types*. Types are used in variable declarations, explicit coercions and conversions, object creation,

array creation, class literals, static state and method accessors, and `instanceof` expressions.¹

X10 has a unified type system. The top of the type hierarchy is the interface `x10.lang.Any`. This specifies the following signature:

```
package x10.lang;
public interface Any {
    property def home():Place;
    property def at(p:Object):Boolean;
    property def at(p:Place):Boolean;
    global safe def toString():String;
    global safe def typeName():String;
    global safe def equals(Any):Boolean;
    global safe def hashCode():Int;
}
```

Properties are described in (§4.1.3): in brief they are immutable instance fields of objects and structs that can be used to construct types through constraints. Property methods (§4.1.3) permit more complex expressions to be defined using properties and other property methods. A method is `safe` if it has certain behavioral characteristics §9.6.3. It is `global` if it can be invoked at any place.

Types in X10 are specified through declarations and through type constructors, described in the remainder of the chapter:

- A class declaration defines a *class type* (§4.1).
- An interface declaration defines an *interface type* (§4.1.2).
- Classes and interface have *type parameters*. A class or interface with one or more type parameters is a *generic class* or *generic interface* (§4.2.1).
- New type constructors may be defined with *type definitions* (§4.3).
- Methods, constructors, closures, and type definitions may have *type parameters*, which are instantiated with concrete types at invocation (§4.2).

¹In order to allow this version of the language to focus on the core new ideas, X10 v2.0 does not have user-definable class loaders, though there is no technical reason why they could not have been added.

- *Function type* constructors are used to define function types; functions and method selectors have function type (§4.5).
- A *constrained type* constrains the properties of a base type (§4.4).
- Types may be marked with user-defined annotations. *Annotated types* (§4.6) may be processed by compiler plugins.

```

      Type ::= FunctionType
            | ConstrainedType
      FunctionType ::= TypeParameters? ( Formals? ) Constraint? Throws? => Type
      TypeParameters ::= [ TypeParameter ( , TypeParameter ) * ]
      TypeParameter ::= Identifier
      Throws ::= throws TypeName ( , TypeName ) *
      ConstrainedType ::= Annotation* BaseType Constraint? PlaceConstraint?
      BaseType ::= ClassBaseType
                | InterfaceBaseType
                | PathType
                | ( Type )
      ClassType ::= Annotation* ClassBaseType Constraint? PlaceConstraint?
      InterfaceType ::= Annotation* InterfaceBaseType Constraint? PlaceConstraint?
      PathType ::= Expression . Identifier
      Annotation ::= @ InterfaceBaseType Constraint?
      ClassOrInterfaceType ::= ClassType
                          | InterfaceType
      ClassBaseType ::= TypeName
      InterfaceBaseType ::= TypeName

```

4.1 Classes and interfaces

4.1.1 Class types

A *class declaration* (§9) introduces a *class type* containing all instances of the class.

Class instances are created via constructor calls. Class instances have fields and methods, type members, and value properties bound at construction time. In ad-

dition, classes have static members: constant fields, type definitions, and member classes and member interfaces.

A class with type parameters is *generic*. A class type is legal only if all of its parameters are instantiated on concrete types.

X10 does not permit mutable static state, so the role of static methods and initializers is quite limited. Instead programmers should use singleton classes to carry mutable static state.

Classes are structured in a single-inheritance hierarchy. All classes extend the class `x10.lang.Object`. Classes are declared to extend a single superclass (except for `Object`, which extends no other class).

Variables of class type may contain the value `null`.

4.1.2 Interface types

An *interface declaration* (§8) defines an *interface type*, which specifies a set of methods, type members, and properties to be implemented by any class declared to implement the interface. Interfaces can also have static members: constant fields, type definitions, and member classes and interfaces.

An interface may extend multiple interfaces.

Classes may be declared to implement multiple interfaces. Semantically, the interface type is the set of all objects that are instances of classes that implement the interface. A class implements an interface if it is declared to and if it implements all the methods and properties defined in the interface.

4.1.3 Properties

Classes and interfaces may have *properties*, public final instance fields bound on object creation. For example, the following code declares a class named `Point` with properties `x` and `y` and a `move` method. The properties are bound using the `property` statement in the constructor.

```
class Point(x: Int, y: Int) {  
  def this(x: Int, y: Int) { property(x, y); }  
  def move(dx: Int, dy: Int) = new Point(x+dx, y+dy);  
}
```

The properties of a class or interface may be constrained with a boolean expression. The type `Point{x==0}` is the set of all points whose `x` property is `0`.

4.2 Type parameters

A class, interface, method, or closure may have type parameters whose scope is the signature and body of the declaring class, interface, method, or closure.

Similarly, a type definition may have type parameters that scope over the body of the type definition.

Type parameters may be constrained by a *guard* on the declaration (§9, §4.3, §9.6, §11.2). The type parameters of classes and interfaces must be bound to concrete types (possibly to a type parameter) for the type to be legal; thus `List[int]` and `List[C]` are legal types, but `List` alone is not. The type parameters of methods and closures must be bound to concrete types at invocation. Parametrized type definitions specify new type constructors; the type parameters of a type definition must be bound to yield a type.

4.2.1 Generic types

A *generic class* is a class declared with one or more type parameters. Generic classes can be instantiated by instantiating the type parameters of the base type.

Consider the following declaration of a `Cell` class.

```
class Cell[X] {  
  var x: X;  
  def this(x: X) { this.x = x; }  
  def get(): X = x;  
  def set(x: X) = { this.x = x; }  
}
```

This declares a class `Cell` with a type parameter `X`. `Cell` may be used as a type by instantiating `X`.

`Cell[Int]` is the type of all `Cell` containing an `Int`. The `get` method returns an `Int`; the `set` method takes an `Int` as argument. Note that `Cell` alone is not a legal type because the parameter is not bound.

Parameters may be declared as *invariant*, *covariant*, or *contravariant*. The *X* parameter of *Cell* above is invariant. Consider the following classes:

```
class Get[+X] {
  var x: X;
  def this(x: X) { this.x = x; }
  def get(): X = x;
}

class Set[-X] {
  var x: X;
  def this(x: X) { this.x = x; }
  def set(x: X) = { this.x = x; }
}
```

The *X* parameter of the *Get* class is covariant; the *X* parameter of the *Set* class is contravariant.

A covariant type parameter is permitted to appear only in covariant type positions, and a contravariant type parameter in contravariant positions.

- The return type of a method is a covariant position.
- The argument types of a method are contravariant positions.
- Whether a type argument position of a generic class, interface or struct type *C* is covariant or contravariant is determined by the + or - annotation at that position in the declaration of *C*.

Given types *S* and *T*.

- If the parameter of *Get* is covariant, then *Get*[*S*] is a subtype of *Get*[*T*] if *S* is a *subtype* of *T*.
- If the parameter of *Set* is contravariant, then *Set*[*S*] is a subtype of *Set*[*T*] if *S* is a *supertype* of *T*.
- If the parameter of *Cell* is invariant, then *Cell*[*S*] is a subtype of *Cell*[*T*] if *S* is *equal* to *T*.

4.3 Type definitions

With value arguments, type arguments, and constraints, the syntax for X10 types can often be verbose; X10 therefore provides *type definitions* to allow users to define new type constructors.

Type definitions have the following syntax:

$$\text{TypeDefinition} ::= \text{type Identifier} ([\text{TypeParameters}])^? \\ ((\text{Formals}))^? \text{Constraint}^? = \text{Type}$$

A type definition can be thought of as a type-valued function, mapping type parameters and value parameters to a concrete type. The following examples are legal type definitions:

```
type StringSet = Set[String];
type MapToList[K,V] = Map[K,List[V]];
type Nat = UInt{self!=0};
type Int(x: Int) = Int{self==x};
type Dist(r: Int) = Dist{self.rank==r};
type Dist(r: Region) = Dist{self.region==r};
```

As the two definitions of `Dist` demonstrate, type definitions may be overloaded: two type definitions with different numbers of type parameters or with different types of value parameters, according to the method overloading rules (§9.6.2), define distinct type constructors.

Type definitions may appear as (static) class or interface member or in a block statement.

Type definitions are applicative, not generative; that is, they define aliases for types but do not introduce new types. Thus, the following code is legal:

```
type A = Int;
type B = String;
type C = String;
a: A = 3;
b: B = new C("Hi");
c: C = b + ", Mom!";
```

If a type definition has no type parameters and no value parameters and is an alias for a class type, a `new` expression may be used to create an instance of the class using the type definition's name. Given the following type definition:

```
type A = C[T1, ..., Tk]{c};
```

where $C[T_1, \dots, T_k]$ is a class type, a constructor of C may be invoked with $\text{new } A(e_1, \dots, e_n)$, if the invocation $\text{new } C[T_1, \dots, T_k](e_1, \dots, e_n)$ is legal and if the constructor return type is a subtype of A .

The collection of type definitions in `x10.lang._` is automatically imported in every compilation unit.

4.4 Constrained types

Given a type T , a *constrained type* $T\{e\}$ may be constructed by constraining its properties with a boolean expression e .

$T\{e\}$ is a *dependent type*, that is, a type dependent on values. The type T is called the *base type* and e is called the *constraint*. For reference types, the constraint may specify the places at which the object resides.

For brevity, the constraint may be omitted and interpreted as `true`.

Constraints on properties induce a natural subtyping relationship: $C\{c\}$ is a subtype of $D\{d\}$ if C is a subclass of D and c entails d .

Type parameters cannot be constrained.

4.4.1 Constraints

Expressions used as constraints are restricted by the constraint system in use to ensure that the constraints can be solved at compile time. The X10 compiler allows compiler plugins to be installed to extend the constraint language and the constraint system. Constraints must be of type `Boolean`. The compiler supports the following constraint syntax.

$$\begin{aligned} \text{Constraint} & ::= \text{ValueArguments } \text{Guard}^? \\ & \quad | \quad \text{ValueArguments}^? \text{Guard} \end{aligned}$$

$$\begin{aligned} \text{ValueArguments} & ::= (\text{ArgumentList}^?) \\ \text{ArgumentList} & ::= \text{Expression } (, \text{Expression })^* \\ \text{Guard} & ::= \{ \text{DepExpression } \} \\ \text{DepExpression} & ::= (\text{Formal } ;)^* \text{ArgumentList} \end{aligned}$$

In X10 v2.0 value constraints may be equalities (`==`), disequalities (`!=`) and conjunctions thereof. The terms over which these constraints are specified include literals and (accessible, immutable) variables and fields, property methods, and the special constants `here`, `self`, and `this`. Additionally, place types are permitted (§4.4.2).

Type constraints may be subtyping and supertyping (`<:` and `:>`) expressions over types.

Subsequent implementations are intended to support boolean algebra, arithmetic, relational algebra, etc., to permit types over regions and distributions. We envision this as a major step towards removing most, if not all, dynamic array bounds and place checks from X10.

Acyclicity restriction

To ensure that type-checking is decidable, we require that property graphs be acyclic. That is, it should not be the case at runtime that a set of objects can be created such that the graph formed by taking objects as nodes and adding an edge from m to n if m has a property whose value is n has a cycle in it.

Currently this restriction is not checked by the compiler. Future versions of the compiler will check this restriction by introducing rules on escaping of `this` (§9.4) before the invocation of property calls.

4.4.2 Place constraints

An X10 computation spans multiple places (§14). Each place contains data and activities that operate on that data. X10 v2.0 does not permit the dynamic creation of a place. Each X10 computation is initiated with a fixed number of places, as determined by a configuration parameter. In this section we discuss how the programmer may supply place type information, thereby allowing the compiler to check data locality, i.e., that data items being accessed in an atomic section are local.

$$\begin{aligned} \textit{PlaceConstraint} & ::= \textit{! Place}^? \\ \textit{Place} & ::= \textit{Expression} \end{aligned}$$

Because of the importance of places in the X10 design, special syntactic support is provided for constrained types involving places.

All X10 classes extend the class `x10.lang.Object`, which defines a property `home` of type `Place`.

If a constrained reference type `T` has an `!p` suffix, the constraint for `T` is implicitly assumed to contain the clause `self.home==p`; that is, `C{c}!p` is equivalent to `C{self.home==p && c}`.

The place `p` may be omitted. It defaults to `this` for types in field declarations, and to `here` elsewhere.

4.4.3 Constraint semantics

STATIC SEMANTICS RULE (Variable occurrence): In a dependent type `T = C{c}`, the only variables that may occur in `c` are (a) `self`, (b) properties visible at `T`, (c) final local variables, final method parameters or final constructor parameters visible at `T`, (d) final fields visible at `T`'s lexical place in the source program.

STATIC SEMANTICS RULE (Restrictions on `this`): The special variable `this` may be used in a dependent clause for a type `T` only if (a) `T` occurs in a property declaration for a class, (b) `T` occurs in an instance method, (c) `T` occurs in an instance field, (d) `T` occurs in an instance initializer.

In particular, `this` may not be used in types that occur in a static context, or in the arguments, body or return type of a constructor or in the extends or implements clauses of class and interface definitions. In these contexts, the object that `this` would correspond to is not defined.

STATIC SEMANTICS RULE (Variable visibility): If a type `T` occurs in a field, method or constructor declaration, then all variables used in `T` must have at least the same visibility as the declaration. The relation “at least the same visibility as” is given by the transitive closure of:

`public > protected > package > private`

All inherited properties of a type `T` are visible in the property list of `T`, and the body of `T`.

In general, variables (i.e., local variables, parameters, properties, fields) are visible at `T` if they are defined before `T` in the program. This rule applies to types in property lists as well as parameter lists (for methods and constructors). A formal parameter is visible in the types of all other formal parameters of the same method, constructor, or type definition, as well as in the method or constructor body itself.

Properties are accessible via their containing object—`this` within the body of their class declaration. The special variable `this` is in scope at each property declaration, constructor signatures and bodies, instance method signatures and bodies, and instance field signatures and initializers, but not in scope at `static` method or field declarations or `static` initializers.

We permit variable declarations $v: T$ where T is obtained from a dependent type $C\{c\}$ by replacing one or more occurrences of `self` in c by v . (If such a declaration $v: T$ is type-correct, it must be the case that the variable v is not visible at the type T . Hence we can always recover the underlying dependent type $C\{c\}$ by replacing all occurrences of v in the constraint of T by `self`.)

For instance, $v: \text{Int}\{v == 0\}$ is shorthand for $v: \text{Int}\{\text{self} == 0\}$.

STATIC SEMANTICS RULE (Constraint type): The type of a constraint c must be `Boolean`.

A variable occurring in the constraint c of a dependent type, other than `self` or a property of `self`, is said to be a *parameter* of c .

An instance o of C is said to be of type $C\{c\}$ (or: *belong to* $C\{c\}$) if the predicate c evaluates to `true` in the current lexical environment, augmented with the binding $\text{self} \mapsto o$. We shall use the function $\llbracket C\{c\} \rrbracket$ to denote the set of objects that belong to $C\{c\}$.

4.4.4 Consistency of dependent types

A dependent type $C\{c\}$ may contain zero or more parameters. We require that a type never be empty—so that it is possible for a variable of the type to contain a value. This is accomplished by requiring that the constraint c must be satisfiable *regardless* of the value assumed by parameters to the constraint (if any). Formally, consider a type $T = C\{c\}$, with the variables $f_1: F_1, \dots, f_k: F_k$ free in c . Let $S = \{f_1: F_1, \dots, f_k: F_k, f_{k+1}: F_{k+1}, \dots, f_n: F_n\}$ be the smallest set of declarations containing $f_1: F_1, \dots, f_k: F_k$ and closed under the rule: $f: F$ in S if a reference to variable f (which is declared as $f: F$) occurs in a type in S .

(NOTE: The syntax rules for the language ensure that S is always finite. The type for a variable v cannot reference a variable whose type depends on v .)

We say that $T = C\{c\}$ is *parametrically consistent* (in brief: *consistent*) if:

- Each type F_1, \dots, F_n is (recursively) parametrically consistent, and

- It can be established that $\forall f_1: F_1, \dots, f_n: F_n. \exists \text{self}: C. c \ \&\& \ \text{inv}(C)$.

where $\text{inv}(C)$ is the invariant associated with the type C (§9.1). Note by definition of S the formula above has no free variables.

STATIC SEMANTICS RULE: For a declaration $v: T$ to be type-correct, T must be parametrically consistent. The compiler issues an error if it cannot determine the type is parametrically consistent.

Example 4.4.1 A class that represents a line has two distinct points:

```
class Line(start: Point,
           end: Point{self != this.start}) {...}
```

□

One can use dependent type to define other closed geometric figures as well.

Example 4.4.2 Here is an example:

```
class Point(x: Int, y: Int) {...}
```

To see that the declaration `end: Point{self != start}` is parametrically consistent, note that the following formula is valid:

$$\forall \text{this}: \text{Line}. \exists \text{self}: \text{Point}. \text{self} \neq \text{this.start}$$

since the set of all Points has more than one element.

□

Example 4.4.3 A triangle has three lines sharing three vertices.

```
class Triangle
(a: Line,
 b: Line{a.end == b.start},
 c: Line{b.end == c.start && c.end == a.start})
{ ... }
```

Given $a: \text{Line}$, the type $b: \text{Line}\{a.\text{end} == b.\text{start}\}$ is consistent, and given the two, the type $c: \text{Line}\{b.\text{end} == c.\text{start}, c.\text{end} == a.\text{start}\}$ is consistent.

□

4.5 Function types

Function types are defined via the \Rightarrow type constructor. Closures (§11) and method selectors (§11.3) are of function type. The general form of a function type is:

$$(\mathbf{x}_1: T_1, \dots, \mathbf{x}_n: T_n) \{c\} \Rightarrow T \\ \text{throws } S_1, \dots, S_k$$

This is the type of functions that take value parameters \mathbf{x}_i of types T_i such that the guard c holds and returns a value of type T or throws exceptions of types S_i .

The value parameters are in scope throughout the function signature—they may be used in the types of other formal parameters and in the return type. Value parameters names may be omitted if they are not used. The guard specifies a condition that must hold for an application to be well-typed.

$$\begin{aligned} \text{FunctionType} &::= \text{TypeParameters}^? (\text{Formals}^?) \text{Constraint}^? \Rightarrow \text{Type Throws}^? \\ \text{TypeParameters} &::= [\text{TypeParameter} (, \text{TypeParameter})^*] \\ \text{TypeParameter} &::= \text{Identifier} \\ \text{Formals} &::= \text{Formal} (, \text{Formal})^* \end{aligned}$$

For every sequence of types T_1, \dots, T_n, T , and n distinct variables $\mathbf{x}_1, \dots, \mathbf{x}_n$ and constraint c , the expression $(\mathbf{x}_1:T_1, \dots, \mathbf{x}_n:T_n) \{c\} \Rightarrow T$ is a *function type*. It stands for the set of all functions f which can be applied in a place p to a list of values (v_1, \dots, v_n) provided that the constraint $c[v_1, \dots, v_n, p/\mathbf{x}_1, \dots, \mathbf{x}_n, \text{here}]$ is true, and which returns a value of type $T[v_1, \dots, v_n/\mathbf{x}_1, \dots, \mathbf{x}_n]$. When c is true, the clause $\{c\}$ can be omitted. When $\mathbf{x}_1, \dots, \mathbf{x}_n$ do not occur in c or T , they can be omitted. Thus the type $(T_1, \dots, T_n) \Rightarrow T$ is actually shorthand for $(\mathbf{x}_1:T_1, \dots, \mathbf{x}_n:T_n) \{\text{true}\} \Rightarrow T$, for some variables $\mathbf{x}_1, \dots, \mathbf{x}_n$.

Juxtaposition is used to express function application: the expression $f(a_1, \dots, a_n)$ expresses the application of a function f to the argument list a_1, \dots, a_n .

Note that function invocation may throw unchecked exceptions.

A function type is covariant in its result type and contravariant in each of its argument types. That is, let $S_1, \dots, S_n, S, T_1, \dots, T_n, T$ be any types satisfying $S_i <: T_i$ and $S <: T$. Then $(\mathbf{x}_1:T_1, \dots, \mathbf{x}_n:T_n) \{c\} \Rightarrow S$ is a subtype of $(\mathbf{x}_1:S_1, \dots, \mathbf{x}_n:S_n) \{c\} \Rightarrow T$.

A value f of a function type $(\mathbf{x}_1:T_1, \dots, \mathbf{x}_n:T_n) \{c\} \Rightarrow T$ also has all the methods of `Any` associated with it (see §11.5).

A function type $F = (x_1:T_1, \dots, x_n:T_n) \{c\} \Rightarrow T$ can be used as the declared type of local variables, parameters, loop variables, return types of methods and in `_ instanceof F` and `_ as F` expressions.

A class or struct definition may use a function type F in its `implements` clause; this declares an abstract method `def apply(x1:T1, ..., xn:Tn){c}:T` on that class. Similarly, an interface definition may specify a function type " F " in its "`extends`" clause. A class or struct implementing such an interface implicitly defines an abstract method `def apply(x1:T1, ..., xn:Tn){c}:T`. Expressions of such a struct, class or interface type can be assigned to variables of type F and can be applied via juxtaposition to an argument list of the right type.

Thus, objects and structs in X10 may behave like functions.

A function type F is not a class type in that it does not extend any type or implement any interfaces, or support equality tests. F cannot be extended by any type. It is not an interface type in that it is not a subtype of `x10.lang.Object`. (Values of type F cannot be assigned to variables of type `x10.lang.Object`.) It is not a struct type in that it has no defined fields and hence no notion of structural equality.

`null` is a legal value for a function type.

4.6 Annotated types

Any X10 type may be annotated with zero or more user-defined *type annotations* (§18).

Annotations are defined as (constrained) interface types and are processed by compiler plugins, which may interpret the annotation symbolically.

A type T is annotated by interface types A_1, \dots, A_n using the syntax `@A1 ... @An T`.

4.7 Subtyping and type equivalence

Subtyping is relation between types. It is the reflexive, transitive closure of the *direct subtyping* relation, defined as follows.

Class types. A class type is a direct subtype of any class it is declared to extend. A class type is direct subtype of any interfaces it is declared to implement.

Interface types. An interface type is a direct subtype of any interfaces it is declared to extend.

Function types. Function types are covariant on their return type and contravariant on their argument types. For instance, a function type $(S1) \Rightarrow T1$ is a subtype of another function type $(S2) \Rightarrow T2$ if $S2$ is a subtype of $S1$ and $T1$ is a subtype of $T2$.

Constrained types. Two dependent types $C\{c\}$ and $C\{d\}$ are said to be *equivalent* if c is true whenever d is, and vice versa. Thus, $\llbracket C\{c\} \rrbracket = \llbracket C\{d\} \rrbracket$.

Note that two dependent type that are syntactically different may be equivalent. For instance, $\text{Int}\{\text{self} \geq 0\}$ and $\text{Int}\{\text{self} == 0 \mid \mid \text{self} > 0\}$ are equivalent though they are syntactically distinct. The Java type system is essentially a nominal system—two types are the same if and only if they have the same name. The X10 type system extends the nominal type system of Java to permit constraint-based equivalence.

A dependent type $C\{c\}$ is a subtype of a type $C\{d\}$ if c implies d . When this subtyping relationship holds, $\llbracket C\{c\} \rrbracket$ is a subset of $\llbracket C\{d\} \rrbracket$. All dependent types defined on a class C refine the unconstrained class type C ; C is equivalent to $C\{\text{true}\}$.

Type parameters. A type parameter X of a class or interface C is a subtype of a type T if the class invariant of C implies that X is a subtype of T . Similarly, T is a subtype of parameter X if the class invariant implies the relationship.

A type parameter X of a method m is a subtype of a type T if the guard of m implies that X is a subtype of T . Similarly, T is a subtype of parameter X if the guard implies the relationship.

4.8 Least common ancestor of types

To compute the type of conditional expressions (§12.20), and of rail constructors (§12.27), the least common ancestor of types must be computed.

The least common ancestor of two types T_1 and T_2 is the unique most-specific type that is a supertype of both T_1 and T_2 .

If the most-specific type is not unique (which can happen when T_1 and T_2 both implement two or more incomparable interfaces), then least common ancestor type is `x10.lang.Any`.

4.9 Coercions and conversions

X10 v2.0 supports the following coercions and conversions

4.9.1 Coercions

A *coercion* does not change object identity; a coerced object may be explicitly coerced back to its original type through a cast. A *conversion* may change object identity if the type being converted to is not the same as the type converted from. X10 permits user-defined conversions (§4.9.2).

Subsumption coercion. A subtype may be implicitly coerced to any supertype.

Explicit coercion (casting with `as`) A reference type may be explicitly coerced to any other reference type using the `as` operation. If the value coerced is not an instance of the target type, a `ClassCastException` is thrown. Casting to a constrained type may require a run-time check that the constraint is satisfied.

Effects of explicit numeric coercion Coercing a number of one type to another type gives the best approximation of the number in the result type, or a suitable disaster value if no approximation is good enough.

- Casting a number to a *wider* numeric type is safe and effective, and can be done by an implicit conversion as well as an explicit coercion. For example, `4 as Long` produces the Long value of 4.

- Casting a floating-point value to an integer value truncates the digits after the decimal point, thereby rounding the number towards zero. `54.321 as Int` is `54`, and `-54.321 as Int` is `-54`. If the floating-point value is too large to represent as that kind of integer, the coercion returns the largest or smallest value of that type instead: `1e110 as Int` is `Int.MAX_VALUE`, `2147483647`.
- Casting a `Double` to a `Float` normally truncates digits: `0.12345678901234567890` as `Float` is `0.12345679f`. This can turn a nonzero `Double` into `0.0f`, the zero of type `Float`: `1e-100 as Float` is `0.0f`. Since `Doubles` can be as large as about `1.79E308` and `Floats` can only be as large as about `3.4E38f`, a large `Double` will be converted to the special `Float` value of `Infinity`: `1e100 as Float` is `Infinity`.
- Integers are coerced to smaller integer types by truncating the high-order bits. If the value of the large integer fits into the smaller integer's range, this gives the same number in the smaller type: `12 as Byte` is the `Byte`-sized `12`, `-12 as Byte` is `-12`. However, if the larger integer *doesn't* fit in the smaller type, the numeric value and even the sign can change: `254 as Byte` is `Byte`-sized `-2`.

4.9.2 Conversions

Widening numeric conversion. A numeric type may be implicitly converted to a wider numeric type. In particular, an implicit conversion may be performed between a numeric type and a type to its right, below:

`Byte < Short < Int < Long < Float < Double`

String conversion. Any object that is an operand of the binary `+` operator may be converted to `String` if the other operand is a `String`. A conversion to `String` is performed by invoking the `toString()` method of the object.

User defined conversions. The user may define conversion operators from type `A` to a container type `B` by specifying a method on `B` as follows:

`public static operator (r: A): T = ...`

The return type T should be a subtype of B . The return type need not be specified explicitly; it will be computed in the usual fashion if it is not. However, it is good practice for the programmer to specify the return type for such operators explicitly.

For instance, the code for `x10.lang.Point` contains:

```
public static global safe operator (r: Rail[int])
    : Point(r.length) = make(r);
```

The compiler looks for such operators on the container type B when it encounters an expression of the form `r as B` (where `r` is of type A). If it finds such a method, it sets the type of the expression `r as B` to be the return type of the method. Thus the type of `r as B` is guaranteed to be some subtype of B .

Example 4.9.1 Consider the following code:

```
val p = [2, 2, 2, 2, 2] as Point;
val q = [1, 1, 1, 1, 1] as Point;
val a = p - q;
```

This code fragment compiles successfully, given the above operator definition. The type of `p` is inferred to be `Point(5)` (i.e. the type `Point{self.rank==5}`). Similarly for `q`. Hence the application of the operator “-” is legal (it requires both arguments to have the same rank). The type of `a` is computed as `Point(5)`. \square

4.10 Built-in types

The package `x10.lang` provides a number of built-in class and interface declarations that can be used to construct types.

4.10.1 The class `Object`

The class `x10.lang.Object` is the supertype of all classes. A variable of this type can hold a reference to any object. The code for this class (with annotations removed) is:

```

public class Object (home: Place)
    implements Any
{
    public native def this();
    public property def home() = home;
    public property def at(p:Place) = home==p;
    public property def at(r:Object) = home==r.home;
    public global safe native def toString() : String;
    public global safe native def typeName() : String;
    public global safe def equals(x:Any) = this == x;
    public global safe native def hashCode():Int;
}

```

4.10.2 The class String

All strings in X10 are instances of the class `x10.lang.String`. A string object is immutable, and has a concatenation operator (+) available on it.

4.10.3 Array types

Arrays in X10 are instances of the class `x10.lang.Array`. Because of the importance of arrays in X10, the language supports more concise syntax for accessing array elements and performing operations on arrays.

The array type `Array[T]` is the type of all reference arrays of base type `T`. Such an array can take on any distribution, over any region.

Both array classes implement the function type `(Point) => T`; the element of array `A` at point `p` may be accessed using the syntax `A(p)`. The `Array` class also implements the `Settable[Point, T]` interface permitting assignment to an array element using the syntax `A(p) = v`.

X10 also supports dependent types for arrays, e.g., `Array[Double]{rank==3}` is the type of all arrays of `Double` of rank 3. The `Array` class has `distribution`, `region`, and `rank` properties. X10 v2.0 defines type definitions that allows a `distribution`, `region`, or `rank` to be specified with on the array type.

```

package x10.lang;
type Array[T](n: Int) = Array[T]{rank==n};

```

```

type Array[T](d: Dist) = Array[T]{dist==d};
type Array[T](r: Region) = Array[T]{region==r};

```

4.10.4 Rails

A *rail* is a one-dimensional, zero-based, local array. It is more primitive than the `Array` class. Rails are indexed by integers rather than multi-dimensional points. Rails have a single `length` property of type `Int`. Rails can be mutable or immutable and are defined by the following class definitions:

```

package x10.lang;
public class ValRail[T](length: Int) implements (Int)=>T, Iterable[T] { }
public class Rail[T](length: Int) implements (Int)=>T, Settable[Int,T], Itera

```

X10 supports shorthand syntax for rail construction (§12.27).

4.10.5 Future types

The interface `x10.lang.Future[T]` is the type of all future expressions. The type represents a value which when forced will return a value of type `T`. The interface makes available the following methods:

```

package x10.lang;
public interface Future[T] implements () => T {
    public def apply(): T;
    public def force(): T;
    public def forced(): Boolean;
}

```

4.11 Type inference

X10 v2.0 supports limited local type inference, permitting variable types and return types to be elided. It is a static error if an omitted type cannot be inferred or uniquely determined.

4.11.1 Variable declarations

The type of a variable declaration can be omitted if the declaration has an initializer. The inferred type of the variable is the computed type of the initializer.

4.11.2 Return types

The return type of a method can be omitted if the method has a body (i.e., is not `abstract` or `extern`). The inferred return type is the computed type of the body.

The return type of a closure can be omitted. The inferred return type is the computed type of the body.

The return type of a constructor can be omitted if the constructor has a body (i.e., is not `extern`). The inferred return type is the enclosing class type with properties bound to the arguments in the constructor's `property` statement, if any, or to the unconstrained class type.

The inferred type of a method or closure body is the least common ancestor of the types of the expressions in `return` statements in the body. If the method does not return a value, the inferred type is `Void`.

4.11.3 Type arguments

A call to a polymorphic method may omit the explicit type arguments. If the method has a type parameter `T`, the type argument corresponding to `T` is inferred to be the least common ancestor of the types of any formal parameters of type `T`.

Consider the following method:

```
def choose[T](a: T, b: T): T { ... }
```

Given `Set[T] <: Collection[T]`, `List[T] <: Collection[T]`, and `SubClass <: SuperClass`, in the following snippet, the algorithm will infer the type `Collection[Any]` for `x`.

```
def m(intSet: Set[Int], stringList: List[String]) {
  val x = choose(intSet, stringList);
  ...
}
```

And in this snippet, the algorithm should infer the type `Collection[Int]` for `y`.

```
def m(intSet: Set[Int], intList: List[Int]) {  
  val y = choose(intSet, intList);  
  ...  
}
```

Finally, in this snippet, the algorithm should infer the type `Collection{T <: SuperClass}` for `z`.

```
def m(intSet: Set[SubClass], numList: List{T <: SuperClass}) {  
  val z = choose(intSet, numList);  
  ...  
}
```

5 Variables

A variable is a storage location. X10 supports seven kinds of variables: constant *class variables* (static variables), *instance variables* (the instance fields of a class), *array components*, *method parameters*, *constructor parameters*, *exception-handler parameters* and *local variables*.

Variables are declared thus:

```
VarDeclaratorWithType ::= VarDeclaratorId ResultType
    VarDeclarator ::= VarDeclaratorId ResultType?
    VarDeclaratorId ::= Annotation* Identifier
                    | Annotation* Identifier ( VarDeclaratorList )
                    | Annotation* ( VarDeclaratorList )
    VarDeclaratorList ::= VarDeclarator ( , VarDeclarator )*
    ResultType ::= : Type
```

All variables are initialized with a value and cannot be observed without a value.

Variables whose value may not be changed after initialization are said to be *final* or *immutable*, or *constants* (§5.1). The programmer indicates that a variable is final by declaring it with the `val` keyword rather than the `var` keyword. Variables that are annotated neither `val` nor `var` are considered final.

A variable of a reference data type `T` where `T` is the name of a reference class (possibly with type arguments) always holds a reference to an instance of the class `T` or a class that is a subclass of `T`, or a `null` reference.

A variable of a rail type `Rail[T]` has as many variables as the size of the rail. A variable of a rail type `ValRail[T]` has as many variables as the size of the rail. Each of these variables is immutable and has the type `T`.

A variable of an interface type `I` always holds either a reference to a reference class implementing `I` (including possibly a boxed struct that implements `I`), or a `null` reference.

A variable of a struct data type *T* always takes as much space as necessary to represent an instance of *T*. That is structs are implemented “inline” or in an unboxed fashion.

A variable of a function type always holds a reference to enough memory to represent the constants accessed within the body of the function.

5.1 Final variables

A final variable satisfies two conditions:

- it can be assigned to at most once,
- it must be assigned to before use.

X10 follows Java language rules in this respect [5, §4.5.4,8.3.1.2,16]. Briefly, the compiler must undertake a specific analysis to statically guarantee the two properties above.

Final local variables and fields are defined by the `val` keyword. Elements of value arrays are also final.

5.2 Initial values of variables

Every variable declared to be a given type must always contain a value of that type.

Every class variable must be initialized before it is read, through the execution of an explicit initializer or a static block. Every instance variable must be initialized before it is read, through the execution of an explicit initializer or a constructor. Non-final instance variables of class type are initialized to `null`. Non-final instance variables of struct type are assumed to have an initializer that sets the value to the result of invoking the nullary constructor on the class. An initializer is required if the default initial value is not assignable to the variable’s type.

Each method and constructor parameter is initialized to the corresponding argument value provided by the invoker of the method. An exception-handling parameter is initialized to the object thrown by the exception. A local variable must be explicitly given a value by initialization or assignment, in a way that the compiler can verify using the rules for definite assignment [5, § 16].

5.3 Destructuring syntax

X10 permits a *destructuring* syntax for local variable declarations and formal parameters. At present, X10 v2.0 supports this feature only for variables of type `Point`; future versions of the language may support general pattern matching. Intuitively, this syntax allows a point to be “destructured” into its corresponding `Int` indices in a pattern-matching style. The k th declarator in a `Point VarDeclaratorList` is treated as a variable of type `Int` that is initialized with the value of the k th index of the point. The second form of the syntax permits the specification of only the index variables.

For example, the following code binds the `Int` variable `i` to 0 and `j` to 1, and the variable `p` to the point object.

```
p(i,j): Point = new Point(0,1);
```

5.4 Formal parameters

Formal parameters are always declared with a type. The variable name can be omitted if it is not to be used in the scope of the declaration.

```
Formal ::= FormalModifier* var VarDeclaratorWithType
        | FormalModifier* val VarDeclaratorWithType
        | FormalModifier* VarDeclaratorWithType
        | Type
FormalModifier ::= Annotation
                | shared
```

5.5 Local variables

Local variable declarations may have optional initializer expressions. The initializer must be a subtype of the declared type of the variable. If the variable is `final` (`val`) the type may be omitted and inferred from the initializer type (§4.11).


```

LocalDeclaration ::= LocalModifier* var LocalDeclaratorsWithType
                    ( , LocalDeclaratorsWithType )*
                    | LocalModifier* val LocalDeclarators
                    ( , LocalDeclarators )*
                    | LocalModifier* LocalDeclaratorsWithType
                    ( , LocalDeclaratorsWithType )*
LocalDeclarators ::= LocalDeclaratorsWithType
                    ::= LocalDeclaratorWithInit
LocalDeclaratorWithInit ::= VarDeclarator Init
LocalDeclaratorsWithType ::= VarDeclaratorId ( , VarDeclaratorId )* ResultType
LocalModifier ::= Annotation
                  | shared
Init ::= = Expression

```

5.6 Fields

Fields are declared either **var** (non-final, non-static), **val** (final, non-static), or **const** (final, static); the default is **val**. Field declarations may have optional initializer expressions. The initializer must be a subtype of the declared type of the variable. For **var** fields, if the initializer is omitted, the constructor must initialize the field, or else the field is initialized with **null** if a reference type, **0** if an **Int**, **0L** if a **Long**, **0.0F** if a **Float**, **0.0** if a **Double**, or **false** if a **Boolean**. It is a static error if the default value is not a member of the type (e.g., it is a static error to elide the initializer for **Int{self==1}**).

If the variable is final, the type may be omitted and inferred from the initializer type (§4.11). Mutable fields must be declared with a type.

```

FieldDeclaration ::= FieldModifier* var FieldDeclaratorsWithType
                    ( , FieldDeclaratorsWithType )*
                    | FieldModifier* const FieldDeclarators
                    ( , FieldDeclarators )*
                    | FieldModifier* val FieldDeclarators
                    ( , FieldDeclarators )*
                    | FieldModifier* FieldDeclaratorsWithType
                    ( , FieldDeclaratorsWithType )*
FieldDeclarators ::= FieldDeclaratorsWithType
                    ::= FieldDeclaratorWithInit
FieldDeclaratorId ::= Identifier
FieldDeclaratorWithInit ::= FieldDeclaratorId Init
                        | FieldDeclaratorId ResultType Init
FieldDeclaratorsWithType ::= FieldDeclaratorId ( , FieldDeclaratorId )* ResultType
FieldModifier ::= Annotation
                | static

```

5.7 Properties

Property declarations are always declared with a type and are always final (either declared **val** or by default).

```

Property ::= PropertyModifier* val Identifier ResultType
            | PropertyModifier* Identifier ResultType
PropertyModifier ::= Annotation

```

6 Objects

6.1 Basic Design

An object is an instance of a scalar class or an array type. It is created by using an allocation expression (§12.22) or an array creation (§17.4) expression, such as an array initializer.

All classes subclass from `x10.lang.Object`. This class has one property `home` of type `x10.lang.Place`. Thus all objects in X10 are located (have a place).

In X10 v2.0 an object stays resident at the place at which it was created for its entire lifetime. However, the programmer may designate certain immutable field of an object as `global`. The value of these fields is accessible at every place the object can be referenced.

X10 has no operation to dispose of a reference. Instead, the collection of all objects across all places is globally garbage collected.

Unlike Java, X10 objects do not have any synchronization information (e.g., a lock) associated with them. Instead, programmers should use atomic blocks (§15.11) for mutual exclusion and clocks (§16) for sequencing multiple parallel operations.

An object may have many references, stored in fields of objects or components of arrays. A change to an object made through one reference is visible through another reference.

Note that the creation of a remote async activity (§15.2) `A` at `P` may cause the automatic creation of references to remote objects at `P`. (A reference to a remote object is called a *remote object reference*, to a local object a *local object reference*.) For instance `A` may be created with a reference to an object at `P` held in a variable referenced by the statement in `A`. Similarly the return of a value by a `future` may cause the automatic creation of a remote object reference, incurring

some communication cost. An X10 implementation should try to ensure that the creation of a second or subsequent reference to the same remote object at a given place does not incur any (additional) communication cost.

A reference to an object carries with it the values of final global fields of the object. The implementation should try to ensure that the cost of communicating the values of final fields of an object from the place where it is hosted to any other place is not incurred more than once for each target place.

X10 does not have an operation (such as Pascal’s “dereference” operation) which returns an object given a reference to the object. Rather, most operations on object references are transparently performed on the bound object, as indicated below. The operations on objects and object references include:

- Field access (§12.4). An activity holding a reference to an object may perform this operation only if the object is local. (By contrast, an activity holding a reference to a struct may perform this operation regardless of the location of the struct, since structs can be copied freely from place to place.). The implementation should try to ensure that the cost of copying the field from the place where the object was created to the referencing place will be incurred at most once per referencing place, according to the rule for final fields discussed above.
- Method invocation (§12.6). A method may be marked `global`. A `global` method may be invoked at any place. It may access only the global fields of the object. The mutable fields of an object may be accessed only by activities operating in its home. Any activity may use an `at` statement to place-shift to the place of the object. Methods may also be marked `pinned`, `nonblocking`, `sequential`, `safe`, and `pure` (§9.6.3)..
- Casting (§12.23). An activity can perform this operation on local or remote objects, and should not incur communication costs (to bring over type information) more than once per place.
- `instanceof` operator (§12.24). An activity can perform this operation on local or remote objects, and should not incur communication costs (to bring over type information) more than once per place.
- The equality operators `==` and `!=` (§12.21). On creation, each object is associated with a globally unique identifier (`guid`). Two object references are `==` iff they refer to objects with the same `guid`.

X10 has a rather simple *distributed object model*.

The state of an object is partitioned into *global* state (a programmer defined subset of `val` fields, §9.5.2) and *non-global* state.

- Field definitions are marked with the qualifier `global` if they are intended to be included in the global state.
- If the `global` qualifier is omitted, the field is considered non-global.
- Properties and static fields are implicitly marked `global`.
- `var` fields cannot be marked `global`.

Similarly, the methods of an object may be qualified as `global`(§9.6.3) ; if they are not global they are said to be *non-global*. Global methods cannot be overridden by non-global methods.

Consider the execution of an `at (P) S` statement at a place `Q` different from `P`. Suppose `x` is an in-scope immutable local variable and contains a reference to an object `o` created at `Q`. Then within `S`, `x` is said to be a *remote reference* to `o`. References to `o` from place `Q` are said to be *local references*. X10 permits `global` fields to be read and `global` methods to be invoked through a remote reference.

Like local references, remote references are first-class entities: they may be passed as arguments to methods, returned from methods, stored in fields of objects.

Remote references may also be compared for equality (`==`). Two remote reference are equal if they are references to the same object. Equality is guaranteed to be a constant-time operation and not involve any communication.

When a remote reference to an object `o` located at place `P` is transmitted to `P` it automatically becomes a local reference to `o`. Therefore the situation in which a local reference can be compared to a remote reference simply cannot arise.

The X10 compiler ensures that non-global methods on `o` can only be invoked in a place where `here == o.home()`, i.e. the place where `o` was created.

Implementation notes Remote references to an object `o` are intended to be implemented by serializing the global state of `o` across the network, together with a globally unique id (guid). The data is deserialized at the receiver to create an implementation-level entity that is the remote reference. There is no requirement

that the implementation intern such entities; however the implementation must correctly implement equality (see below).

There is no requirement that a remote reference use only as much space as a local reference.

Local execution The semantics of `atomic` and `when` constructs requires that their bodies do not execute any `at` operations, implicitly or explicitly. Hence the compiler must establish that if a non-global method `m` is being invoked on a reference `o` in the body of such a construct, then `o` is a local reference. This can be done using place types (§4.4.2).

6.2 Examples

Assume the class declarations.

```
class C { ...}
class D {
    var f:C=null;
}
```

Now the code below must behave as described.

```
val x = new C(..);
// C object o created, reference stored in x.
at (P) {
    // In the body x contains a remote reference to o
    val f = new D();
    f.x1 = x; // remote reference stored in f.x1
    Console.OUT.println(f.x1 == x);           // must print true
    Console.OUT.println(x == x);             // must print true
    at (Q) {
        // x continues to be a remote reference to o1.
        at (P) {
            Console.OUT.println(f.x1 == x);    // must print true
            Console.OUT.println(x == x);      // must print true
        }
    }
}
```

Here is another example.

```
val x = new C(..);
// C object o created, reference stored in x.
// type of x is C{c}! if the return type
// of the constructor is C{c}.
at (P) {
    val x1 = x;
    // type of x is C{c} because of the place shift
    // introduced by at(P)
    at (x.home()) {
        // x is now bound to o through a local reference. So is x1.
        Console.out.println(x1==x); // Must print true.
        // non-global methods can be invoked on x or x1 and will
        // execute locally on o
        // type of both x and x1 is C{c}!.
    }
}
```

6.2.1 Programming Methodology

A programmer wishing to ensure that a `val` field is not serialized when the containing object is serialized (e.g. because it contains a large cache which makes sense only in the current place) must ensure the field is *not* marked global.

7 Names and packages

X10 supports Java’s mechanisms for names and packages [5, §6,§7], including `public`, `protected`, `private` and package-specific access control.

```
TypeName ::= Identifier
           | TypeName . Identifier
           | PackageName . Identifier
PackageName ::= Identifier
               | PackageName . Identifier
```

While not enforced by the compiler, classes and interfaces in the X10 library support the following naming conventions. Names of types—including classes, type parameters, and types specified by type definitions—are in CamelCase and begin with an uppercase letter. For backward compatibility with languages such as C and Java, type definitions are provided to allow primitive types such as `int` and `boolean` to be written in lowercase. Names of methods, fields, value properties, and packages are in camelCase and begin with a lowercase letter. Names of `const` fields are in all uppercase with words separated by an “_”.

8 Interfaces

X10 v2.0 interfaces are essentially the same Java interfaces [5, §9]. An interface primarily specifies signatures for public methods. It may extend multiple interfaces.

X10 permits interfaces to have properties and specify an interface invariant. This is necessary so that programmers can build dependent types on top of interfaces and not just classes.

$$\begin{aligned} \textit{NormalInterfaceDeclaration} ::= & \textit{InterfaceModifiers}^? \textbf{interface} \textit{Identifier} \\ & \textit{TypePropertyList}^? \textit{PropertyList}^? \textit{Constraint}^? \\ & \textit{ExtendsInterfaces}^? \textit{InterfaceBody} \end{aligned}$$

The invariant associated with an interface is the conjunction of the invariants associated with its superinterfaces and the invariant defined at the interface.

STATIC SEMANTICS RULE: The compiler declares an error if this constraint is not consistent (§4.4.4).

Each interface implicitly defines a nullary getter method `def p(): T` for each property `p: T`.

STATIC SEMANTICS RULE: The compiler issues a warning if an interface body contains an explicit definition for a method with this signature.

A class `C` (with properties) is said to implement an interface `I` if

- its properties contains all the properties of `I`,
- its class invariant $\textit{inv}(C)$ implies $\textit{inv}(I)$.

9 Classes

The *class declaration* has a list of type parameters, properties, a constraint (the *class invariant*), a single superclass, one or more interfaces, and a class body containing the the definition of fields, properties, methods, and member types. Each such declaration introduces a class type (§4.1).

```
NormalClassDeclaration ::= ClassModifiers? class Identifier
                           TypeParameterList? PropertyList? Guard?
                           Super? Interfaces? ClassBody

TypeParameterList ::= [ TypeParameters ]
TypeParameters   ::= TypeParameter ( , TypeParameter )*
TypeParameter   ::= Variance? Annotation* Identifier
Variance        ::= +
                   -

PropertyList     ::= ( Properties )
Properties       ::= Property ( , Property )*
Property        ::= Annotation* val? Identifier : Type

Super           ::= extends ClassType
Interfaces     ::= implements InterfaceType ( , InterfaceType )*

ClassBody      ::= ClassMember*
ClassMember    ::= ClassDeclaration
                   | InterfaceDeclaration
                   | FieldDeclaration
                   | MethodDeclaration
                   | ConstructorDeclaration
```

A type parameter declaration is given by an optional variance tag and an identifier. A type parameter must be bound to a concrete type when an instance of the class is created.

A property has a name and a type. Properties are accessible in the same way as `public final` fields.

STATIC SEMANTICS RULE: It is a compile-time error for a class defining a property `x: T` to have an ancestor class that defines a property with the name `x`.

Each class `C` defining a property `x: T` implicitly has a field

```
public val x : T;
```

and a getter method

```
public final def x()=x;
```

Each interface `I` defining a property `x: T` implicitly has a getter method

```
public def x(): T;
```

STATIC SEMANTICS RULE: It is a compile-time error for a class or interface defining a property `x: T` to have an existing method with the signature `x(): T`.

Properties are used to build dependent types from classes, as described in §4.4.

Properties are initialized by the invocation of a special `property` call in each constructor of the class:

```
property(e1,..., en);
```

The number and type of arguments to the `property` call must match the number and type of properties in the class declaration, in left to right lexical order. Each constructor is required to initialize its properties before normal termination.

The *Guard* in a class or interface declaration specifies an explicit condition on the properties of the type, and is discussed further in §9.1.

STATIC SEMANTICS RULE: Every constructor for a class defining properties `x1: T1, ..., xn: Tn` must ensure that each of the fields corresponding to the properties is definitely initialized (cf. requirement on initialization of final fields in Java) before the constructor returns.

Type parameters are used to define generic classes and interfaces, as described in §4.2.1.

Classes are structured in a single-inheritance code hierarchy, may implement multiple interfaces, may have static and instance fields, may have static and instance methods, may have constructors, may have static and instance initializers, may have static and instance inner classes and interfaces. X10 does not permit mutable static state.

Method signatures may specify checked exceptions. Method definitions may be overridden by subclasses; the overriding definition may have a declared return type that is a subclass of the return type of the definition being overridden. Multiple methods with the same name but different signatures may be provided on a class (ad hoc polymorphism). The `public/private/protected/default-protected` access modification framework may be used.

Class declarations may be used to construct class types (§4.1). Classes may have mutable fields. Instances of a class are always created in a fixed place and in X10 v2.0 stay there for the lifetime of the object. Variables declared at a class type always store a reference to the object, regardless of whether the object is local or remote.

9.1 Type invariants

There is a general recipe for constructing a list of parameters or properties $\mathbf{x}_1 : T_1\{c_1\}$, ..., $\mathbf{x}_k : T_k\{c_k\}$ that must satisfy a given (satisfiable) constraint c .

```
class Foo( $\mathbf{x}_1 : T_1\{\mathbf{x}_2 : T_2; \dots; \mathbf{x}_k : T_k; c\}$ ,
           $\mathbf{x}_2 : T_2\{\mathbf{x}_3 : T_3; \dots; \mathbf{x}_k : T_k; c\}$ ,
          ...
           $\mathbf{x}_k : T_k\{c\}$ ) {
  ...
}
```

The first type $\mathbf{x}_1 : T_1\{\mathbf{x}_2 : T_2; \dots; \mathbf{x}_k : T_k; c\}$ is consistent iff $\exists \mathbf{x}_1 : T_1, \mathbf{x}_2 : T_2, \dots, \mathbf{x}_k : T_k. c$ is consistent. The second is consistent iff

$$\forall \mathbf{x}_1 : T_1\{\mathbf{x}_2 : T_2; \dots; \mathbf{x}_k : T_k; c\} \\ \exists \mathbf{x}_2 : T_2. \exists \mathbf{x}_3 : T_3, \dots, \mathbf{x}_k : T_k. c$$

But this is always true. Similarly for the conditions for the other properties.

Thus logically every satisfiable constraint c on a list of parameters x_1, \dots, x_k can be expressed using the dependent types of x_i , provided that the constraint language is rich enough to permit existential quantifiers.

Nevertheless we will find it convenient to permit the programmer to explicitly specify a depclause after the list of properties, thus:

```
class Point(i: Int, j: Int) { ... }
class Line(start: Point, end: Point){end != start} { ... }
class Triangle (a: Line, b: Line, c: Line)
  {a.end == b.start, b.end == c.start,
   c.end == a.start} { ... }
```

Consider the definition of the class `Line`. This may be thought of as saying: the class `Line` has two fields, `start: Point` and `end: Point`. Further, every instance of `Line` must satisfy the constraint that `end != start`. Similarly for the other class definitions.

In the general case, the production for *NormalClassDeclaration* specifies that the list of properties may be followed by a *Guard*.

$$\begin{aligned} \text{NormalClassDeclaration} \quad ::= \quad & \text{ClassModifiers}^? \text{ class Identifier} \\ & \text{TypeParameterList}^? \text{ PropertyList}^? \text{ Guard}^? \\ & \text{Extends}^? \text{ Interfaces}^? \text{ ClassBody} \end{aligned}$$

$$\begin{aligned} \text{NormalInterfaceDeclaration} \quad ::= \quad & \text{InterfaceModifiers}^? \text{ interface Identifier} \\ & \text{TypeParameterList}^? \text{ PropertyList}^? \text{ Guard}^? \\ & \text{ExtendsInterfaces}^? \text{ InterfaceBody} \end{aligned}$$

All the properties in the list, together with inherited properties, may appear in the *Guard*. A guard c with property list $x_1: T_1, \dots, x_n: T_n$ for a class C is said to be consistent if each of the T_i are consistent and the constraint

$$\exists x_1: T_1, \dots, x_n: T_n, \text{ self: } C. c$$

is valid (always true).

The guard is an invariant on all instances of the class or interface.

With every defined class or interface T we associate a *type invariant* $inv(T)$ as follows. The type invariant associated with `x10.lang.Object` is `true`.

The type invariant associated with any interface I that extends interfaces I_1, \dots, I_k and defines properties $x_1: P_1, \dots, x_n: P_n$ and specifies a guard c is given by:

$$\begin{array}{l} \text{inv}(\mathbf{I}_1), \dots, \text{inv}(\mathbf{I}_k), \\ \text{self.x}_1: \mathbf{P}_1, \dots, \text{self.x}_n: \mathbf{P}_n, \mathbf{c} \end{array}$$

Similarly the type invariant associated with any class \mathbf{C} that implements interfaces $\mathbf{I}_1, \dots, \mathbf{I}_k$, extends class \mathbf{D} and defines properties $\mathbf{x}_1: \mathbf{P}_1, \dots, \mathbf{x}_n: \mathbf{P}_n$ and specifies a guard \mathbf{c} is given by:

$$\begin{array}{l} \text{inv}(\mathbf{D}), \text{inv}(\mathbf{I}_1), \dots, \text{inv}(\mathbf{I}_k), \\ \text{self.x}_1: \mathbf{P}_1, \dots, \text{self.x}_n: \mathbf{P}_n, \mathbf{c} \end{array}$$

It is required that the type invariant associated with a class entail the type invariants of each interface that it implements.

It is guaranteed that for any variable \mathbf{v} of type $\mathbf{T}\{\mathbf{c}\}$ (where \mathbf{T} is an interface name or a class name) the only objects \mathbf{o} that may be stored in \mathbf{v} are such that \mathbf{o} satisfies $\text{inv}(\mathbf{T}[\mathbf{o}/\mathbf{this}]) \wedge \mathbf{c}[\mathbf{o}/\mathbf{self}]$.

9.2 implements and extends clauses

Consider a class definition

$$\begin{array}{l} \text{ClassModifiers}^? \\ \text{class } \mathbf{C}(\mathbf{x}_1: \mathbf{P}_1, \dots, \mathbf{x}_n: \mathbf{P}_n) \text{ extends } \mathbf{D}\{\mathbf{d}\} \\ \quad \text{implements } \mathbf{I}_1\{\mathbf{c}_1\}, \dots, \mathbf{I}_k\{\mathbf{c}_k\} \\ \text{ClassBody} \end{array}$$

Each of the following static semantics rules must be satisfied:

STATIC SEMANTICS RULE (Int-implements): The type invariant $\text{inv}(\mathbf{C})$ of \mathbf{C} must entail $\mathbf{c}_i[\mathbf{this}/\mathbf{self}]$ for each i in $\{1, \dots, k\}$

STATIC SEMANTICS RULE (Super-extends): The return type \mathbf{c} of each constructor in ClassBody must entail \mathbf{d} .

9.3 Constructor definitions

A constructor for a class \mathbf{C} is guaranteed to return an object of the class on successful termination. This object must satisfy $\text{inv}(\mathbf{C})$, the class invariant associated with \mathbf{C} (§9.1). However, often the objects returned by a constructor may satisfy

stronger properties than the class invariant. X10's dependent type system permits these extra properties to be asserted with the constructor in the form of a constrained type (the “return type” of the constructor):

```

ConstructorDeclarator ::= def this TypeParameterList? ( FormalParameterList? )
                        ReturnType? Guard? Throws?
    ReturnType ::= : Type
    Guard ::= "{ " DepExpression " } "
    Throws ::= throws ExceptionType ( , ExceptionType ) *
    ExceptionType ::= ClassBaseType Annotation*

```

The parameter list for the constructor may specify a *guard* that is to be satisfied by the parameters to the list.

Example 9.3.1 Here is another example, constructed as a simplified version of `x10.lang.Region`.

```

type MyRegion(n:Int)=MyRegion{self.rank==n};
class MyRegion(rank:Int) {
  def this(r:Int):MyRegion(r) {
    property(r);
  }
  def this(diag:ValRail[Int]):MyRegion(diag.length){
    ...
  }
  def union(r:MyRegion(n)):MyRegion(n) { ...}
  ...
}

```

The first constructor returns the empty region of rank `r`. The second constructor takes a `ValRail[Int]` of arbitrary length `n` and returns a `MyRegion(n)` (intended to represent the set of points in the rectangular parallelopiped between the origin and the `diag`.)

Now the following code type checks:

```

val R1 = new MyRegion([4,4,4]); // R1's type is MyRegion(3)
val R2 = new MyRegion([5,4,1]); // R2's type is MyRegion(3)

```

Hence the following code type checks and infers that `R3`'s type is `MyRegion(3)`:

```
val R3 = R1.union(R2);           // R3's type is MyRegion(3)
```

□

STATIC SEMANTICS RULE (Super-invoke): Let C be a class with properties $p_1 : P_1, \dots, p_n : P_n$, invariant c extending the constrained type $D\{d\}$ (where D is the name of a class).

For every constructor in C the compiler checks that the call to super invokes a constructor for D whose return type is strong enough to entail d . Specifically, if the call to super is of the form $\text{super}(e_1, \dots, e_k)$ and the static type of each expression e_i is S_i , and the invocation is statically resolved to a constructor $\text{def this}(x_1 : T_1, \dots, x_k : T_k) \{c\} : D\{d_1\}$ then it must be the case that

$$\begin{aligned} x_1 : S_1, \dots, x_i : S_i &\vdash x_i : T_i \quad (\text{for } i \in \{1, \dots, k\}) \\ x_1 : S_1, \dots, x_k : S_k &\vdash c \\ d_1[a/\text{self}], x_1 : S_1, \dots, x_k : S_k &\vdash d[a/\text{self}] \end{aligned}$$

where a is a constant that does not appear in $x_1 : S_1 \wedge \dots \wedge x_k : S_k$.

STATIC SEMANTICS RULE (Constructor return): The compiler checks that every constructor for C ensures that the properties p_1, \dots, p_n are initialized with values which satisfy $t(C)$, and its own return type c' as follows. In each constructor, the compiler checks that the static types T_i of the expressions e_i assigned to p_i are such that the following is true:

$$p_1 : T_1, \dots, p_n : T_n \vdash t(C) \wedge c'$$

(Note that for the assignment of e_i to p_i to be type-correct it must be the case that $p_i : T_i \wedge p_i : P_i$.)

STATIC SEMANTICS RULE (Constructor invocation): The compiler must check that every invocation $C(e_1, \dots, e_n)$ to a constructor is type correct: each argument e_i must have a static type that is a subtype of the declared type T_i for the i th argument of the constructor, and the conjunction of static types of the argument must entail the *Guard* in the parameter list of the constructor.

9.4 proto qualifier on types

X10 ensures that every variable must have a value consistent with its type before it is read.

For local variables, this is ensured by using a pre-specified static analysis to ensure that every local variable is written into before it is read. Type-checking of assignment ensures the value written is consistent with the static type of the variable.

For fields, this is ensured by introducing a form of ownership types called *incomplete types* to address the *escaping-this* problem. To permit flexibility in writing constructors, X10 v1.7 permits `this` to be used in a constructor as a reference to the object currently being constructed. Unfortunately there are no restrictions on the usage of `this`. In particular, this reference can be permitted to escape: it may be stored in variables on the heap (thereby permitting concurrently executing activities to read the value of fields that may not yet have been initialized), passed as an argument to method invocations, or used as the target for a method invocation. Indeed, the method may be invoked in a super constructor, and may have been overridden at a subclass, guaranteeing that accesses to fields defined in the subclass are accesses to uninitialized variables. For instance an immutable field may be observed containing a value (the value the field was initialized with) which may be different from the value it will contain once the constructor has returned.

Incomplete types are designed with the following goals:

- Guarantee that fields are not read before they are initialized.
- Allow the creation of immutable cyclic object graphs.¹ This requires that it be possible to pass an object under construction into a constructor invocation.
- Allow appropriate user-defined methods can be called during object creation (so that the transformation between the values supplied as parameters to a constructor and the values actually placed in fields is determined by arbitrary user-defined code).
- Keep the design minimally invasive. Most programmers should not have to be concerned about this problem.
- Ensure that there is no runtime overhead.

These goals are met by introducing incomplete types through the type qualifier `proto`. Types of the form `proto T` are said to be *incomplete types*; types that do

¹(Mutable graphs can be created without escaping `this` by initializing the backpointer to `null` and then changing it later.

not have the qualifier are said to be *complete*. Say that an object *o* is *confined* to a given activity *A* if it can be reached only from stack frames of *A* or from objects which are, recursively, confined to *A*. Thus confined objects cannot be accessed by activities other than *A*.

Incomplete types ensure that objects whose constructors have not exited are confined. Further, all references to such objects on the stack are contained in variables of incomplete types. The compiler does not permit the fields of variables of incomplete types to be read. Thus incomplete types permit the construction of graphs of objects while ensuring that these objects are confined and their fields are not read during construction.

The return value of a constructor for class *C* that takes no incomplete arguments is (a subtype of) *C*, that is, a complete type. It will point only to completed objects. It can now be assigned to any (type-consistent) field of any object, that is, it is now allowed to escape.

proto Rules

For every type *T* (where *T* is not a type variable), we introduce the type `proto T`.

There is no relationship between types *T* and `proto T` – neither is a subtype of the other.²

Incomplete types are permitted to occur only as types of method parameters or local variables or as return types for methods and constructors. They may not occur in (the source or target of) cast statements, `extends` or `implements` clauses, `catch` clauses, or as types of class fields.

Within the body of a class *C* the type of `this` in constructors, instance initializers and instance variable initializers is `proto C`.

Let *v* be a value of type `proto C`, for some class *C*.

No fields of *v* can be read. (This is the defining property of `proto` types.) However, *v*'s (accessible) instance fields can be assigned.

²Clearly, a value of type `proto T` cannot be used anywhere that a *T* is needed, since its fields cannot be read. As discussed below, an incomplete value *v* can be assigned to a field *f* of an object *o* only if *o* is incomplete. This ensures that *v* cannot escape through this assignment. A completed value *p* cannot be substituted for *o* – it may permit *v* to escape through an assignment to its field. Therefore *T* cannot be a subtype of `proto T`.

v can be assigned to an instance field $o.f$ only if f is of some type S such that $T <: S$ and o has an incomplete type.

v can be assigned to local variables only if they are of some type `proto S` (such that $T <: S$).

Instance methods of class C may be qualified with `proto` (these methods are called *incomplete methods*). The type of `this` in incomplete methods is `proto C`. Incomplete methods can be overridden only by incomplete methods. Only incomplete methods can be invoked on v . Incomplete methods which do not take an argument of incomplete type can be invoked on completed values.

v can be passed as argument into a constructor or method call, or returned from a method. The return type of a method taking an argument at an incomplete type must be `void` or incomplete. The return type of a constructor taking an argument at a `proto` type must be incomplete.

A generic class (method) type parameter T can be instantiated with the type `proto S` (where S is not a type parameter itself), provided that the class (method) body satisfies the conditions above for `proto S`.

During code generation, the type `proto T` is treated as if it were T . That is, there is no run-time cost to `proto` types.

The invariants maintained by the design are as follows. Say that an object field or stack variable (local variable) contains an incomplete value if a value of type `proto T` (for some T) was written into it.

- If an object o has a field containing an incomplete value v , then either v 's constructor has exited or o is confined. Further, every reference to o on the stack is held at an incomplete type.
- If a stack variable contains an incomplete value, then the variable's type is incomplete.

Say that a constructor invocation for a class C on the call stack is a *root* if it takes no incomplete arguments. Such a constructor invocation will return an object of type C whose fields may point to an arbitrary graph of newly created objects (objects created by the activity after the constructor invocation). Since the object returned is at type C – and not `proto C` – It may be assigned to any field of any object on the heap of type D such that $C <: D$. It is no longer confined. Thus the “magic moment” when an incomplete value becomes complete is when the last constructor for any incomplete value it references (including itself) returns.

Example

Example 9.4.1 This example shows how to create a fixed-size circular buffer. (Its pointer structure is immutable, though the contents of each field are mutable.)

```
class CircularBuffer[A] {
  var a: A;
  val next: CircularBuffer[A];
  private def this(x: proto CircularBuffer[A]): proto CircularBuffer[A] {
    next = x;
  }
  def this(var n: Int) {
    var temp: proto CircularBuffer[A] = this;
    while (--n > 0)
      temp = new CircularBuffer[A](temp);
    next = temp;
  }
}
```

□

9.5 Field definitions

A class may have zero or more mutable or immutable fields. No two fields declared in a class may have the same name.

Fields may be marked `static`. Only one instance of such a field exists, and it may be accessed through the name of the class in which it is defined (§9.7). Fields not marked `static` are said to be *instance* fields. One copy of such a field exists for every instance of the class.

To avoid an ambiguity, it is a static error for a class to declare a field with a function type (§4.5) with the same name and signature as a method of the same class.

9.5.1 Field hiding

A subclass that defines a field `f` hides any field `f` declared in a superclass, regardless of their types. The superclass field `f` may be accessed within the body of the subclass via the reference `super.f`.

9.5.2 Field qualifiers

global qualifier

A field may be declared `global`.

FieldModifier ::= `global`

A `global` field must be immutable. It may be read from any place. Properties and static fields are implicitly marked `global`. Fields not marked `global` cannot be overridden by fields marked `global`.

9.6 Method definitions

X10 permits guarded method definitions.

```

MethodDeclaration ::= MethodHeader ;
                  | MethodHeader = ClosureBody
MethodHeader ::= MethodModifiers? def Identifier TypeParameters?
                ( FormalParameterList? ) Guard?
                ReturnType? Throws?

```

A formal parameter may optionally have a `val` or `var` modifier (default: `val`). The body of the method is executed in an environment in which each formal parameter corresponds to a local variable and is initialized with the value of the actual parameter. The local variable is mutable if and only if the parameter is a `var` parameter.

The guard (specified by *Guard*) specifies a constraint *c* on the properties of the class *C* on which the method is being defined. The method exists only for those instances of *C* which satisfy *c*. It is illegal for code to invoke the method on objects whose static type is not a subtype of *C*{*c*}.

STATIC SEMANTICS RULE: The compiler checks that every method invocation `o.m(e1, ..., en)` for a method is type correct. Each argument *e_i* must have a static type *S_i* that is a subtype of the declared type *T_i* for the *i*th argument of the method, and the conjunction of static types of the arguments must entail the guard in the parameter list of the method.

The compiler checks that in every method invocation `o.m(e1, ..., en)` the static type of *o*, *S*, is a subtype of *C*{*c*}, where the method is defined in class *C* and the guard for *m* is equivalent to *c*.

Finally, if the declared return type of the method is $D\{d\}$, the return type computed for the call is $D\{a: S; x_1: S_1; \dots; x_n: S_n; d[a/\text{this}]\}$, where a is a new variable that does not occur in d, S, S_1, \dots, S_n , and x_1, \dots, x_n are the formal parameters of the method. The method body is either an expression, a block of statements, or a block ending with an expression.

Example 9.6.1 Consider the program:

```
type Point(r:Int)=Point{self.rank==r};
final public class Point(rank: Int) implements (Int) => Int {
  public global val coords: ValRail[Int](rank);
  public global safe def apply(i: Int) = coords(i);
  public global safe def coords() = coords;
  public global safe operator - this: Point(rank)
    = Point.make(rank, (i:Int)=>-this.coords(i));
  public global safe operator this + (that: Point(rank)): Point(rank)
    = Point.make(rank, (i:Int)=> this.coords(i) + that.coords(i));
  ...
}
```

The following code fragment will typecheck:

```
s: Point(3) = new Point([1,2,3]);
t: Point(3) = new Point([-1,-1,-1]);
u: Point(3) = s + u;
```

□

9.6.1 Property methods

A method declared with the modifier `property` may be used in constraints. A property method declared in a class must have a body and must not be `void`. The body of the method must consist of only a single `return` statement or a single expression. It is a static error if the expression cannot be represented in the constraint system.

The expression may contain invocations of other properties. It is the responsibility of the programmer to ensure that the evaluation of a property terminates at compile-time, otherwise the type-checker will not terminate and the program will fail to compile.

Property methods in classes are implicitly `final`; they cannot be overridden. Property methods are also implicitly `global`.

A property method definition may omit the formal parameters and the `def` keyword. That is, the following are equivalent:

```
property def rail(): Boolean = rect && onePlace == here && zeroBased;
property rail: Boolean = rect && onePlace == here && zeroBased;
```

9.6.2 Method overloading, overriding, hiding, shadowing and obscuring

The definitions of method overloading, overriding, hiding, shadowing and obscuring in X10 are the same as in Java, modulo the following considerations motivated by type parameters and dependent types.

Two or more methods of a class or interface may have the same name if they have a different number of type parameters, or they have value parameters of different types.

X10 v2.0 does not permit overloading based on constraints.

The definition of a method declaration m_1 “having the same signature as” a method declaration m_2 involves identity of types.

The *constraint erasure* of a type T is defined as follows. The constraint erasure of (a) a class, interface or struct type T is T ; (b) a type $T\{c\}$ is the constraint erasure of T ; (b) a type $T[S_1, \dots, S_n]$ is $T'[S_1', \dots, S_n']$ where each primed type is the erasure of the corresponding unprimed type. Two methods are said to have *the same signature* if (a) they have the same number of type parameters, (b) they have the same number of formal (value) parameters, and (c) for each formal parameter the constraint erasure of its types are equivalent. It is a compile-time error for there to be two methods with the same name and same signature in a class (either defined in that class or in a superclass).

STATIC SEMANTICS RULE: A class C may not have two declarations for a method named m —either defined at C or inherited:

```
def m[X1, ..., Xm](v1: T1, ..., vn: Tn){tc}: T {...}
def m[X1, ..., Xm](v1: S1, ..., vn: Sn){sc}: S {...}
```

if it is the case that the constraint erasures of the types T_1, \dots, T_n are equivalent to the constraint erasures of the types S_1, \dots, S_n respectively.

In addition, the guard of a overriding method must be no stronger than the guard of the overridden method. This ensures that any virtual call to the method satisfies the guard of the callee.

STATIC SEMANTICS RULE: If a class C overrides a method of a class or interface B , the guard of the method in B must entail the guard of the method in C .

A class C inherits from its direct superclass and superinterfaces all their methods visible according to the access modifiers of the superclass/superinterfaces that are not hidden or overridden. A method M_1 in a class C overrides a method M_2 in a superclass D if M_1 and M_2 have the same signature. Methods are overridden on a signature-by-signature basis.

A method invocation $o.m(e_1, \dots, e_n)$ is said to have the *static signature* $\langle T, T_1, \dots, T_n \rangle$ where T is the static type of o , and T_1, \dots, T_n are the static types of e_1, \dots, e_n , respectively. As in Java, it must be the case that the compiler can determine a single method defined on T with argument type T_1, \dots, T_n ; otherwise, a compile-time error is declared. However, unlike Java, the X10 type T may be a dependent type $C\{c\}$. Therefore, given a class definition for C we must determine which methods of C are available at a type $C\{c\}$. But the answer to this question is clear: exactly those methods defined on C are available at the type $C\{c\}$ whose guard d is implied by c .

9.6.3 Method qualifiers

atomic qualifier

A method may be declared **atomic**.

MethodModifier ::= **atomic**

Such a method is treated as if the statement in its body is wrapped implicitly in an **atomic** statement.

global qualifier

A method may be declared **global**.

MethodModifier ::= **global**

A **global** method can be invoked on an object o in any place. The body of such a method is type-checked without assuming that `here==this.home`. This permits

global fields of `o` to be accessed, but not local fields. The programmer must insert an explicit `at(this)...` to get to the place where the object lives and access the field.

global methods can be overridden only by methods also marked `global`.

pinned qualifier

A method may be declared `pinned`.

MethodModifier ::= `pinned`

A `pinned` method may not contain any `at` statement or expression whose place argument is not statically equivalent to `here`. It must call only `pinned` methods. That is, a `pinned` method does not cause any communication.

`pinned` methods can be overridden only by methods marked `pinned`.

nonblocking qualifier

A method may be declared `nonblocking`.

MethodModifier ::= `nonblocking`

A `nonblocking` method may not contain any `when` statement whose condition is not statically equivalent to `true`. It must call only `nonblocking` methods. That is, a `nonblocking` method does not block.

`nonblocking` methods can be overridden only by methods marked `nonblocking`.

sequential qualifier

A method may be declared `sequential`.

MethodModifier ::= `sequential`

A `sequential` method may not contain any `async` statement. It must call only `sequential` methods. That is, a `sequential` method does not spawn any activity.

`sequential` methods can be overridden only by methods marked `sequential`.

safe qualifier

A method may be declared *safe*.

MethodModifier ::= *safe*

The *safe* annotation is considered shorthand for *pinned nonblocking sequential*.

9.7 Static initialization

The X10 runtime implements the following procedure to ensure reliable initialization of the static state of classes.

Execution commences with a single thread executing the *initialization* phase of an X10 computation at place 0. This phase must complete successfully before the body of the main method is executed.

The initialization phase must be thought of as if it is implemented in the following fashion: (The implementation may do something more efficient as long as it is faithful to this semantics.)

```

Within the scope of a new finish
for every static field f of every class C
  (with type T and initializer e):
  async {
    val l = e;
    ateach (Dist.makeUnique()) {
      assign l to the static f field of
        the local C class object;
      mark the f field of the local C
        class object as initialized;
    }
  }

```

During this phase, any read of a static field *C.f* (where *f* is of type *T*) is replaced by a call to the method *C.read_f():T* defined on class *C* as follows

```

def read_f():T {
  await (initialized(C.f));
  return C.f;
}

```

If all these activities terminate normally, all static field have a legal value (per their type), and the finish terminates normally. If any activity throws an exception, the finish throws an exception. Since no user code is executing which can catch exceptions thrown by the finish, the exceptions are printed on the console, and computation aborts.

If the activities deadlock, the implementation deadlocks.

In all cases, the main method is executed only once all static fields have been initialized correctly.

Since static state is immutable and is replicated to all places via the initialization phase as described above, it can be accessed from any place.

10 Structs

An instance of a class *C* (an *object*) is represented in X10 as a contiguously allocated chunk of words in the heap, containing the fields of the object as well as one or more words used in method lookup (itable/vtable). Variables with base type *C* (or a supertype of *C*) are implemented as cells with enough memory to hold a *reference* to the object. The size of a reference (32 bits or 64 bits) depends on the underlying operating system.

For many high-performance programming idioms, the overhead of one extra level of indirection represented by an object is not acceptable. For instance, a programmer may wish to define a type *Complex* (consisting of two double fields) and require that instances of this type be represented precisely as these two fields. A variable or field of type *complex* should, therefore, contain enough space to store two doubles. An array of *complex* of size *N* should store $2*N$ doubles. Method invocations should be resolved statically so that there is no need to store vtable/itable words with each instance. Parameters of type *complex* should be passed inline to a method as two doubles. If a method's return type is *complex* the method should return two doubles on the stack. Two values of this type should be equal precisely when the two doubles are equal (structural equality).

X10 supports the notion of *structs* which are precisely objects that can be implemented inline with a contiguous chunk of memory representing their fields, without any vtable/itable. Structs are introduced by struct definitions. struct definitions look very similar to class definitions, but have additional restrictions.

Implementation Note In X10 v2.0, structs are only implemented as inline headerless objects when using the C++ backend. In the Java backend, structs are mapped directly to Java classes and therefore have the same space and indirection characteristics as classes.

10.1 Struct declaration

X10 supports user-defined primitives (called *structs*). Like classes, structs define zero or more fields and zero or more methods, and may implement zero or more interfaces. A struct has the same modifiers as a class. However, structs are implicitly final and do *not* participate in any code inheritance relation. (This makes structs very easy to implement, without vtables.)

```

StructModifiers?
struct C[X1, ..., Xn](p1:T1, ..., pn:Tn){c}
    implements I1, ..., Ik {
    StructBody
}

```

Each field and method in a struct is implicitly marked `global`.

The size of a variable of struct type C is the size of the fields defined at C (up to alignment considerations). No extra space is allocated for a vtable or an itable. This means that unlike classes, structs cannot be defined recursively. That is, a struct S cannot contain a field of type S, or a field of struct type T which, recursively, contains a field of type S.

- More precisely, we require that the set of *size equations* for all structs and classes must have a unique solution. A size equation for a struct S is defined as follows. Assume S has m fields of type S_i (for i in $0, \dots, m-1$), and n fields of type (class) C_j (for j in $0, \dots, n-1$). Then the size equation for S is

$$\text{size}(S) = \text{size}(S_0) + \dots + \text{size}(S_{m-1}) + \text{size}(C_0) + \dots + \text{size}(C_{n-1})$$

The size equation for a class C is just $\text{size}(C) = \text{AddressSize}$, where `AddressSize` is a compile-time parameter.

Values of a struct C type can be created by invoking a constructor defined in C, but without prefixing it with `new`.

Constrained types can be built on top of the base C in the same way as they can be built on top of a class D. In struct `C[T1, ..., Tn]{c}`, the type of `self` in `c` is `C[T1, ..., Tn]`.

10.2 Boxing of structs

If a struct *S* implements an interface *I* (e.g., *Any*), a value *v* of type *S* can be assigned to a variable of type *I*. The implementation creates an object *o* that is an instance of an anonymous class implementing *I* and containing *v*. The result of invoking a method of *I* on *o* is the same as invoking it on *v*. This operation is termed *auto-boxing*.

In a generic class or struct obtained by instantiating a type parameter *T* with a struct *S*, variables declared at type *T* in the body of the class are not boxed. Rather they are implemented as if they were declared at type *S*.

10.3 Implementation of Any methods

Unlike objects, structs do not have global identity. Instead, two structs are equal (==) if and only if their corresponding fields are equal (==). This is the central property of structs.

All structs implement `x10.lang.Any`. All structs have the following methods implicitly defined on them:

```
property def home()=here;  
property def at(Place)=true;  
property def at(Object)=true;
```

It is an error for a programmer to attempt to define them.

Structs are required to implement the following methods:

```
public global safe def equals(Any):Boolean;  
public global safe def hashCode():Int;  
public global safe def typeName():String;  
public global safe def toString():String;
```

These methods are defined automatically if they are not supplied by the programmer. A programmer who provides an explicit implementation of `equals(Any)` for a struct *S* should also consider supplying a definition for `equals(S):Boolean`. This will often yield better performance since the cost of an upcast to *Any* and then a downcast to *S* can be avoided.

Expressions of a struct type may be used in `instanceof` and as expressions.

10.4 “Primitives”

The package `x10.lang` provides the following structs. Most of the functionality of these structs is implemented natively.

```
boolean, char,  
byte, short, int, long  
float, double  
ubyte, ushort, uint, ulong
```

10.5 Generic programming with structs

An unconstrained type variable `X` can be instantiated with `Object` or its subclasses or structs or functions.

Within a generic class, all the operations of `Any` are available on a variable of type `X`. Additionally, variables of type `X` may be used with `==`, `!=`, in `instanceof`, and casts.

The programmer must be aware of the different interpretations of equality for structs and classes and ensure that the code is correctly written for both cases. If necessary the programmer can write code that distinguishes between the two cases (a type parameter `X` is instantiated to a struct or not) as follows:

```
val x:X = ...;  
if (x instanceof Object) { // x is a real object  
    val x2 = x as Object; // this cast will always succeed.  
    ...  
} else { // x is a struct  
    ...  
}
```

10.6 Programming Methodology

A programmer should by default organize his/her code in a class hierarchy, providing structs only in those well-thought situations where concrete types are appropriate.

10.6.1 Compatibility Note

A value class in X10 v1.7 can often be translated into a struct in X10 2.0. The crucial conditions to be checked manually are:

- A struct is of bounded size.
- Each method is global.
- The class is final.

If these conditions are not met, the value class should be converted into a class with global fields and methods.

10.6.2 Examples

An example illustrating pairing:

```
struct Pair[S,T] {  
  val x: S;  
  val y: T;  
  def this(x: S, y: T) {  
    this.x=x;  
    this.y=y;  
  }  
  def x()=x;  
  def y()=y;  
  final def hashCode() = x.hashCode() + y.hashCode();  
  final def equalsX[U](o:Pair[S,U]) = x==o.x;  
  final def equalsY[U](o:Pair[U,T]) = y==o.y;  
  final def equals(that:Any) = this == that;  
  final def equals(that:Pair[X,Y]) = this==that;  
}
```

The following types all make sense:

- `Pair[Complex, String]`: A struct with two fields, one inlined field of type `Complex` and another of type `String`.

- `Pair[Complex, Int]`: A struct with two fields, one inlined field of type `Complex` and another of type `Int`.

The definition of `x10.lang.Complex` provides a good example of the use of structs.

11 Functions

11.1 Overview

The runtime entities in X10 are of three kinds: *structs*, *objects*, and *functions*. This section is concerned with functions and their types – how they are created, and what operations can be performed on them.

Intuitively, a function is a piece of code which can be applied to a set of arguments to produce a value. The application may not terminate, or may terminate abruptly. Functions may throw checked and unchecked exceptions. The body of a function may be any X10 expression: hence a function evaluation may spawn multiple activities, read and write mutable locations, wait until memory locations contain a desired value, and execute over multiple places. In particular, function evaluation may be non-deterministic. When applied to the same input twice, a function may yield two different results.

It is a limitation of X10 v2.0 that functions do not support type arguments. This limitation may be removed in future versions of the language.

A *function literal* $(x_1:T_1, \dots, x_n:T_n)\{c\}:T \Rightarrow e$ creates a function of type $(x_1:T_1, \dots, x_n:T_n)\{c\} \Rightarrow T$ (§4.5). The body e of such an expression is type-checked in an environment in which c is true. At runtime, function execution results in the evaluation of e in an environment in which each formal is bound to the given actual parameter.

The *method selector expression* $e.m.(x_1:T_1, \dots, x_n:T_n)$ (§11.3) permits the specification of the function underlying the method m , which takes arguments of type $(x_1:T_1, \dots, x_n:T_n)$. Within this function, *this* is bound to the result of evaluating e .

Function types may be used in `implements` clauses of class definitions. Instances of such classes may be used as functions of the given type. Indeed, an object may

behave like any (fixed) number of functions, since the class it is an instance of may implement any (fixed) number of function types.

11.2 Function Literals

X10 provides first-class, typed functions, including *closures*, *operator functions*, and *method selectors*.

```

ClosureExpression ::= ( Formals? )
                   Guard? ReturnType? Throws? => ClosureBody
ClosureBody      ::= Expression
                   | { Statement* }
                   | { Statement* Expression }

```

Functions have zero or more formal parameters, an optional return type and optional set of exceptions throws by the body. The body has the same syntax as a method body; it may be either an expression, a block of statements, or a block terminated by an expression to return. In particular, a value may be returned from the body of the function using a return statement (§13.15). The type of a function is a function type (§4.5). In some cases the return type *T* is also optional and defaults to the type of the body. If a formal *x_i* does not occur in any *T_j*, *c*, *T* or *e*, the declaration *x_i : T_i* may be replaced by just *T_i*.

As with methods, a function may declare a guard to constrain the actual parameters with which it may be invoked. The guard may refer to the type parameters, formal parameters, and any final variables in scope at the function expression.

The body of the function is evaluated when the function is invoked by a call expression (§12.6), not at the function's place in the program text.

As with methods, a function with return type *Void* cannot have a terminating expression. If the return type is omitted, it is inferred, as described in §4.11. It is a static error if the return type cannot be inferred.

Example 11.2.1 The following method takes a function parameter and uses it to test each element of the list, returning the first matching element.

```

def find[T](f: (T) => Boolean, xs: List[T]): T = {
  for (x: T in xs)
    if (f(x)) return x;
}

```

```

    null
  }

```

The method may be invoked thus:

```

xs: List[Int] = ...;
x: Int = find((x: Int) => x>0, xs);

```

□

As with a normal method, the function may have a `throws` clause. It is a static error if the body of the function throws a checked exception that is not declared in the function's `throws` clause.

11.2.1 Outer variable access

In a function $(x_1: T_1, \dots, x_n: T_n)\{c\} \Rightarrow \{s\}$ the types T_i , the guard c and the body s may access fields of enclosing classes and local variables and type parameters declared in an outer scope.

Recall that languages such as Java require that methods may access only those local variables declared in an enclosing scope (“outer variables”) which are final. This is valuable in preventing accidental races between multiple functions reading and writing the same outer variable. At the same time, it is desirable to support the following common idiom of expression:

```

def allPositive(c: Collection): Boolean {
  shared var result: Boolean = true;
  c.applyToAll((x: Int) => { if (x < 0) atomic {result=false;}});
  return result;
}

```

This motivates the following rule:

STATIC SEMANTICS RULE: In an expression $(x_1: T_1, \dots, x_n: T_n) \Rightarrow e$, any outer local variable accessed by e must be final or must be declared as `shared` (§15.10).

The function body may refer to instances of enclosing classes using the syntax `C.this`, where C is the name of the enclosing class.

NOTE: The main activity may run in parallel with any functions it creates. Hence even the read of an outer variable by the body of a function may result in a race condition. Since functions are first-class, the analysis of whether a function may execute in parallel with the activity that created it may be difficult.

11.3 Methods selectors

A method selector expression allows a method to be used as a first-class function.

$$\begin{array}{l} \text{MethodSelector} ::= \text{Primary} . \text{MethodName} . \text{TypeParameters}^? (\text{Formals}^?) \\ \quad \quad \quad | \quad \text{TypeName} . \text{MethodName} . \text{TypeParameters}^? (\text{Formals}^?) \end{array}$$

The *method selector expression* $e.m.(T_1, \dots, T_n)$ is type correct only if it is the case that the static type of e is a class or struct or interface with a method $m(x_1:T_1, \dots, x_n:T_n)\{c\}:T$ defined on it (for some x_1, \dots, x_n, c, T). At run-time the evaluation of this expression evaluates e to a value v and creates a function f which, when applied to an argument list (a_1, \dots, a_n) (of the right type) yields the value obtained by evaluating $v.m(a_1, \dots, a_n)$.

Thus, the method selector

$$e.m.[X_1, \dots, X_m](T_1, \dots, T_n)$$

behaves as if it were the function

$$\begin{array}{l} ((v:T) \Rightarrow \\ \quad [X_1, \dots, X_m](x_1:T_1, \dots, x_n:T_n) \Rightarrow v.m[X_1, \dots, X_m](x_1, \dots, x_n)) \\ (e) \end{array}$$

NOTE: Because of overloading, a method name is not sufficient to uniquely identify a function for a given class (in Java-like languages). One needs the argument type information as well. The selector syntax (dot) is used to distinguish $e.m()$ (a method invocation on e of method named m with no arguments) from $e.m.()$ (the function bound to the method).

A static method provides a binding from a name to a function that is independent of any instance of a class; rather it is associated with the class itself. The static function selector $T.m.(T_1, \dots, T_n)$ denotes the function bound to the static method named m , with argument types (T_1, \dots, T_n) for the type T . The return type of the function is specified by the declaration of $T.m$.

Users of a function type do not care whether a function was defined directly (using the function syntax), or obtained via (static or instance) function selectors.

NOTE: Design note: The function selector syntax is consistent with the reinterpretation of the usual method invocation syntax $e.m(e_1, \dots, e_n)$ into a function specifier, $e.m$, applied to a tuple of arguments (e_1, \dots, e_n) . Note that the receiver is not treated as “an extra argument” to the function. That would break the above approach.

11.4 Operator functions

Every operator (e.g., $+$, $-$, $*$, $/$, ...) has a family of functions, one for each type on which the operator is defined. The function can be selected using the “.” syntax:

$$\begin{array}{lcl} \text{OperatorFunction} & ::= & \text{TypeName} . \text{Operator} (\text{Formals}^?) \\ & | & \text{TypeName} . \text{Operator} \end{array}$$

If an operator has more than one arity (e.g., unary and binary $-$), the appropriate version may be selected by giving the formal parameter types. The binary version is selected by default. For example, the following equivalences hold:

<code>String.+</code>	\equiv	<code>(x: String, y: String): String => x + y</code>
<code>Long.-</code>	\equiv	<code>(x: Long, y: Long): Long => x - y</code>
<code>Float.-(Float,Float)</code>	\equiv	<code>(x: Float, y: Float): Float => x - y</code>
<code>Int.-(Int)</code>	\equiv	<code>(x: Int): Int => -x</code>
<code>Boolean.&</code>	\equiv	<code>(x: Boolean, y: Boolean): Boolean => x & y</code>
<code>Boolean.!</code>	\equiv	<code>(x: Boolean): Boolean => !x</code>
<code>Int.<(Int,Int)</code>	\equiv	<code>(x: Int, y: Int): Boolean => x < y</code>
<code>Dist. (Place)</code>	\equiv	<code>(d: Dist, p: Place): Dist => d p</code>

Unary and binary promotion (§12.9) is not performed when invoking these operations; instead, the operands are coerced individually via implicit coercions (§4.9), as appropriate.

Additionally, for every expression e of a type T at which a binary operator OP is defined, the expression $e.OP$ or $e.OP(T)$ represents the function defined by:

$$(x: T): T \Rightarrow \{ e \text{ OP } x \}$$

Primary ::= *Expr . Operator (Formals²)*
 | *Expr . Operator*

For example, one may write an expression that adds one to each member of a list `xs` by:

```
xs.map(1.+)
```

11.5 Functions as objects of type Any

Two functions `f` and `g` are equal (“==”) if both are instances of classes and the same object, or if both were obtained by the same evaluation of a function literal.¹ Further, it is guaranteed that if two functions are equal then they refer to the same locations in the environment and represent the “same code” (so their executions in an identical environment are comparable).

Every function type implements all the methods of `Any`. For a value `f` of a function type $(x_1:T_1, \dots, x_n:T_n) \Rightarrow T$, the expression `f.equals(g)` is of type `Boolean`. It succeeds if and only if `f==g` succeeds. Similarly, the expression `f.hashCode()` is of type `Int` and returns an implementation defined hash code which is guaranteed to be the same for two values that are equal. The expression `f.toString()` returns an implementation-dependent string. Two strings returned on different evaluations are equal to each other. Similarly, `f.typeName()` returns an implementation-dependent string. Two strings returned on different evaluations are equal to each other.

The method `f.home()` always returns `here`. The call `f.at(p)` for `p` a value of type `Place` always returns `true`. The call `f.at(o)` for `o` a value of type `Object` always returns `true`.

¹A literal may occur in program text within a loop, and hence may be evaluated multiple times.

12 Expressions

X10 supports a rich expression language similar to Java's. Evaluating an expression produces a value, which may be either an instance of a reference class, struct, or function. Expressions may also be `void`; that is, they produce no value. Expression evaluation may have side effects: assignment to a variable, allocation, method calls, or exceptional control-flow. Evaluation is strict and is performed left to right.

12.1 Literals

X10 supports the following literal expressions:

- An 32-bit integer literal is a value of type `x10.lang.Int`.
- An 64-bit long literal is a value of type `x10.lang.Long`.
- A 32-bit floating-point literal is a value of type `x10.lang.Float`.
- A 64-bit floating-point literal is a value of type `x10.lang.Double`.
- A character literal is a value of type `x10.lang.Char`.
- A string literal is a value of type `x10.lang.String`.
- The boolean literals `true` and `false` are of type `x10.lang.Boolean`.
- The `null` literal is of the null type, a subtype of all reference types.

12.2 this

```
ThisExpression ::= this
                  | ClassName . this
```

The expression **this** is a final local variable containing a reference to an instance of the lexically enclosing class. It may be used only within the body of an instance method, a constructor, or in the initializer of a instance field.

Within an inner class, **this** may be qualified with the name of a lexically enclosing class. In this case, it represents an instance of that enclosing class.

The type of a **this** expression is the innermost enclosing class, or the qualifying class, constrained by the class invariant and the method guard, if any.

The **this** expression may also be used within constraints in a class or interface header (the class invariant and **extends** and **implements** clauses). Here, the type of **this** is restricted so that only properties declared in the class header itself, and specifically not any members declared in the class body or in supertypes, are accessible through **this**.

12.3 Local variables

```
LocalExpression ::= Identifier
```

A local variable expression consists simply of the name of the local variable.

12.4 Field access

```
FieldExpression ::= Expression . Identifier
                   | super . Identifier
                   | ClassName . Identifier
                   | ClassName . super . Identifier
```

A field of an object instance may be accessed with a field access expression.

The type of the access is the declared type of the field with the actual target substituted for **this** in the type.

The field accessed is selected from the fields and value properties of the static type of the target and its superclasses.

If the field target is given by the keyword `super`, the target's type is the superclass of the enclosing class, as constrained by the superclass's class invariant, if any.

If the field target is `null`, a `NullPointerException` is thrown.

If the field target is a class name, a static field is selected.

It is illegal to access a field that is not visible from the current context. It is illegal to access a non-static field through a static field access expression.

12.5 Function Literals

Please see §11 for a description of `FunctionExpressions`.

12.6 Calls

$$\begin{aligned}
 \textit{MethodCall} & ::= \textit{TypeName} . \textit{Identifier} \textit{TypeArguments}^? (\textit{ArgumentList}^?) \\
 & \quad | \text{super} . \textit{Identifier} \textit{TypeArguments}^? (\textit{ArgumentList}^?) \\
 & \quad | \textit{ClassName} . \text{super} . \textit{Identifier} \textit{TypeArguments}^? (\textit{ArgumentList}^?) \\
 \textit{Call} & ::= \textit{Primary} \textit{TypeArguments}^? (\textit{ArgumentList}^?) \\
 \textit{TypeArguments} & ::= [\textit{Type} (, \textit{Type})^*]
 \end{aligned}$$

A *MethodCall* may be to either `static` or to instance methods. A *Call* may to either a method or a closure. The syntax is ambiguous; the target must be type-checked to determine if it is the name of a method or if it refers to a closure.

It is a static error if a call may resolve to both a closure call or to a method call.

A closure call `e(...)` is shorthand for a method call `e.apply(...)`.

Method selection rules are similar to that of Java. For a call with no explicit type arguments, a method with no parameters is considered more specific than a method with one or more type parameters that would have to be inferred.

Type arguments may be omitted and inferred, as described in §4.11.

It is a static error if a method's *Guard* is not satisfied by the caller.

12.7 Assignment

```

Expression ::= Assignment
Assignment ::= SimpleAssignment
                | OpAssignment
SimpleAssignment ::= LeftHandSide = Expression
OpAssignment ::= LeftHandSide += Expression
                ::= LeftHandSide -= Expression
                ::= LeftHandSide *= Expression
                ::= LeftHandSide /= Expression
                ::= LeftHandSide %= Expression
                ::= LeftHandSide &= Expression
                ::= LeftHandSide |= Expression
                ::= LeftHandSide ^= Expression
                ::= LeftHandSide <<= Expression
                ::= LeftHandSide >>= Expression
                ::= LeftHandSide >>>= Expression
LeftHandSide ::= Identifier
                | Primary . Identifier
                | Primary ( Expression )

```

The assignment expression `x = e` assigns a value given by expression `e` to a mutable variable `x`. There are three forms of assignment: `x` may be a local variable, it may be a field `y.f`, or it may be the variable obtained by evaluating the expression `a(i)`. In the last case, `a` must be an instance of a class implementing the interface `x10.lang.Settable[S,T]`, where `S` and `T` are the types of `i` and `e`, respectively.

This interface is defined thus:

```

package x10.lang;
public interface Settable[S,T] {
    def set[S,T](i: S, v: T): T;
}

```

The assignment `a(i) = e` is equivalent to the call `a.set(i, e)`.

For a binary operator `op`, the `op`-assignment expression `x op= e` evaluates `x` to a memory location, then evaluates `e`, applies the operation `x op e`, and assigns the result into the location computed for `x`. The expression is equivalent to `x = x op e` except that any subexpressions of `x` are evaluated only once.

12.8 Increment and decrement

The operators `++` and `--` increment and decrement a variable, respectively. The variable must be non-final and of numeric type.

When the operator is prefix, the variable is incremented or decremented by 1 and the result of the expression is the new value of the variable. When the operator is postfix, the variable is incremented or decremented by 1 and the result of the expression is the old value of the variable.

The new value of the variable v is identical to the result of the expressions $v+1$ or $v-1$, as appropriate.

12.9 Numeric promotion

The unary and binary operators promote their operands as follows. Values are sign extended and converted to instances of the promoted type.

- The unary promotion of `Byte`, `Short`, `Int` is `Int`.
- The unary promotion of `Long` is `Long`.
- The unary promotion of `Float` is `Float`.
- The unary promotion of `Double` is `Double`.
- The binary promotion of two types is the greater of the unary promotion of each type according to the following order: `Int`, `Long`, `Float`, `Double`.

12.10 Unary plus and unary minus

The unary `+` operator applies unary numeric promotion to its operand. The operand must be of numeric type.

The unary `-` operator applies unary numeric promotion to its operand and then subtracts the promoted operand from `0`. The operand must be of numeric type. The type of the result is promoted type.

12.11 Bitwise complement

The unary `~` operator applies unary numeric promotion to its operand and then evaluates to the bitwise complement of the promoted operand. The operand must be of integral type. The type of the result is promoted type.

12.12 Binary arithmetic operations

The binary arithmetic operations apply binary numeric promotion to their operands. The operands must be of numeric type. The type of the result is the promoted type. The `+` operator adds the promoted operands. The `-` operator subtracts the second operand from the first. The `*` operator multiplies the promoted operands. The `/` operator divides the first operand by the second. The `%` operator evaluates to the remainder of the division of the first operand by the second.

Floating point operations are determined by the IEEE 754 standard. The integer `/` and `%` throw a `DivideByZeroException` if the right operand is zero.

12.13 Binary shift operations

Unary promotion is performed on each operand separately. The operands must be of integral type. The type of the result is the promoted type of the left operand.

If the promoted type of the left operand is `Int`, the right operand is masked with `0x1f` using the bitwise AND (`&`) operator. If the promoted type of the left operand is `Long`, the right operand is masked with `0x3f` using the bitwise AND (`&`) operator.

The `<<` operator left-shifts the left operand by the number of bits given by the right operand.

The `>>` operator right-shifts the left operand by the number of bits given by the right operand. The result is sign extended; that is, if the right operand is k , the most significant k bits of the result are set to the most significant bit of the operand.

The `>>>` operator right-shifts the left operand by the number of bits given by the right operand. The result is not sign extended; that is, if the right operand is k , the most significant k bits of the result are set to `0`.

12.14 Binary bitwise operations

The binary bitwise operations apply binary numeric promotion to their operands. The operands must be of integral type. The type of the result is the promoted type. The `&` operator performs the bitwise AND of the promoted operands. The `|` operator performs the bitwise inclusive OR of the promoted operands. The `^` operator performs the bitwise exclusive OR of the promoted operands.

12.15 String concatenation

The `+` operator is used for string concatenation as well as addition. If either operand is of static type `x10.lang.String`, the other operand is converted to a `String`, if needed, and the two strings are concatenated.

String conversion of a non-null value is performed by invoking the `toString()` method of the value. If the value is `null`, the value is converted to `"null"`.

The type of the result is `String`.

12.16 Logical negation

The operand of the unary `!` operator must be of type `x10.lang.Boolean`. The type of the result is `Boolean`. If the value of the operand is `true`, the result is `false`; if if the value of the operand is `false`, the result is `true`.

12.17 Boolean logical operations

Operands of the binary boolean logical operators must be of type `Boolean`. The type of the result is `Boolean`

The `&` operator evaluates to `true` if both of its operands evaluate to `true`; otherwise, the operator evaluates to `false`.

The `|` operator evaluates to `false` if both of its operands evaluate to `false`; otherwise, the operator evaluates to `true`.

12.18 Boolean conditional operations

Operands of the binary boolean conditional operators must be of type `Boolean`. The type of the result is `Boolean`.

The `&&` operator evaluates to `true` if both of its operands evaluate to `true`; otherwise, the operator evaluates to `false`. Unlike the logical operator `&`, if the first operand is `false`, the second operand is not evaluated.

The `||` operator evaluates to `false` if both of its operands evaluate to `false`; otherwise, the operator evaluates to `true`. Unlike the logical operator `||`, if the first operand is `true`, the second operand is not evaluated.

12.19 Relational operations

The relational operations apply binary numeric promotion to their operands. The operands must be of numeric type. The type of the result is `Boolean`.

The `<` operator evaluates to `true` if the left operand is less than the right. The `<=` operator evaluates to `true` if the left operand is less than or equal to the right. The `>` operator evaluates to `true` if the left operand is greater than the right. The `>=` operator evaluates to `true` if the left operand is greater than or equal to the right.

Floating point comparison is determined by the IEEE 754 standard. Thus, if either operand is NaN, the result is `false`. Negative zero and positive zero are considered to be equal. All finite values are less than positive infinity and greater than negative infinity.

12.20 Conditional expressions

ConditionalExpression ::= Expression ? Expression : Expression

A conditional expression evaluates its first subexpression (the condition); if `true` the second subexpression (the consequent) is evaluated; otherwise, the third subexpression (the alternative) is evaluated.

The type of the condition must be `Boolean`. The type of the conditional expression is the least common ancestor (§4.8) of the types of the consequent and the alternative.

12.21 Stable equality

$$\begin{array}{lcl} \textit{EqualityExpression} & ::= & \textit{Expression} == \textit{Expression} \\ & | & \textit{Expression} != \textit{Expression} \end{array}$$

The `==` and `!=` operators provide *stable equality*; that is, the result of the equality operation is not affected by the mutable state of the program.

Two operands may be compared with the infix predicate `==`. The operation evaluates to `true` if and only if no action taken by any user program can distinguish between the two operands. In more detail, the rules are as follows.

If the operands both have reference type, then the operation evaluates to `true` if both are references to the same object.

If one operand evaluates to `null` then the predicate evaluates to `true` if and only if the other operand is also `null`.

If the operands both have struct type, then they must be structurally equal; that is, they must be instances of the same struct and all their fields or components must be `==`.

The definition of equality for function types is specified in §11.5.

If the operands both have numeric type, binary promotion (§12.9) is performed on the operands before the comparison.

The predicate `!=` returns `true` (`false`) on two arguments if and only if the operand `==` returns `false` (`true`) on the same operands.

The predicates `==` and `!=` may not be overridden by the programmer.

12.22 Allocation

$$\begin{array}{lcl} \textit{NewExpression} & ::= & \textit{new } \textit{ClassName } \textit{TypeArguments}^? (\textit{ArgumentList}^?) \textit{ClassBody}^? \\ & | & \textit{new } \textit{InterfaceName } \textit{TypeArguments}^? (\textit{ArgumentList}^?) \textit{ClassBody} \end{array}$$

An allocation expression creates a new instance of a class and invokes a constructor of the class. The expression designates the class name and passes type and value arguments to the constructor.

The allocation expression may have an optional class body. In this case, an anonymous subclass of the given class is allocated. An anonymous class allocation may

also specify a single super-interface rather than a superclass; the superclass of the anonymous class is `x10.lang.Object`.

If the class is anonymous—that is, if a class body is provided—then the constructor is selected from the superclass. The constructor to invoke is selected using the same rules as for method invocation (§12.6).

The type of an allocation expression is the return type of the constructor invoked, with appropriate substitutions of actual arguments for formal parameters, as specified in §12.6.

It is illegal to allocate an instance of an `abstract` class. It is illegal to allocate an instance of a class or to invoke a constructor that is not visible at the allocation expression.

12.23 Casts

The cast operation may be used to cast an expression to a given type:

```
UnaryExpression ::= CastExpression
CastExpression ::= UnaryExpression as Type
```

The result of this operation is a value of the given type if the cast is permissible at run time.

(B: This section need serious whomping. The Java mention needs to go. The rules for coercions are given in §??. If the Type has a constraint, the constraint will be checked at runtime. We need to give examples. :B)

Type conversion is checked according to the rules of the Java language (e.g., [5, §5.5]). For constrained types, both the base type and the constraint are checked. If the value cannot be cast to the appropriate type, a `ClassCastException` is thrown.

Conversions of numeric values Can't do (a as T) if a can't be a T.

12.24 instanceof

X10 permits types to be used in an `instanceof` expression to determine whether an object is an instance of the given type:

RelationalExpression ::= *RelationalExpression instanceof Type*

In the above expression, *Type* is any type. At run time, the result of this operator is `true` if the *RelationalExpression* can be coerced to *Type* without a `ClassCastException` being thrown. Otherwise the result is `false`. This determination may involve checking that the constraint, if any, associated with the type is true for the given expression.

12.25 Subtyping expressions

SubtypingExpression ::= *Expression <: Expression*
 | *Expression :> Expression*
 | *Expression == Expression*

The subtyping expression $T_1 <: T_2$ evaluates to `true` T_1 is a subtype of T_2 .

The expression $T_1 :> T_2$ evaluates to `true` T_2 is a subtype of T_1 .

The expression $T_1 == T_2$ evaluates to `true` T_1 is a subtype of T_2 and if T_2 is a subtype of T_1 .

Subtyping expressions are used in subtyping constraints for generic types.

12.26 Contains expressions

ContainsExpression ::= *Expression in Expression*

The expression $p \text{ in } r$ tests if a value p is in a collection r ; it evaluates to $r.\text{contains}(p)$. The collection r must be of type `Collection[T]` and the value p must be of type T .

12.27 Rail constructors

RailConstructor ::= *[Expressions]*
Expressions ::= *Expression (, Expression)**

The rail constructor $[a_0, \dots, a_{k-1}]$ creates an instance of `ValRail` with length k where the i th element is a_i . The element type of the array (T) is bound to the least common ancestor of the types of the a_i (§4.8).

Since arrays are subtypes of `(Point) => T`, rail constructors can be passed into the `Array` and `ValArray` constructors as initializer functions.

Rail constructors of type `ValRail[Int]` and length n may be implicitly converted to type `Point{rank==n}`. Rail constructors of type `ValRail[Region]` and length n may be implicitly converted to type `Region{rank==n}`.

13 Statements

This chapter describes the statements in the sequential core of X10. Statements involving concurrency and distribution are described in §15.

13.1 Empty statement

Statement ::= ;

The empty statement ; does nothing. It is useful when a loop header is evaluated for its side effects. For example, the following code sums the elements of an array.

```
var sum: Int = 0;
for (var i: Int = 0; i < a.length; i++, sum += a[i])
    ;
```

13.2 Local variable declaration

Statement ::= *LocalVariableDeclarationStatement*
LocalVariableDeclarationStatement ::= *LocalVariableDeclaration* ;

The syntax of local variables declarations is described in §5.

Local variables may be declared only within a block statement (§13.3). The scope of a local variable declaration is the statement itself and the subsequent statements in the block.

13.3 Block statement

Statement ::= *BlockStatement*
BlockStatement ::= { *Statement** }

A block statement consists of a sequence of statements delimited by “{” and “}”. Statements are evaluated in order. The scope of local variables introduced within the block is the remainder of the block following the variable declaration.

13.4 Expression statement

Statement ::= *ExpressionStatement*
ExpressionStatement ::= *StatementExpression* ;
StatementExpression ::= *Assignment*
 | *Allocation*
 | *Call*

The expression statement evaluates an expression, ignoring the result. The expression must be either an assignment, an allocation, or a call.

13.5 Labeled statement

Statement ::= *LabeledStatement*
LabeledStatement ::= *Identifier* : *Statement*

Statements may be labeled. The label may be used as the target of a break or continue statement. The scope of a label is the statement labeled.

13.6 Break statement

Statement ::= *BreakStatement*
BreakStatement ::= break *Identifier*[?]

An unlabeled break statement exits the currently enclosing loop or switch statement.

An labeled break statement exits the enclosing loop or switch statement with the given label.

It is illegal to break out of a loop not defined in the current method, constructor, initializer, or closure.

The following code searches for an element of a two-dimensional array and breaks out of the loop when found:

```
var found: Boolean = false;
outer: for (var i: Int = 0; i < a.length; i++)
    for (var j: Int = 0; j < a(i).length; j++)
        if (a(i)(j) == v) {
            found = true;
            break outer;
        }
```

13.7 Continue statement

Statement ::= *ContinueStatement*
ContinueStatement ::= `continue` *Identifier*[?]

An unlabeled continue statement branches to the top of the currently enclosing loop.

An labeled break statement branches to the top of the enclosing loop with the given label.

It is illegal to continue a loop not defined in the current method, constructor, initializer, or closure.

13.8 If statement

```

Statement ::= IfThenStatement
           | IfThenElseStatement
IfThenStatement ::= if ( Expression ) Statement
IfThenElseStatement ::= if ( Expression ) Statement else Statement

```

An if statement comes in two forms: with and without an else clause.

The if-then statement evaluates a condition expression and evaluates the consequent expression if the condition is `true`. If the condition is `false`, the if-then statement completes normally.

The if-then-else statement evaluates a condition expression and evaluates the consequent expression if the condition is `true`; otherwise, the alternative statement is evaluated.

The condition must be of type `Boolean`.

13.9 Switch statement

```

Statement ::= SwitchStatement
SwitchStatement ::= switch ( Expression ) { Case+ }
Case ::= case Expression : Statement*
      | default : Statement*

```

A switch statement evaluates an index expression and then branches to a case whose value equal to the value of the index expression. If no such case exists, the switch branches to the `default` case, if any.

Statements in each case branch evaluated in sequence. At the end of the branch, normal control-flow falls through to the next case, if any. To prevent fall-through, a case branch may be exited using a `break` statement.

The index expression must be of type `Int`.

Case labels must be of type `Int` and must be compile-time constants. Case labels cannot be duplicated within the `switch` statement.

13.10 While statement

Statement ::= *WhileStatement*
WhileStatement ::= **while** (*Expression*) *Statement*

A while statement evaluates a condition and executes a loop body if **true**. If the loop body completes normally (either by reaching the end or via a **continue** statement with the loop header as target), the condition is reevaluated and the loop repeats if **true**. If the condition is **false**, the loop exits.

The condition must be of type **Boolean**.

13.11 Do-while statement

Statement ::= *DoWhileStatement*
DoWhileStatement ::= **do** *Statement* **while** (*Expression*) ;

A do-while statement executes a loop body, and then evaluates a condition expression. If **true**, the loop repeats. Otherwise, the loop exits.

The condition must be of type **Boolean**.

13.12 For statement

Statement ::= *ForStatement*
 | *EnhancedForStatement*
ForStatement ::= **for** (*ForInit*[?] ; *Expression*[?] ; *ForUpdate*[?]) *Statement*
 ForInit ::= *StatementExpression* (, *StatementExpression*)^{*}
 | *LocalVariableDeclaration*
EnhancedForStatement ::= **for** (*Formal in Expression*) *Statement*

X10 provides two forms of for statement: a basic for statement and an enhanced for statement.

A basic for statement consists of an initializer, a condition, an iterator, and a body. First, the initializer is evaluated. The initializer may introduce local variables that are in scope throughout the for statement. An empty initializer is permitted.

Next, the condition is evaluated. If `true`, the loop body is executed; otherwise, the loop exits. The condition may be omitted, in which case the condition is considered `true`. If the loop completes normally (either by reaching the end or via a `continue` statement with the loop header as target), the iterator is evaluated and then the condition is reevaluated and the loop repeats if `true`. If the condition is `false`, the loop exits.

The condition must be of type `Boolean`. The initializer and iterator are statements, not expressions and so do not have types.

An enhanced `for` statement is used to iterate over a collection. If the formal parameter is of type `T`, the collection expression must be of type `Iterable[T]`. Destructuring syntax may be used for the formal parameter (§5). Each iteration of the loop binds the parameter to another element of the collection. The formal parameter must be immutable.

In a common case, the collection is intended to be of type `Region` and the formal parameter is of type `Point`. Expressions `e` of type `Dist` and `Array` are also accepted, and treated as if they were `e.region`. If the collection is a region, the `for` statement enumerates the points in the region in canonical order.

13.13 Throw statement

Statement ::= *ThrowStatement*
ThrowStatement ::= `throw Expression` ;

The `throw` statement throws an exception. The exception must be a subclass of the class `x10.lang.Throwable`.

Example 13.13.1 The following statement checks if an index is in range and throws an exception if not.

```
if (i < 0 || i > x.length)
    throw new IndexOutOfBoundsException();
```

□

13.14 Try-catch statement

```

Statement ::= TryStatement
TryStatement ::= try BlockStatement Catch+ Finally?
                | try BlockStatement Catch* Finally
  Catch ::= catch ( Formal ) BlockStatement
  Finally ::= finally BlockStatement

```

Exceptions are handled with a **try** statement. A **try** statement consists of a **try** block, zero or more **catch** blocks, and an optional **finally** block.

First, the **try** block is evaluated. If the block throws an exception, control transfers to the first matching **catch** block, if any. A **catch** matches if the value of the exception thrown is a subclass of the **catch** block's formal parameter type.

The **finally** block, if present, is evaluated on all normal and exceptional control-flow paths from the **try** block. If the **try** block completes normally or via a **return**, a **break**, or a **continue** statement, the **finally** block is evaluated, and then control resumes at the statement following the **try** statement, at the branch target, or at the caller as appropriate. If the **try** block completes exceptionally, the **finally** block is evaluated after the matching **catch** block, if any, and then the exception is rethrown.

13.15 Return statement

```

Statement ::= ReturnStatement
ReturnStatement ::= return Expression ;
                   | return ;

```

Methods and closures may return values using a **return** statement. If the method's return type is explicitly declared **Void**, the method must return without a value; otherwise, it must return a value of the appropriate type.

14 Places

An X10 place is a repository for data and activities. Each place is to be thought of as a locality boundary: the activities running in a place may access data items located at that place with the efficiency of on-chip access. Accesses to remote places may take orders of magnitude longer.

X10 provides a built-in struct, `x10.lang.Place`; all places are instances of this struct.

In X10 v2.0, the set of places available to a computation is determined at the time that the program is run and remains fixed through the run of the program. The number of places available may be determined by reading (`Place.MAX_PLACES`). (This number is specified from the command line/configuration information; see associated README documentation.)

All scalar objects created during program execution are located in one place, though they may be referenced from other places. Aggregate objects (arrays) may be distributed across multiple places using distributions.

The set of all places in a running instance of an X10 program may be obtained through the `const` field `Place.places`. (This set may be used to define distributions, for instance, §17.3.)

The set of all places is totally ordered. The first place may be obtained by reading `Place.FIRST_PLACE`. The initial activity for an X10 computation starts in this place (§15.5). For any place, the operation `next()` returns the next place in the total order (wrapping around at the end). Further details on the methods and fields available on this class may be obtained by consulting the API documentation.

NOTE: Future versions of the language may permit user-definable places, and the ability to dynamically create places.

STATIC SEMANTICS RULE: Variables of type `Place` must be initialized and are implicitly `final`.

14.1 Place expressions

Any expression of type `Place` is called a place expression. Examples of place expressions are `this.home` (the place at which the current object lives), `here` (the place where the current activity is executing), etc.

Place expressions are used in the following contexts:

- As a target for an `async` activity or a future (§15.2).
- In a cast expression (§12.23).
- In an `instanceof` expression (§12.24).
- In stable equality comparisons, at type `Place`.

Like values of any other type, places may be passed as arguments to methods, returned from methods, stored in fields etc.

14.2 here

X10 supports a special indexical constant¹ `here`:

ExpressionName ::= `here`

The constant evaluates to the place at which the current activity is running. Unlike other place expressions, this constant cannot be used as the `placetype` of fields, since the type of a field should be independent of the activity accessing it.

Example 14.2.1 The code:

```
public class F {
  public def m(a: F) {
    val OldHere: Place = here;
    async (a) {
      Console.OUT.println("OldHere == here:"
                          + (OldHere == here));
    }
  }
}
```

¹ An indexical constant is one whose value depends on its context of use.

```
    }  
  }  
  public static def main(Rail[String]) {  
    new F().m( (at(Place.FIRST_PLACE.next()) new F()));  
  }  
}
```

will print out `true` iff the computation was configured to start with the number of places set to 1. □

15 Activities

An X10 computation may have many concurrent *activities* “in flight” at any give time. We use the term activity to denote a dynamic execution instance of a piece of code (with references to data). An activity is intended to execute in parallel with other activities. An activity may be thought of as a very light-weight thread. In X10 v2.0, an activity may not be interrupted, suspended or resumed as the result of actions taken by any other activity.

An activity is spawned in a given place and stays in that place for its lifetime. An activity may be *running*, *blocked* on some condition or *terminated*. When the statement associated with an activity terminates normally, the activity terminates normally; when it terminates abruptly with some reason *R*, the activity terminates with the same reason (§15.1).

An activity may be long-running and may invoke recursive methods (thus may have a stack associated with it). On the other hand, an activity may be short-running, involving a fine-grained operation such as a single read or write.

An activity may asynchronously and in parallel launch activities at other places.

X10 distinguishes between *local* termination and *global* termination of a statement. The execution of a statement by an activity is said to terminate locally when the activity has finished all its computation related to that statement. (For instance the creation of an asynchronous activity terminates locally when the activity has been created.) It is said to terminate globally when it has terminated locally and all activities that it may have spawned at any place (if any) have, recursively, terminated globally.

An X10 computation is initiated as a single activity from the command line. This activity is the *root activity* for the entire computation. The entire computation terminates when (and only when) this activity globally terminates. Thus X10 does not permit the creation of so called “daemon threads”—threads that outlive the lifetime of the root activity. We say that an X10 computation is *rooted* (§15.5).

Future Extensions. *We may permit the initial activity to be a daemon activity to permit reactive computations, such as web servers, that may not terminate.*

15.1 The X10 rooted exception model

The rooted nature of X10 computations permits the definition of a *rooted* exception model. In multi-threaded programming languages there is a natural parent-child relationship between a thread and a thread that it spawns. Typically the parent thread continues execution in parallel with the child thread. Therefore the parent thread cannot serve to catch any exceptions thrown by the child thread.

The presence of a root activity permits X10 to adopt a different model. In any state of the computation, say that an activity *A* is a *root of* an activity *B* if *A* is an ancestor of *B* and *A* is suspended at a statement (such as the `finish` statement §15.4) awaiting the termination of *B* (and possibly other activities). For every X10 computation, the *root-of* relation is guaranteed to be a tree. The root of the tree is the root activity of the entire computation. If *A* is the nearest root of *B*, the path from *A* to *B* is called the *activation path* for the activity.¹

We may now state the exception model for X10. An uncaught exception propagates up the activation path to its nearest root activity, where it may be handled locally or propagated up the *root-of* tree when the activity terminates (based on the semantics of the statement being executed by the activity).² Thus, unlike concurrent languages such as Java, no exception is “thrown on the floor”.

15.2 Spawning an activity

Asynchronous activities serve as a single abstraction for supporting a wide range of concurrency constructs such as message passing, threads, DMA, streaming, data prefetching. (In general, asynchronous operations are better suited for supporting scalability than synchronous operations.)

An activity is created by executing the statement:

¹Note that depending on the state of the computation the activation path may traverse activities that are running, suspended or terminated.

²In X10 v2.0 the `finish` statement is the only statement that marks its activity as a root activity. Future versions of the language may introduce more such statements.

$$\begin{aligned}
\textit{Statement} & ::= \textit{AsyncStatement} \\
\textit{AsyncStatement} & ::= \textbf{async } \textit{PlaceExpressionSingleList}^? \textit{Statement} \\
\textit{PlaceExpressionSingleList} & ::= (\textit{PlaceExpression}) \\
\textit{PlaceExpression} & ::= \textit{Expression}
\end{aligned}$$

The place expression e is expected to be of type `Place`, e.g., `here` or `d(p)` for some distribution d and point p (§14). If not, the compiler replaces e with `e.home` if e is of type `x10.lang.Object`. Otherwise the compiler reports a type error.

Note specifically that the expression `a(i)` when used as a place expression may evaluate to `a(i).home`, which may not be the same place as `a.dist(i)`. The programmer must be careful to choose the right expression, appropriate for the statement. Accesses to `a(i)` within *Statement* should typically be guarded by the place expression `a.dist(i)`.

In many cases the compiler may infer the unique place at which the statement is to be executed by an analysis of the types of the variables occurring in the statement. (The place must be such that the statement can be executed safely, without generating a `BadPlaceException`.) In such cases the programmer may omit the place designator; the compiler will throw an error if it cannot determine the unique designated place.³

An activity A executes the statement `async (P) S` by launching a new activity B at the designated place, to execute the specified statement. The statement terminates locally as soon as B is launched. The activation path for B is that of A , augmented with information about the line number at which B was spawned. B terminates normally when S terminates normally. It terminates abruptly if S throws an (uncaught) exception. The exception is propagated to A if A is a root activity (see §15.4), otherwise through A to A 's root activity. Note that while an activity is running, exceptions thrown by activities it has already generated may propagate through it up to its root activity.

Multiple activities launched by a single activity at another place are not ordered in any way. They are added to the pool of activities at the target place and will be executed in sequence or in parallel based on the local scheduler's decisions. If the programmer wishes to sequence their execution s/he must use X10 constructs, such as `clocks` and `finish` to obtain the desired effect. Further, the X10 implementations are not required to have fair schedulers, though every implementation should make a best faith effort to ensure that every activity eventually gets a chance to make forward progress.

³X10 v2.0 does not specify a particular algorithm; this will be fixed in future versions.

STATIC SEMANTICS RULE: The statement in the body of an `async` is subject to the restriction that it must be acceptable as the body of a `void` method for an anonymous inner class declared at that point in the code, which throws no checked exceptions. As such, it may reference variables in lexically enclosing scopes (including `clock` variables, §16) provided that such variables are (implicitly or explicitly) `val`.

15.3 Place changes

An activity may change place using the `at` statement or `at` expression:

```

Statement ::= AtStatement
AtStatement ::= at PlaceExpressionSingleList Statement
Expression ::= AtExpression
AtExpression ::= at PlaceExpressionSingleList ClosureBody

```

The statement `at (p) S` executes the statement `S` synchronously at place `p`. The expression `at (p) E` executes the statement `E` synchronously at place `p`, returning the result to the originating place.

15.4 Finish

The statement `finish S` converts global termination to local termination and introduces a root activity.

```

Statement ::= FinishStatement
FinishStatement ::= finish Statement

```

An activity `A` executes `finish S` by executing `S`. The execution of `S` may spawn other asynchronous activities (here or at other places). Uncaught exceptions thrown or propagated by any activity spawned by `S` are accumulated at `finish S`. `finish S` terminates locally when all activities spawned by `S` terminate globally (either abruptly or normally). If `S` terminates normally, then `finish S` terminates normally and `A` continues execution with the next statement after `finish S`. If `S` terminates abruptly, then `finish S` terminates abruptly and throws a single exception, `x10.lang.MultipleExceptions` formed from the collection of exceptions accumulated at `finish S`.

Thus a `finish S` statement serves as a collection point for uncaught exceptions generated during the execution of `S`.

Note that repeatedly finishing a statement has no effect after the first `finish`: `finish finish S` is indistinguishable from `finish S`.

Interaction with clocks. `finish S` interacts with clocks (§16).

While executing `S`, an activity must not spawn any `clocked` `asyncs`. (`Asyncs` spawned during the execution of `S` may spawn `clocked` `asyncs`.) A `ClockUseException` is thrown if (and when) this condition is violated.

In X10 v2.0 this condition is checked dynamically; future versions of the language will introduce type qualifiers which permit this condition to be checked statically.

Future Extensions. *The semantics of `finish S` is conjunctive; it terminates when all the activities created during the execution of `S` (recursively) terminate. In many situations (e.g., nondeterministic search) it is natural to require a statement to terminate when any one of the activities it has spawned succeeds. The other activities may then be safely aborted. Future versions of the language may introduce a `finishone S` construct to support such speculative or nondeterministic computation.*

15.5 Initial activity

An X10 computation is initiated from the command line on the presentation of a classname `C`. The class must have a `public static def main(a: array[String])` method, otherwise an exception is thrown and the computation terminates. The single statement

```
finish async (Place.FIRST_PLACE) {
    C.main(s);
}
```

is executed where `s` is an array of strings created from command line arguments. This single activity is the root activity for the entire computation. (See §14 for a discussion of places.)

15.6 Foreach statements

+

Statement ::= *ForEachStatement*

ForEachStatement ::= **foreach** (*Formal in Expression*) *Statement*

The **foreach** statement is similar to the enhanced **for** statement (§13.12).

An activity executes a **foreach** statement in a similar fashion except that separate **async** activities are launched in parallel in the local place of each object returned by the iteration. The statement terminates locally when all the activities have been spawned. It never throws an exception, though exceptions thrown by the spawned activities are propagated through to the root activity.

In a common case, the the collection is intended to be of type **Region** and the formal parameter is of type **Point**. Expressions *e* of type **Dist** and **Array** are also accepted, and treated as if they were *e.region*.

15.7 Ateach statements

Statement ::= *AtEachStatement*

AtEachStatement ::= **ateach** (*Formal in Expression*) *Statement*

The **ateach** statement is similar to the **foreach** statement. The collection must be of type **Dist** and the formal parameter of type **Point**. Expressions *e* of type **Array** are also accepted, and treated as if they were *e.dist*. The compiler reports a type error in all other cases.

This statement differs from **foreach** only in that each activity is spawned at the place specified by the distribution for the point. That is, **ateach**(*p*(*i*₁, ..., *i*_k): **Point** in *D*) *S* may be thought of as standing for:

```
foreach (p(i1,...,ik): point in D.region)
  async (D(p)) S
```

15.8 Futures

X10 provides syntactic support for *asynchronous expressions*, also known as futures:

Primary ::= *FutureExpression*
FutureExpression ::= `future` *PlaceExpressionSingleList*² *ClosureBody*

Intuitively such an expression evaluates its body asynchronously at the given place. The resulting value may be obtained from the future returned by this expression, by using the `force` operation.

In more detail, in an expression `future (Q) e`, the place expression *Q* is treated as in an `async` statement. *e* is an expression of some type *T*. *e* may reference only those variables in the enclosing lexical environment which are declared to be `final`.

If the type of *e* is *T* then the type of `future (Q) e` is `Future[T]`. This type `Future[T]` is defined as if by:

```
package x10.lang;
public interface Future[T] implements () => T {
    global def forced(): Boolean;
    global def force(): T;
}
```

Evaluation of `future (Q) e` terminates locally with the creation of a value *f* of type `Future[T]`. This value may be stored in objects, passed as arguments to methods, returned from method invocation etc.

At any point, the method `forced` may be invoked on *f*. This method returns without blocking, with the value `true` if the asynchronous evaluation of *e* has terminated globally and with the value `false` if it has not.

`Future[T]` is a subtype of the function type `() => T`. Invoking—*forcing*—the future *f* blocks until the asynchronous evaluation of *e* has terminated globally. If the evaluation terminates successfully with value *v*, then the method invocation returns *v*. If the evaluation terminates abruptly with exception *z*, then the method throws exception *z*. Multiple invocations of the function (by this or any other activity) do not result in multiple evaluations of *e*. The results of the first evaluation are stored in the future *f* and used to respond to all queries.

```
promise: Future[T] = future (a.dist(3)) a(3);
value: T = promise();
```

15.9 At expressions

Expression ::= **at** (*Expression*) *Expression*

An **at** expression evaluates an expression synchronously at the given place and returns its value. Note that expression evaluation may spawn asynchronous activities. The **at** expression will return without waiting for those activities to terminate. That is, **at** does not have built-in **finish** semantics.

15.10 Shared variables

Compiler Limitation: Shared variables are not currently implemented.

A *shared local variable* is declared with the annotation **shared**. It can be accessed within any control construct in its scope, including **async**, **at**, **future** and closures.

Note that the lifetime of some of these constructs may outlast the lifetime of the scope – requiring the implementation to allocate them outside the current stack frame.

15.11 Atomic blocks

Languages such as Java use low-level synchronization locks to allow multiple interacting threads to coordinate the mutation of shared data. X10 eschews locks in favor of a very simple high-level construct, the *atomic block*.

A programmer may use atomic blocks to guarantee that invariants of shared data-structures are maintained even as they are being accessed simultaneously by multiple activities running in the same place.

15.11.1 Unconditional atomic blocks

The simplest form of an atomic block is the *unconditional atomic block*:

Statement ::= **AtomicStatement**
AtomicStatement ::= **atomic** *Statement*
MethodModifier ::= **atomic**

For the sake of efficient implementation X10 v2.0 requires that the atomic block be *analyzable*, that is, the set of locations that are read and written by the *Block-Statement* are bounded and determined statically.⁴ The exact algorithm to be used by the compiler to perform this analysis will be specified in future versions of the language.

Such a statement is executed by an activity as if in a single step during which all other concurrent activities in the same place are suspended. If execution of the statement may throw an exception, it is the programmer's responsibility to wrap the atomic block within a `try/finally` clause and include undo code in the finally clause. Thus the `atomic` statement only guarantees atomicity on successful execution, not on a faulty execution.

We allow methods of an object to be annotated with `atomic`. Such a method is taken to stand for a method whose body is wrapped within an `atomic` statement.

Atomic blocks are closely related to non-blocking synchronization constructs [6], and can be used to implement non-blocking concurrent algorithms.

STATIC SEMANTICS RULE: In `atomic S`, `S` may include method calls, conditionals, etc.

It may *not* include an `async` activity (such as creation of a `future`).

It may *not* include any statement that may potentially block at runtime (e.g., `when`, `force` operations, `next` operations on clocks, `finish`).

All locations accessed in an atomic block must statically satisfy the *locality condition*: they must belong to the place of the current activity.

The compiler checks for this condition by checking whether the statement could be the body of a `void` method annotated with `safe` at that point in the code (§9.6.3).

Consequences. Note an important property of an (unconditional) atomic block:

$$\text{atomic } \{s1; \text{atomic } s2\} = \text{atomic } \{s1; s2\} \quad (15.1)$$

Further, an atomic block will eventually terminate successfully or thrown an exception; it may not introduce a deadlock.

⁴A static bound is a constant that depends only on the program text, and is independent of any runtime parameters.

Example

The following class method implements a (generic) compare and swap (CAS) operation:

```
// target defined in lexically enclosing environment.
public atomic def CAS(old: Object, new: Object): Boolean {
  if (target.equals(old)) {
    target = new;
    return true;
  }
  return false;
}
```

15.11.2 Conditional atomic blocks

Conditional atomic blocks are of the form:

$$\begin{aligned} \textit{Statement} & ::= \textit{WhenStatement} \\ \textit{WhenStatement} & ::= \textit{when} (\textit{Expression}) \textit{Statement} \\ & \quad | \quad \textit{WhenStatement} \textit{ or } (\textit{Expression}) \textit{Statement} \end{aligned}$$

In such a statement the one or more expressions are called *guards* and must be Boolean expressions. The statements are the corresponding *guarded statements*. The first pair of expression and statement is called the *main clause* and the additional pairs are called *auxiliary clauses*. A statement must have a main clause and may have no auxiliary clauses.

An activity executing such a statement suspends until such time as any one of the guards is true in the current state. In that state, the statement corresponding to the first guard that is true is executed. The checking of the guards and the execution of the corresponding guarded statement is done atomically.

X10 does not guarantee that a conditional atomic block will execute if its condition holds only intermittently. For, based on the vagaries of the scheduler, the precise instant at which a condition holds may be missed. Therefore the programmer is advised to ensure that conditions being tested by conditional atomic blocks are eventually stable, i.e., they will continue to hold until the block executes (the action in the body of the block may cause the condition to not hold any more).

RATIONALE: The guarantee provided by `wait/notify` in Java is no stronger. Indeed conditional atomic blocks may be thought of as a replacement for Java's `wait/notify` functionality.

We note two common abbreviations. The statement `when (true) S` is behaviorally identical to `atomic S`: it never suspends. Second, `when (c) {;}` may be abbreviated to `await(c);`—it simply indicates that the thread must await the occurrence of a certain condition before proceeding. Finally note that a `when` statement with multiple branches is behaviorally identical to a `when` statement with a single branch that checks the disjunction of the condition of each branch, and whose body contains an `if/then/else` checking each of the branch conditions.

STATIC SEMANTICS RULE: For the sake of efficient implementation certain restrictions are placed on the guards and statements in a conditional atomic block.

Guards are required not to have side-effects, not to spawn asynchronous activities and to have a statically determinable upper bound on their execution. These conditions are expected to be checked statically by the compiler.

The body of a `when` statement must satisfy the conditions for the body of an `atomic` block.

Note that this implies that guarded statements are required to be *flat*, that is, they may not contain conditional atomic blocks. (The implementation of nested conditional atomic blocks may require sophisticated operational techniques such as rollbacks.)

Sample usage. There are many ways to ensure that a guard is eventually stable. Typically the set of activities are divided into those that may enable a condition and those that are blocked on the condition. Then it is sufficient to require that the threads that may enable a condition do not disable it once it is enabled. Instead the condition may be disabled in a guarded statement guarded by the condition. This will ensure forward progress, given the weak-fairness guarantee.

Example 15.11.1 The following class shows how to implement a bounded buffer of size 1 in X10 for repeated communication between a sender and a receiver.

```
class OneBuffer {
  datum: Object = null;
  filled: Boolean = false;
  public def send(v: Object) {
```



```
        when (!filled) {
            this.datum = v;
            this.filled = true;
        }
    }
    public def receive(): Object {
        when (filled) {
            v: Object = datum;
            datum = null;
            filled = false;
            return v;
        }
    }
}
```

□

16 Clocks

The standard library for X10, `x10.lang`, defines a class `Clock` intended for repeated quiescence detection of arbitrary, data-dependent collection of activities. Clocks are a generalization of *barriers*. They permit dynamically created activities to register and deregister. An activity may be registered with multiple clocks at the same time. In particular, nested clocks are permitted: an activity may create a nested clock and within one phase of the outer clock schedule activities to run to completion on the nested clock. Nevertheless the design of clocks ensures that deadlock cannot be introduced by using clock operations, and that clock operations do not introduce any races.

This chapter describes the syntax and semantics of clocks and statements in the language that have parameters of type `Clock`.

The key invariants associated with clocks are as follows. At any stage of the computation, a clock has zero or more *registered* activities. An activity may perform operations only on those clocks it is registered with (these clocks constitute its *clock set*). An activity is registered with one or more clocks when it is created. During its lifetime the only additional clocks it is registered with are exactly those that it creates. In particular it is not possible for an activity to register itself with a clock it discovers by reading a data-structure.

An activity may perform the following operations on a clock `c`. It may *unregister* with `c` by executing `c.drop()`; . After this, it may perform no further actions on `c` for its lifetime. It may *check* to see if it is unregistered on a clock. It may *register* a newly forked activity with `c`. Once registered and "active" (see below), it may also perform the following operations. It may *resume* the clock by executing `c.resume()`; . This indicates to `c` that it has finished posting all statements it wishes to perform in the current phase. Finally, it may *block* (by executing `next;`) on all the clocks that it is registered with. (This operation implicitly *resume*'s all clocks for the activity.) It will resume from this statement only when all these

clocks are ready to advance to the next phase.

A clock becomes ready to advance to the next phase when every activity registered with the clock has executed at least one `resume` operation on that clock and all statements posted for completion in the current phase have been completed.

Though clocks introduce a blocking statement (`next`) an important property of X10 is that clocks cannot introduce deadlocks. That is, the system cannot reach a quiescent state (in which no activity is progressing) from which it is unable to progress. For, before blocking each activity resumes all clocks it is registered with. Thus if a configuration were to be stuck (that is, no activity can progress) all clocks will have been resumed. But this implies that all activities blocked on `next` may continue and the configuration is not stuck. The only other possibility is that an activity may be stuck on `finish`. But the interaction rule between `finish` and clocks (§15.4) guarantees that this cannot cause a cycle in the wait-for graph. A more rigorous proof may be found in [9].

16.1 Clock operations

The special statements introduced for clock operations are listed below.

```

Statement ::= ClockedStatement
ClockedStatement ::= clocked ( ClockList ) Statement
NextStatement ::= next ;

```

Note that `x10.lang.Clock` provides several useful methods on clocks (e.g. `drop`).

16.1.1 Creating new clocks

Clocks are created using a factory method on `x10.lang.Clock`:

```
timeSynchronizer: Clock = Clock.make();
```

The current activity is automatically registered with the newly created clock. It may deregister using the `drop` method on clocks (see the documentation of `x10.lang.Clock`). All activities are automatically deregistered from all clocks they are registered with on termination (normal or abrupt).

16.1.2 Registering new activities on clocks

The programmer may specify which clocks a new activity is to be registered with using the `clocked` clause.

An activity may transmit only those clocks that is registered with and has not quiesced on (§16.1.3). A `ClockUseException` is thrown if (and when) this condition is violated.

An activity may check that it is registered on a clock `c` by executing:

```
c.registered()
```

This call returns the Boolean value `true` iff the activity is registered on `c`; otherwise it returns `false`.

NOTE: X10 does not contain a “register” statement that would allow an activity to discover a clock in a datastructure and register itself on it. Therefore, while clocks may be stored in a datastructure by one activity and read from that by another, the new activity cannot “use” the clock unless it is already registered with it.

16.1.3 Resuming clocks

X10 permits *split phase* clocks. An activity may wish to indicate that it has completed whatever work it wishes to perform in the current phase of a clock `c` it is registered with, without suspending all activity. It may do so by executing the method invocation:

```
c.resume();
```

An activity may invoke this method only on a clock it is registered with, and has not yet dropped (§16.1.5). A `ClockUseException` is thrown if (and when) this condition is violated. Nothing happens if the activity has already invoked a `resume` on this clock in the current phase. Otherwise execution of this statement indicates that the activity will not transmit `c` to an `async` (through a `clocked` clause), until it terminates, drops `c` or executes a `next`.

STATIC SEMANTICS RULE: The compiler should issue an error if any activity has a potentially live execution path from a `resume` statement on a clock `c` to a `async` spawn statement (which registers the new activity on `c`) unless the path goes through a `next` statement. (A `c.drop()` following a `c.resume()` is legal, as is `c.resume()` following a `c.resume()`.)

16.1.4 Advancing clocks

An activity may execute the statement

```
next;
```

Execution of this statement blocks until all the clocks that the activity is registered with (if any) have advanced. (The activity implicitly issues a `resume` on all clocks it is registered with before suspending.)

An X10 computation is said to be *quiescent* on a clock `c` if each activity registered with `c` has resumed `c`. Note that once a computation is quiescent on `c`, it will remain quiescent on `c` forever (unless the system takes some action), since no other activity can become registered with `c`. That is, quiescence on a clock is a *stable property*.

Once the implementation has detected quiescence on `c`, the system marks all activities registered with `c` as being able to progress on `c`. An activity blocked on `next` resumes execution once it is marked for progress by all the clocks it is registered with.

16.1.5 Dropping clocks

An activity may drop a clock by executing:

```
c.drop();
```

The activity is no longer considered registered with this clock. A `ClockUseException` is thrown if the activity has already dropped `c`.

16.2 Program equivalences

From the discussion above it should be clear that the following equivalences hold:

```
c.resume(); next; = next; (16.1)
```

```
c.resume(); d.resume(); = d.resume(); c.resume(); (16.2)
```

```
c.resume(); c.resume(); = c.resume(); (16.3)
```

Note that `next; next;` is not the same as `next;`. The first will wait for clocks to advance twice, and the second once.

17 Arrays

This chapter provides an overview of the class `Array`, and its supporting classes `Point`, `Region` and `Dist`. All these classes are in the `x10.lang` package. For more details, please consider the API documentation.

An array is a mapping from a distribution to a range data type. Multiple arrays may be declared with the same underlying distribution.

Each array has a field `a.dist` which may be used to obtain the underlying distribution.

17.1 Points

Arrays are indexed by points— n -dimensional tuples of integers, implemented by the class `x10.lang.Point`. X10 specifies a simple syntax for the construction of points. A rail constructor (§12.27) of type `ValRail[Int]` of length n can be implicitly coerced to a `Point` of rank n . For example, the following code initializes `p` to a point of rank two using a rail constructor:

CHECK: *Huh? Why do we say rail constructor above? As opposed to just any old valrail?*

```
p: Point(2) = [1,2];
```

The `Point` factory method `make` can take a rail constructor as argument. The assignment above can be written, without implicit coercion, as:

```
p: Point = Point.make([1,2]);
```

Points implement the function type `(Int) => Int`; thus, the i th element of a point `p` may be accessed as `p(i)`.

17.2 Regions

A region is a set of points. X10 provides a built-in class, `x10.lang.Region`, to allow the creation of new regions and to perform operations on regions.

Each region `R` has a constant integer rank, `R.rank`.

Here are several examples of region declarations:

```
val MAX_HEIGHT=20;
val Null = Region.makeUnit(); // Empty 0-dimensional region
val N = 10;
val K = 2;
val R1 = 1..100; // 1-dim region with extent 1..100
val R2 = [1..100] as Region(1); // same as R1
val R3 = (0..99) * (-1..MAX_HEIGHT);
val R4 = [0..99, -1..MAX_HEIGHT] as Region(2); // same as R3
val R5 = Region.makeUpperTriangular(N);
val R7 = R4 && R5; // intersection of two regions
val R8 = R4 || R5; // union of two regions
```

The expression $a_1..a_2$ is shorthand for the rectangular, rank-1 region consisting of the points $\{[a_1], \dots, [a_2]\}$. Each subexpression of a_i must be of type `Int`. If a_1 is greater than a_2 , the region is empty.

A region may be constructed by converting from a rail of regions or a rail of points, typically using a rail constructor (§12.27) (e.g., `R4` above). The region constructed from a rail of points represents the region containing just those points. The region constructed from a rail of regions represents the Cartesian product of each of the arguments. X10 v2.0 does not (yet) support hierarchical regions.

Various built-in regions are provided through factory methods on `Region`. For instance:

- `Region.makeUpperTriangular(N)` returns a region corresponding to the non-zero indices in an upper-triangular $N \times N$ matrix.
- `Region.makeLowerTriangular(N)` returns a region corresponding to the non-zero indices in a lower-triangular $N \times N$ matrix.

All the points in a region are ordered canonically by the lexicographic total order. Thus the points of a region $R=(1..2)*(1..2)$ are ordered as

$(1,1), (1,2), (2,1), (2,2)$

Sequential iteration statements such as `for` (§13.12) iterate over the points in a region in the canonical order.

A region is said to be *rectangular* if it is of the form $(T_1 * \dots * T_k)$ for some set of regions T_i . Such a region satisfies the property that if two points p_1 and p_3 are in the region, then so is every point p_2 between them (that is, it is *convex*). (Note that `||` may produce non-convex regions from convex regions, e.g., `[1,1] || [3,3]` is a non-convex region.)

For each region R , the *rectangular closure* of R is the smallest rectangular region enclosing R . For each integer i less than $R.\text{rank}$, the term $R(i)$ represents the enumeration in the i th dimension of the rectangular closure of R . It may be used in a type expression wherever an enumeration may be used.

17.2.1 Operations on regions

Various non side-effecting operators (i.e., pure functions) are provided on regions. These allow the programmer to express sparse as well as dense regions.

Let R be a region. A subset of R is also called a *sub-region*.

Let R_1 and R_2 be two regions whose type establishes that they are of the same rank. Let S be a region of unrelated rank.

$R_1 \ \&\& \ R_2$ is the intersection of R_1 and R_2 .

$R_1 \ || \ R_2$ is the union of the R_1 and R_2 .

$R_1 \ - \ R_2$ is the set difference of R_1 and R_2 .

$R_1 \ * \ S$ is the Cartesian product of R_1 and S , formed by pairing each point in R_1 with every the point in S . Thus, `([1..2,3..4] as Region 2) * (5..6)` is the region of rank 3 containing the points (x,y,z) where x is 1 or 2, y is 3 or 4, and z is 5 or 6.

For a region R and point p of the same rank $R+p$ and $R-p$ represent the translation of the region with p . That is, point q is in R if and only if point $q+p$ is in $R+p$. (And similarly for $R-p$.)

For more details on the available methods on `Region`, please consult the API documentation.

17.3 Distributions

A *distribution* is a mapping from a region to a set of places. X10 provides a built-in class, `x10.lang.Dist`, to allow the creation of new distributions and to perform operations on distributions. This class is `final` in X10 v2.0; future versions of the language may permit user-definable distributions. Since distributions play a dual role (values as well as types), variables of type `Dist` must be initialized and are implicitly `final`.

The *rank* of a distribution is the rank of the underlying region.

```
R: Region = 1..100;
D: Dist = Dist.makeBlock(R);
D: Dist = Dist.makeCyclic(R);
D: Dist = R -> here;
D: Dist = Dist.random(R);
```

Let D be a distribution. $D.region$ denotes the underlying region. $D.places$ is the set of places constituting the range of D (viewed as a function). Given a point p , the expression $D(p)$ represents the application of D to p , that is, the place that p is mapped to by D . The evaluation of the expression $D(p)$ throws an `ArrayIndexOutOfBoundsException` if p does not lie in the underlying region.

When operated on as a distribution, a region R implicitly behaves as the distribution mapping each item in R to `here` (i.e., $R \rightarrow \text{here}$, see below). Conversely, when used in a context expecting a region, a distribution D should be thought of as standing for $D.region$.

17.3.1 Operations returning distributions

Let R be a region, Q a set of places $\{p_1, \dots, p_k\}$ (enumerated in canonical order), and P a place. All the operations described below may be performed on `Dist.factory`.

Unique distribution The distribution `unique(Q)` is the unique distribution from the region $1..k$ to Q mapping each point i to p_i .

Constant distributions. The distribution $R \rightarrow P$ maps every point in R to P .

Block distributions. The distribution `block(R, Q)` distributes the elements of `R` (in order) over the set of places `Q` in blocks as follows. Let p equal $|R| \div N$ and q equal $|R| \bmod N$, where N is the size of `Q`, and $|R|$ is the size of `R`. The first q places get successive blocks of size $(p + 1)$ and the remaining places get blocks of size p .

The distribution `block(R)` is the same distribution as `block(R, Place.places)`.

Cyclic distributions. The distribution `cyclic(R, Q)` distributes the points in `R` cyclically across places in `Q` in order.

The distribution `cyclic(R)` is the same distribution as `cyclic(R, Place.places)`.

Thus the distribution `cyclic(Place.MAX_PLACES)` provides a 1–1 mapping from the region `Place.MAX_PLACES` to the set of all places and is the same as the distribution `unique(Place.places)`.

Block cyclic distributions. The distribution `blockCyclic(R, N, Q)` distributes the elements of `R` cyclically over the set of places `Q` in blocks of size N .

Arbitrary distributions. The distribution `arbitrary(R, Q)` arbitrarily allocates points in `R` to `Q`. As above, `arbitrary(R)` is the same distribution as `arbitrary(R, Place.places)`.

Domain Restriction. If `D` is a distribution and `R` is a sub-region of `D.region`, then `D | R` represents the restriction of `D` to `R`. The compiler throws an error if it cannot determine that `R` is a sub-region of `D.region`.

Range Restriction. If `D` is a distribution and `P` a place expression, the term `D | P` denotes the sub-distribution of `D` defined over all the points in the region of `D` mapped to `P`.

Note that `D | here` does not necessarily contain adjacent points in `D.region`. For instance, if `D` is a cyclic distribution, `D | here` will typically contain points that are P apart, where P is the number of places. An implementation may find a way to still represent them in contiguous memory, e.g., using a complex arithmetic function to map from the region index to an index into the array.

17.3.2 User-defined distributions

Future versions of X10 may provide user-defined distributions, in a way that supports static reasoning.

17.3.3 Operations on distributions

A *sub-distribution* of D is any distribution E defined on some subset of the region of D , which agrees with D on all points in its region. We also say that D is a *super-distribution* of E . A distribution D_1 is *larger than* D_2 if D_1 is a super-distribution of D_2 .

Let D_1 and D_2 be two distributions.

Intersection of distributions. $D_1 \ \&\& \ D_2$, the intersection of D_1 and D_2 , is the largest common sub-distribution of D_1 and D_2 .

Asymmetric union of distributions. $D_1 \ \text{.overlay}(D_2)$, the asymmetric union of D_1 and D_2 , is the distribution whose region is the union of the regions of D_1 and D_2 , and whose value at each point p in its region is $D_2(p)$ if p lies in $D_2 \ \text{.region}$ otherwise it is $D_1(p)$. (D_1 provides the defaults.)

Disjoint union of distributions. $D_1 \ \|\ D_2$, the disjoint union of D_1 and D_2 , is defined only if the regions of D_1 and D_2 are disjoint. Its value is $D_1 \ \text{.overlay}(D_2)$ (or equivalently $D_2 \ \text{.overlay}(D_1)$). (It is the least super-distribution of D_1 and D_2 .)

Difference of distributions. $D_1 \ - \ D_2$ is the largest sub-distribution of D_1 whose region is disjoint from that of D_2 .

17.3.4 Example

```
def dotProduct(a: Array[T](D), b: Array[T](D)): Array[Double](D) =
  (new Array[T]([1:D.places],
    (Point) => (new Array[T](D | here,
      (i): Point) => a(i)*b(i)).sum()))).sum();
```

This code returns the inner product of two *T* vectors defined over the same (otherwise unknown) distribution. The result is the sum reduction of an array of *T* with one element at each place in the range of *D*. The value of this array at each point is the sum reduction of the array formed by multiplying the corresponding elements of *a* and *b* in the local sub-array at the current place.

17.4 Array initializer

Arrays are instantiated by invoking one of the *make* factory methods of the *Array* class.

An array creation must take either an *Int* as an argument or a *Dist*. In the first case an array is created over the distribution `[0:N-1]->here`; in the second over the given distribution.

An array creation operation may also specify an initializer function. The function is applied in parallel at all points in the domain of the distribution. The array construction operation terminates locally only when the array has been fully created and initialized (at all places in the range of the distribution).

For instance:

```
val data : Array[Int]
  = Array.make[Int](1..1000->here, ((i):Point) => i);
val data2 : Array[Int]
  = Array.make[Int]([1..1000,1..1000]->here, ((i,j):Point) => i*j);
```

The first declaration stores in *data* a reference to a mutable array with 1000 elements each of which is located in the same place as the array. Each array component is initialized to *i*.

The second declaration stores in *data2* a reference to a mutable 2-d array over `[1..1000, 1..1000]` initialized with *i*j* at point `[i,j]`.

Other examples:

```
val D1:Dist(1) = ...; /* An expression that creates a Dist */
val D2:Dist(2) = ...; /* An expression that creates a Dist */

val data : Array[Int]
  = Array.make[Int](1000, ((i):Point) => i*i);

val data2 : Array[Float]
```

```

    = Array.make[Float](D1, ((i):Point) => i*i as Float);

    val result : Array[Float]
      = Array.make[Float](D2, ((i,j):Point) => i+j as Float);;

```

17.5 Operations on arrays

In the following let a be an array with distribution D and base type T .

17.5.1 Element operations

The value of a at a point p in its region of definition is obtained by using the indexing operation $a(p)$. This operation may be used on the left hand side of an assignment operation to update the value. The operator assignments $a(i) \text{ op} = e$ are also available in X10.

For array variables, the right-hand-side of an assignment must have the same distribution D as an array being assigned. This assignment involves control communication between the sites hosting D . Each site performs the assignment(s) of array components locally. The assignment terminates when assignment has terminated at all sites hosting D .

17.5.2 Constant promotion

For a distribution D and a constant or final variable v of type T the expression `new Array[T](D, (p: Point) => v)` denotes the mutable array with distribution D and base type T initialized with v at every point.

17.5.3 Restriction of an array

Let $D1$ be a sub-distribution of D . Then $a \mid D1$ represents the sub-array of a with the distribution $D1$.

Recall that a rich set of operators are available on distributions (§17.3) to obtain sub-distributions (e.g. restricting to a sub-region, to a specific place etc).

17.5.4 Assembling an array

Let a_1, a_2 be arrays of the same base type T defined over distributions D_1 and D_2 respectively. Assume that both arrays are value or reference arrays.

Assembling arrays over disjoint regions If D_1 and D_2 are disjoint then the expression $a_1 \parallel a_2$ denotes the unique array of base type T defined over the distribution $D_1 \parallel D_2$ such that its value at point p is $a_1(p)$ if p lies in D_1 and $a_2(p)$ otherwise. This array is a reference (value) array if a_1 is.

Overlaying an array on another The expression $a_1.\text{overlay}(a_2)$ (read: the array a_1 *overlaid with* a_2) represents an array whose underlying region is the union of that of a_1 and a_2 and whose distribution maps each point p in this region to $D_2(p)$ if that is defined and to $D_1(p)$ otherwise. The value $a_1.\text{overlay}(a_2)(p)$ is $a_2(p)$ if it is defined and $a_1(p)$ otherwise.

This array is a reference (value) array if a_1 is.

The expression $a_1.\text{update}(a_2)$ updates the array a_1 in place with the result of $a_1.\text{overlay}(a_2)$.

17.5.5 Global operations

Pointwise operations The unary `lift` operation applies a function to each element of an array, returning a new array with the same distribution. The `lift` operation is implemented by the following method in `Array[T]`:

```
def lift[S](f: (T) => S): Array[S](dist);
```

The binary `lift` operation takes a binary function and another array over the same distribution and applies the function pointwise to corresponding elements of the two arrays, returning a new array with the same distribution. The `lift` operation is implemented by the following method in `Array[T]`:

```
def lift[S,R](f: (T,S) => R, Array[S](dist)): Array[R](dist);
```

Reductions Let f be a function of type $(T, T) \Rightarrow T$. Let a be a value or reference array over base type T . Let $unit$ be a value of type T . Then the operation $a.reduce(f, unit)$ returns a value of type T obtained by performing f on all points in a in some order, and in parallel. The function f must be associative and commutative. The value $unit$ should satisfy $f(unit, x) == x == f(x, unit)$.

This operation involves communication between the places over which the array is distributed. The X10 implementation guarantees that only one value of type T is communicated from a place as part of this reduction process.

Scans Let f be a reduction operator defined on type T . Let a be a value or reference array over base type T and distribution D . Then the operation $a || f()$ returns an array of base type T and distribution D whose i th element (in canonical order) is obtained by performing the reduction f on the first i elements of a (in canonical order).

This operation involves communication between the places over which the array is distributed. The X10 implementation will endeavour to minimize the communication between places to implement this operation.

Other operations on arrays may be found in `x10.lang.Array` and other related classes.

18 Annotations and compiler plugins

X10 provides an annotation system and compiler plugin system for to allow the compiler to be extended with new static analyses and new transformations.

Annotations are interface types that decorate the abstract syntax tree of an X10 program. The X10 type-checker ensures that an annotation is a legal interface type. In X10, interfaces may declare both methods and properties. Therefore, like any interface type, an annotation may instantiate one or more of its interface's properties. Unlike with Java annotations, property initializers need not be compile-time constants; however, a given compiler plugin may do additional checks to constrain the allowable initializer expressions. The X10 type-checker does not check that all properties of an annotation are initialized, although this could be enforced by a compiler plugin.

18.1 Annotation syntax

The annotation syntax consists of an “@” followed by an interface type.

Annotation ::= @InterfaceBaseType Constraints?

Annotations can be applied to most syntactic constructs in the language including class declarations, constructors, methods, field declarations, local variable declarations and formal parameters, statements, expressions, and types. Multiple occurrences of the same annotation (i.e., multiple annotations with the same interface type) on the same entity are permitted.


```

    ClassModifier ::= Annotation
    InterfaceModifier ::= Annotation
    FieldModifier ::= Annotation
    MethodModifier ::= Annotation
    VariableModifier ::= Annotation
    ConstructorModifier ::= Annotation
    AbstractMethodModifier ::= Annotation
    ConstantModifier ::= Annotation
    Type ::= AnnotatedType
    AnnotatedType ::= Annotation+ Type
    Statement ::= AnnotatedStatement
    AnnotatedStatement ::= Annotation+ Statement
    Expression ::= AnnotatedExpression
    AnnotatedExpression ::= Annotation+ Expression

```

Recall that interface types may have dependent parameters.

The following examples illustrate the syntax:

- Declaration annotations:

```

// class annotation
@Value
class Cons { ... }

// method annotation
@PreCondition(0 <= i && i < this.size)
public def get(i: Int): Object { ... }

// constructor annotation
@Where(x != null)
def this(x: T) { ... }

// constructor return type annotation
def this(x: T): C@Initialized { ... }

// variable annotation
@Unique x: A;

```

- Type annotations:

```
List@Nonempty
```

```
Int@Range(1,4)
```

```
Array[Array[Double]]@Size(n * n)
```

- Expression annotations:

```
m() : @RemoteCall
```

- Statement annotations:

```
@Atomic { ... }
```

```
@MinIterations(0)
```

```
@MaxIterations(n)
```

```
for (var i: Int = 0; i < n; i++) { ... }
```

```
// An annotated empty statement ;
```

```
@Assert(x < y);
```

18.2 Annotation declarations

Annotations are declared as interfaces. They must be subtypes of the interface `x10.lang.annotation.Annotation`. Annotations on types, expressions, statements, classes, fields, methods, constructors, and local variable declarations (or formal parameters) must extend `ExpressionAnnotation`, `StatementAnnotation`, `ClassAnnotation`, `FieldAnnotation`, `MethodAnnotation`, `ConstructorAnnotation`, and `VariableAnnotation`, respectively.

18.3 Compiler plugins

After the base X10 semantic checking is completed, compiler plugins are loaded and run. Plugins may perform any number of compiler passes to implement additional semantic checking and code transformations, including transformations

using the abstract syntax of the annotations themselves. Plugins should output valid X10 abstract syntax trees.

Plugins are implemented in Java as Polyglot [8] passes applied to the AST after normal base X10 type checking. Plugins to run are specified on the command-line. The order of execution is determined by the Polyglot pass scheduler.

To run compiler plugins, add the command-line option:

```
-PLUGINS=P1,P2,...,Pn
```

where P1, P2, ..., Pn are classes that implement the `CompilerPlugin` interface:

```
package polyglot.ext.x10.plugin;

import polyglot.ext.x10.ExtensionInfo;
import polyglot.frontend.Job;
import polyglot.frontend.goals.Goal;

public interface CompilerPlugin {
    public Goal
        register(ExtensionInfo extInfo, Job job);
}
```

The `Goal` object returned by the `register` method specifies dependencies on other passes. Documentation for Polyglot can be found at:

<http://www.cs.cornell.edu/Projects/polyglot>

Most plugins should implement either `SimpleOnePassPlugin` or `SimpleVisitorPlugin`.

The compiler loads plugin classes from the `x10c` classpath.

Plugins are given access to a Polyglot AST and type system. Annotations are represented in the AST as `Nodes` with the following interface:

```
package polyglot.ext.x10.ast;

public interface AnnotationNode extends Node {
    X10ClassType annotation();
}
```

Annotations for a Node object `n` can be accessed through the node's extension object as follows:

```
List<AnnotationNode> annotations =  
    ((X10Ext) n.ext()).annotations();  
List<X10ClassType> annotationTypes =  
    ((X10Ext) n.ext()).annotationInterfaces();
```

In the type system, `X10TypeObject` has the following method for accessing annotations:

```
List<X10ClassType> annotations();
```

19 Linking with native code

On some platforms, X10 v2.0 supports a simple facility to permit the efficient intra-thread communication of an array of primitive type to code written in the language C. The array must be a “local” array. The primary intent of this design is to permit the reuse of native code that efficiently implements some numeric array/matrix calculation.

Future language releases are expected to support similar bindings to FORTRAN, and to support parallel native processing of distributed X10 arrays.

The interface consists of two parts. First, an array intended to be communicated to native code must be created as an unsafe array via the `unsafe` annotation.

```
new unsafe Array[T](dist)
```

Unsafe arrays can be of any dimension. However, X10 v2.0 requires that unsafe arrays be of a primitive type, and local (i.e., with an underlying distribution that maps all elements in its region to here).

Unsafe arrays are allocated in a special array of memory that permits their efficient transmission to natively linked code.

Second, the X10 programmer may specify that certain methods are to be implemented natively by using the modifier `extern`:

```
MethodModifier ::= extern
```

Such a method must be declared `static` and may not have a method body.¹ Primitive types in the method argument are translated to their corresponding JNI type (e.g., `Float` is translated to `jfloat`, `Double` to `jdouble`, etc.). The only non-primitive type permitted in an `extern` method is an unsafe array. This is passed at type `jlong` as an eight-byte address into the unsafe region that contains the data for the array. Note that `jlong` is not the same as `long` on 32-bit machines.

¹This restriction is likely to be lifted in the future.

Since only the starting address of an array is passed, if the array is multidimensional, the user must explicitly communicate (or have a guarantee of) the rank of the passed array, and must either typecast or explicitly code the address calculation. Note that all X10 arrays are created in row-major order, and so any native routine must also access them in the same order.

For each class `C` that contains an `extern` method, the X10 compiler generates a text file `C_x10stub.c`. This file contains generated `C` stub functions that are called from the `extern` routines. The name of the stub function is derived from the name of the `extern` method. If the method is `C.process()`, the stub function will be `Java_C_C_process()`. The name is suffixed with the signature of the method if the method is overloaded.

The programmer must write `C` code to implement the native method, using the methods in the `C` stub file to call the actual native method. The programmer must compile these files and link them into a dynamically linked library (DLL). Note that the `jni.h` header file must be in the include path. The programmer must ensure this library is loaded by the program before the method is called, e.g., by adding a `x10.lang.System.loadLibrary` call (in a static initializer of the X10 class).

Example 19.0.1 The following class illustrates the use of `unsafe` and native linking.

```
public class IntArrayExternUnsafe {
    public static extern
        def process(yy: unsafe Rail[Int], size: Int);
    static { System.loadLibrary("IntArrayExternUnsafe"); }
    public static def main(args: Rail[String]) {
        val b = (new IntArrayExternUnsafe()).run();
        System.out.println("++++++ Test "
                           +(b?"succeeded.":"failed."));
        System.exit(b?0:1);
    }
    public def run() : Boolean {
        val high = 10;
        val d = (0..high) -> here;
        val y: Array[Int] = new unsafe Array[Int](d);
        for (val (j) in y.region) {
            y(j) = j;
        }
    }
}
```

```

    process(y,high);
    for (val (j) in y.region) {
        val expected = j+100;
        if(y(j) != expected) {
            System.out.println("y("+j+")="
                               +y(j)+" != "+expected);
            return false;
        }
    }
    return true;
}
}

```

The programmer may then write the C code thus:

```

void IntArrayExternUnsafe_process(jlong yy, signed int size) {
    int i;
    int* array = (int*) (long) yy;
    for (i = 0; i < size; i++) {
        array[i] += 100;
    }
}
/* automatically generated in C_x10stub.c */
void
Java_IntArrayExternUnsafe_IntArrayExternUnsafe_process
(JNIEnv *env, jobject obj, jlong yy, jint size) {
    IntArrayExternUnsafe_process(yy, size);
}

```

This code may be linked with the stub file (or textually placed in it). The programmer must then compile and link the C code and ensure that the DLL is on the appropriate classpath.

□

References

- [1] David Bacon. Kava: A Java dialect with a uniform object model for lightweight classes. *Concurrency – Practice and Experience*, 15:185–206, 2003.
- [2] Joseph A. Bank, Barbara Liskov, and Andrew C. Myers. Parameterized types and Java. In *Proceedings of the 24th Annual ACM Symposium on Principles of Programming Languages (POPL'97)*, pages 132–145, 1997.
- [3] William Carlson, Tarek El-Ghazawi, Bob Numrich, and Kathy Yelick. Programming in the Partitioned Global Address Space Model, 2003. Presentation at SC 2003, <http://www.gwu.edu/upc/tutorials.html>.
- [4] Bradford L. Chamberlain, Sung-Eun Choi, Steven J. Deitz, and Lawrence Snyder. The high-level parallel language ZPL improves productivity and performance. In *Proceedings of the IEEE International Workshop on Productivity and Performance in High-End Computing*, 2004.
- [5] J. Gosling, W. Joy, G. Steele, and G. Bracha. *The Java Language Specification*. Addison Wesley, 2000.
- [6] Maurice Herlihy. Wait-free synchronization. *ACM Transactions on Programming Languages and Systems*, 13(1):124–149, January 1991.
- [7] Jose E. Moreira, Samuel P. Midkiff, Manish Gupta, Pedro V. Artigas, Marc Snir, and Richard D. Lawrence. Java programming for high-performance numerical computing. *IBM Systems Journal*, 39(1):21–, 2000.
- [8] Nathaniel Nystrom, Michael R. Clarkson, and Andrew C. Myers. Polyglot: An extensible compiler framework for Java. In Görel Hedin, editor, *Compiler Construction, 12th International Conference, CC 2003*, number 2622 in LNCS, pages 138–152, Warsaw, Poland, April 2003. Springer-Verlag.

- [9] V. Saraswat and R. Jagadeesan. Concurrent clustered programming. In *Concur '05, to appear*, 2005.
- [10] A. Skjellum, E. Lusk, and W. Gropp. *Using MPI: Portable Parallel Programming with the Message Passing Interface*. MIT Press, 1999.
- [11] K. A. Yelick, L. Semenzato, G. Pike, C. Miyamoto, B. Liblit, A. Krishnamurthy, P. N. Hilfinger, S. L. Graham, D. Gay, P. Colella, and A. Aiken. Titanium: A high-performance java dialect. *Concurrency - Practice and Experience*, 10(11-13):825–836, 1998.

- `==`, 91
- `Array`, 37
- `Object`, 22, 36
- `String`, 37
- `ValArray`, 37
- `as`, 34
- `do`, 99
- `foreach`, 110
- `for`, 99
- `instanceof`, 92
- `while`, 99
- `x10.lang.Array`, 37
- `x10.lang.Object`, 22, 36
- `x10.lang.String`, 37

- `AnnotationNode`, 135
- `annotations`, 131
 - type annotations, 32
- `Any`
 - structs, 73
- `array`
 - access, 128
 - pointwise operations, 129
 - reductions, 130
 - restriction, 128
 - scans, 130
 - union
 - asymmetric, 129
 - disjoint, 129
- `array types`, 37
- `arrays`, 121
 - constant promotion, 128
 - distribution, 121
- `assignment`, 86
- `atomic`, 67

- atomic blocks, 112
- auto-boxing
 - struct to interface, 73
- casting, 34
- class, 21, 53, 55
 - reference class, 22
- class declaration, 21
- Class invariant, 56
- class invariants, 55
- classcast, 92
- clock
 - clocked statements, 119
 - ClockUseException, 109, 119
 - creation, 118
 - drop, 120
 - next, 120
 - resume, 119
- clocks, 117
- closures
 - parametrized closures, 23
- coercions, 34
 - explicit coercion, 34
 - subsumption coercion, 34
- CompilerPlugin, 134
- constrained types, 26
- constructors
 - parametrized constructors, 23
- conversions, 34
 - numeric conversions, 35
 - string conversion, 35
 - user defined, 35
 - widening conversions, 35
- declaration
 - class declaration, 21
 - interface declaration, 22
 - reference class declaration, 21

- declarations
 - type definitions, 25
- dependent type, consistency, 29
- dependent types, 26
- distribution, 124
 - arbitrary, 125
 - block, 125
 - block cyclic, 125
 - constant, 124
 - cyclic, 125
 - difference, 126
 - intersection, 126
 - restriction
 - range, 125
 - region, 125
 - union
 - asymmetric, 126
 - disjoint, 126
 - unique, 124
 - user-defined, 126
- equality
 - function, 82
- Exception, 100
 - model, 106
 - unchecked, 31
- expressions, 83
- extends clause, 57
- extern, 136
- final variable, 42
- finish, 108
- function
 - ==, 82
 - at(Object), 82
 - at(Place), 82
 - equality, 82
 - equals, 82

- hashCode, 82
- home, 82
- method selector, 80
- operator, 81
- toString, 82
- typeName, 82
- types, 31
- functions, 77
- generic types, 26
- global
 - field, 64
 - method, 67
- Goal, 134
- guards, 55
- here, 103
- identifier, 17
- implements clause, 57
- import,type definitions, 26
- initial activity, 109
- initialization
 - static, 69
- interface, 22
- interface declaration, 22
- interfaces, 52
- literal
 - function, 78
- literals, 18
- locality condition, 113
- method
 - underlying function, 80
- methods
 - parametrized methods, 23
- names, 51
- Node, 134

- nonblocking
 - method, 68
- nullary constructor, 42
- numeric promotion, 87
- Object, 46
- packages, 51
- parameter, 29
- parametrized closures, 23
- parametrized constructors, 23
- parametrized methods, 23
- pinned
 - method, 68
- place types, 27
- place.home, 46
- places, 102
- placetype, 27
- plugins, 133
- point syntax, 121
- Polyglot, 134
- promotion, 87
- properties, 22
- property
 - call, 54
 - initialization, 54
- proto, 59
 - rules, 61
- Qualifier
 - field, 64
 - method, 67
- reference class type, 22
- region, 122
 - banded, 122
 - convex, 123
 - intersection, 123
 - lowerTriangular, 122

- product, 123
- set difference, 123
- sub-region, 123
- union, 123
- upperTriangular, 122
- ReturnStatement, 101
- root activity, 105
- safe
 - method, 69
- sequential
 - method, 68
- statements, 95
- structs, 71
- sub-distribution, 126
- subtyping, 32
- throw, 100
- type equivalence, 32
- type inference, 38
- type invariants, 55
- type-checking
 - extends clause, 57
 - implements clause, 57
- types, 19
 - annotated types, 32
 - class types, 21
 - constrained types, 26
 - dependent types, 26
 - function types, 31
 - generic types, 23, 26
 - inference, 38
 - interface types, 22
 - type definitions, 25
 - type parameters, 23
- variable
 - final, 42
- variable declaration, 41

- variable declarator
 - destructuring, 43
- variables, 41
- Void, 39

A Change Log

A.1 Changes from X10 v2.0

- `Any` is now the top of the type hierarchy (every object, struct and function has a type that is a subtype of `Any`). `Any` defines `home`, `at`, `toString`, `typeName`, `equals` and `hashCode`. `Any` also defines the methods of `Equals`, so `Equals` is not needed any more.
- Revised discussion of incomplete types (§9.4).
- The manual has been revised and brought into line with the current implementation.

A.2 Changes from X10 v1.7

The language has changed in the following way:

- **Type system changes:** There are now three kinds of entities in an X10 computation: objects, structs and functions. Their associated types are class types, struct types and function types.

Class and struct types are called *container types* in that they specify a collection of fields and methods. Container types have a name and a signature (the collection of members accessible on that type). Collection types support primitive equality `==` and may support user-defined equality if they implement the `x10.lang.Equals` interface.

Container types (and interface types) may be further qualified with constraints.

A function type specifies a set of arguments and their type, the result type, and (optionally) a guard. A function application type-checks if the arguments are of the given type and the guard is satisfied, and the return value is of the given type. A function type does not permit `==` checks. Closure literals create instances of the corresponding function type.

Container types may implement interfaces and zero or more function types.

All types support a basic set of operations that return a string representation, a type name, and specify the home place of the entity.

The type system is not unitary. However, any type may be used to instantiate a generic type.

There is no longer any notion of value classes. value classes must be re-written into structs or (reference) classes.

- **Global object model:** Objects are instances of classes. Each object is associated with a globally unique identifier. Two objects are considered identical `==` if their ids are identical. Classes may specify `global` fields and methods. These can be accessed at any place. (`global` fields must be immutable.)
- **Proto types.** For the decidability of dependent type checking it is necessary that the property graph is acyclic. This is ensured by enforcing rules on the leakage of `this` in constructors. The rules are flexible enough to permit cycles to be created with normal fields, but not with properties.
- **Place types.** Place types are now implemented. This means that non-global methods can be invoked on a variable, only if the variable's type is either a struct type or a function type, or a class type whose constraint specifies that the object is located in the current place.

There is still no support for statically checking array access bounds, or performing place checks on array accesses.

The X10 language has been developed as part of the IBM PERCS Project, which is supported in part by the Defense Advanced Research Projects Agency (DARPA) under contract No. NBCH30390004.

Java and all Java-based trademarks are trademarks of Sun Microsystems, Inc. in the United States, other countries, or both.

